

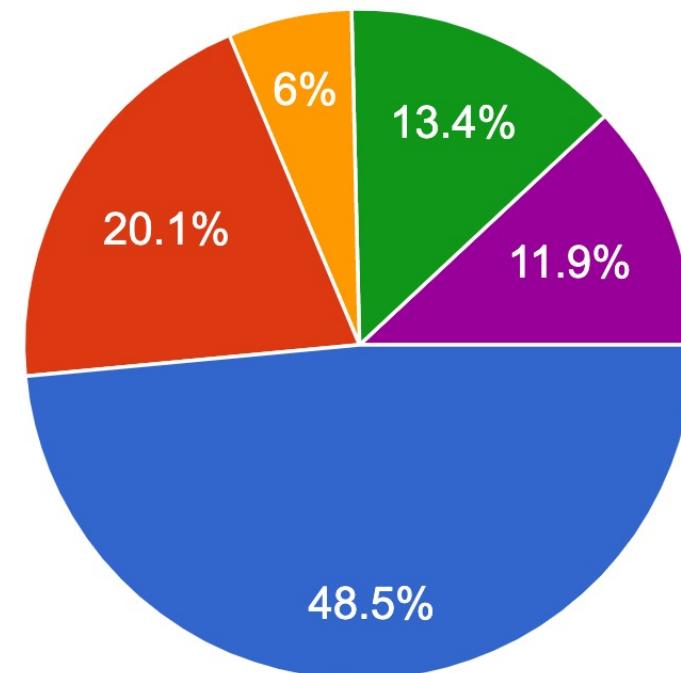
# Lecture 7: Convolutional Networks

# Lecture Format

What is your preferred lecture format?

134 responses

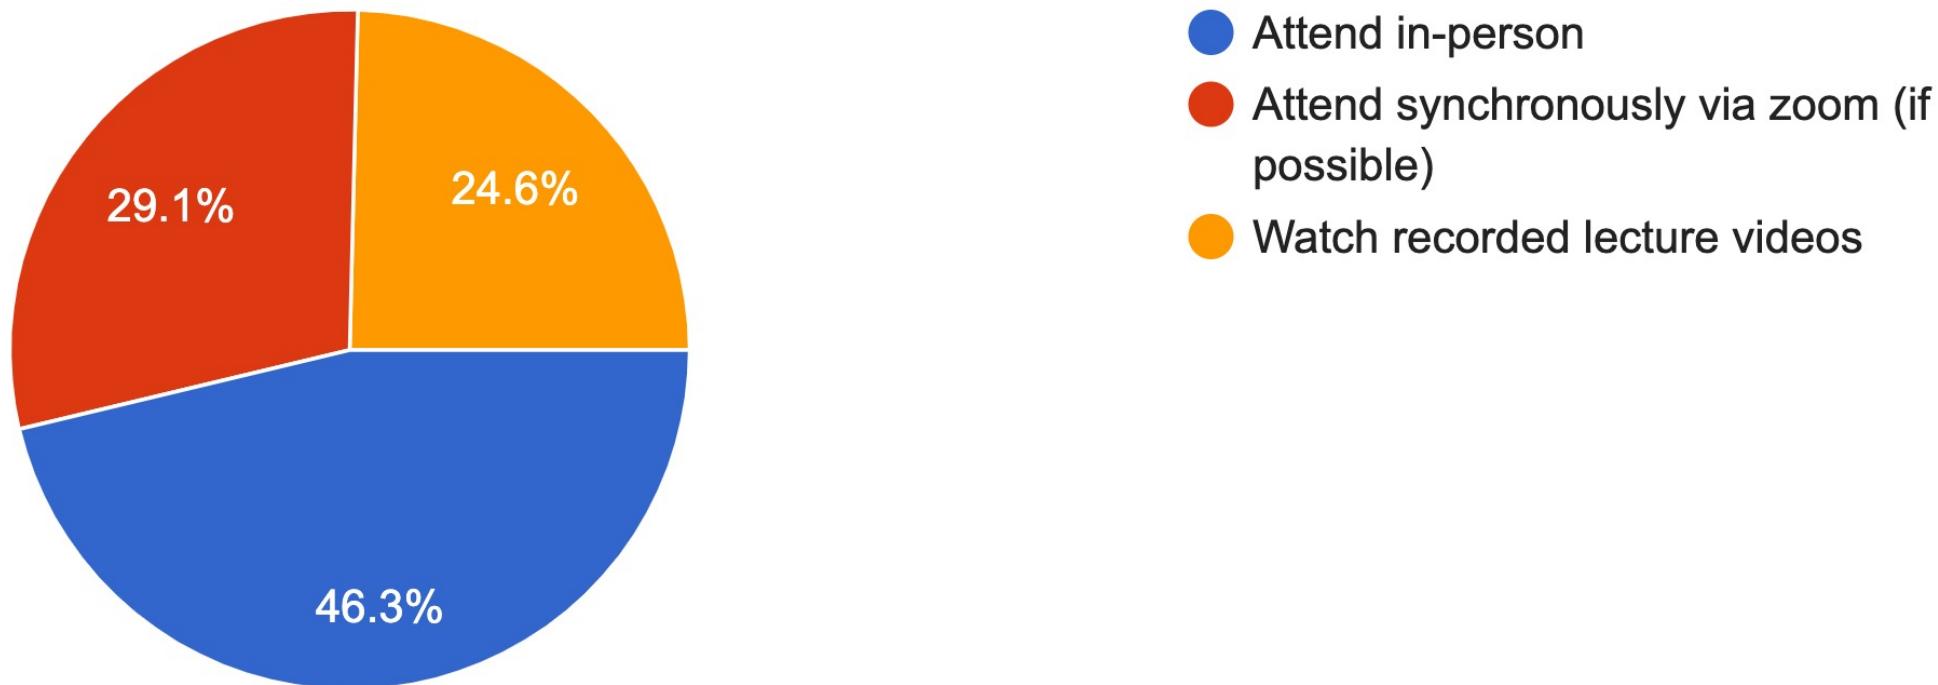
- Strongly prefer remote lectures
- Slightly prefer remote lectures
- Indifferent between in-person and remote lecture
- Slightly prefer in-person lectures
- Strongly prefer in-person lectures



# Lecture Format

If we were to return to in-person lectures, how would you plan to watch lectures?

134 responses



# Lecture Format

- We will remain remote for at least another 2-3 weeks
- Idea: book a conference room for “watch parties?”  
Or just use lecture hall
- COVID in MI have (hopefully!) peaked? If they continue to drop we will consider in-person OH in the next 1-2 weeks
- May revisit after Spring Break
- Feel free to raise hand to ask questions in Zoom!
- Midterm will be remote (but still working on exact format)

Reminder: A2

Due last Friday

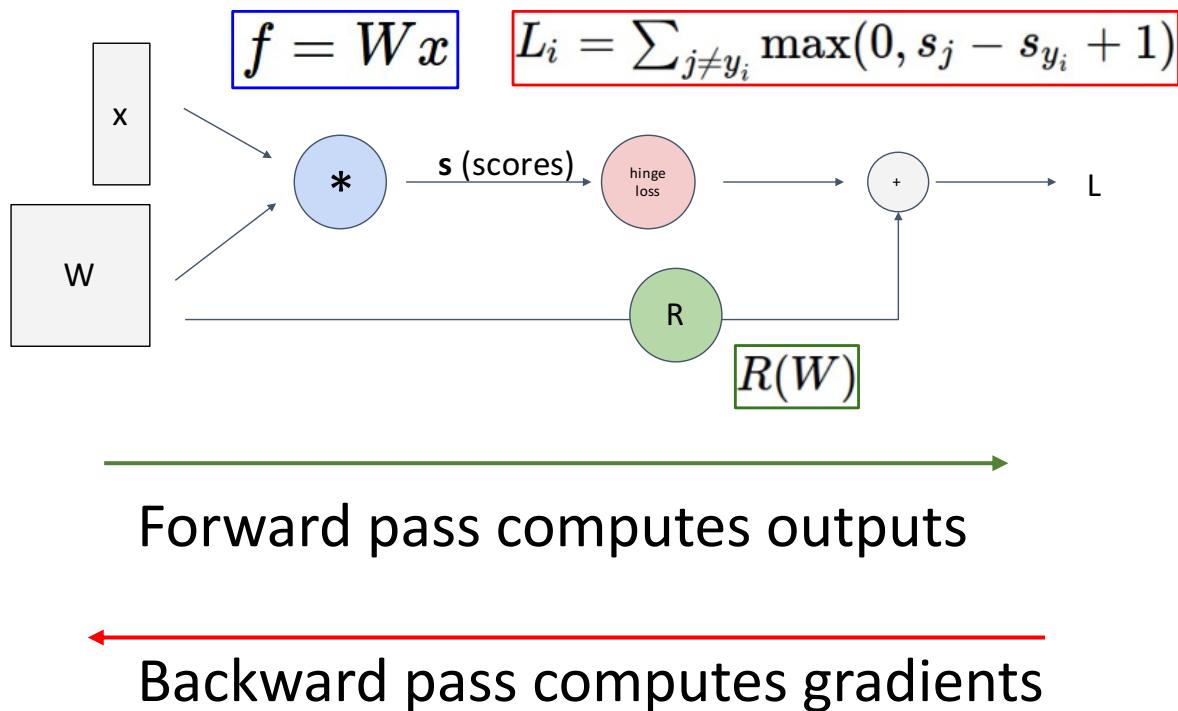
# A3

Will be released tonight, covering:

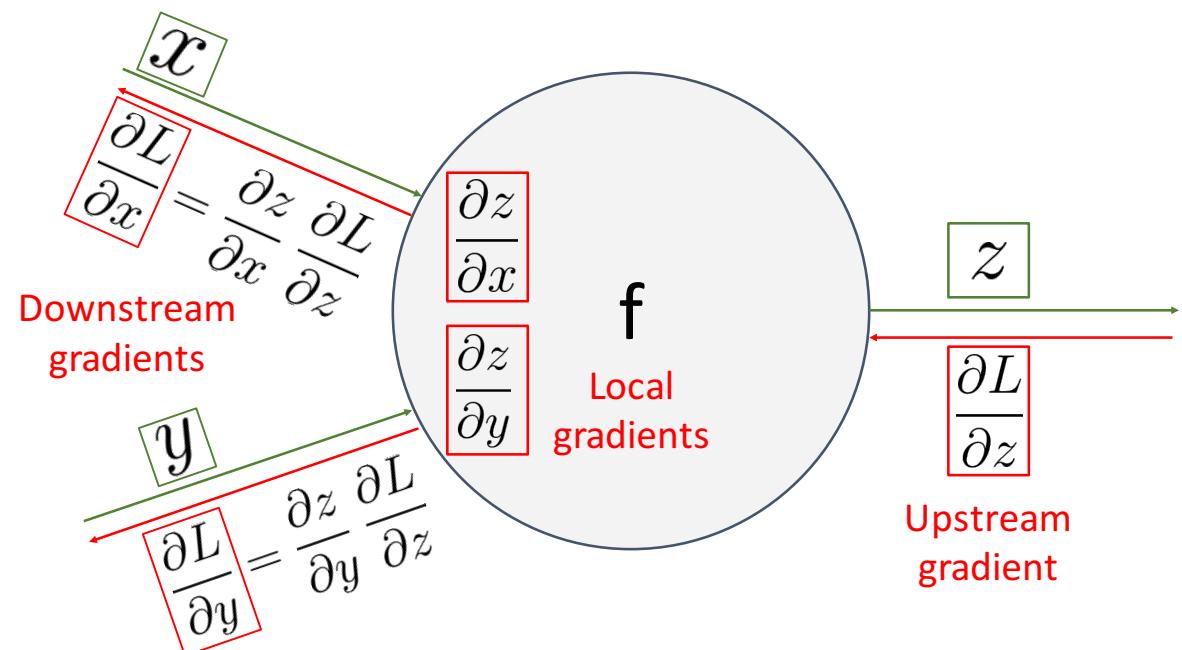
- Backpropagation with modular API
- Different update rules (Momentum, RMSProp, Adam, etc)
- Batch Normalization
- Dropout
- Convolutional Networks

# Last Time: Backpropagation

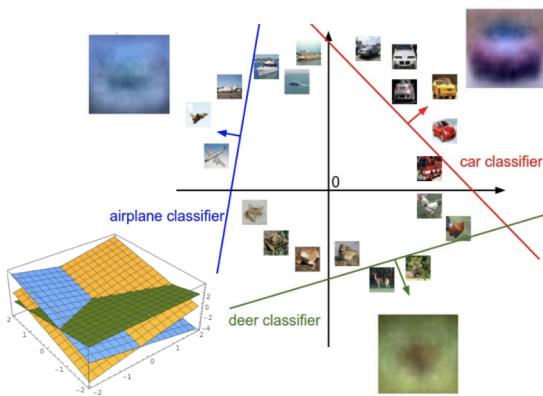
Represent complex expressions  
as **computational graphs**



During the backward pass, each node in the graph receives **upstream gradients** and multiplies them by **local gradients** to compute **downstream gradients**

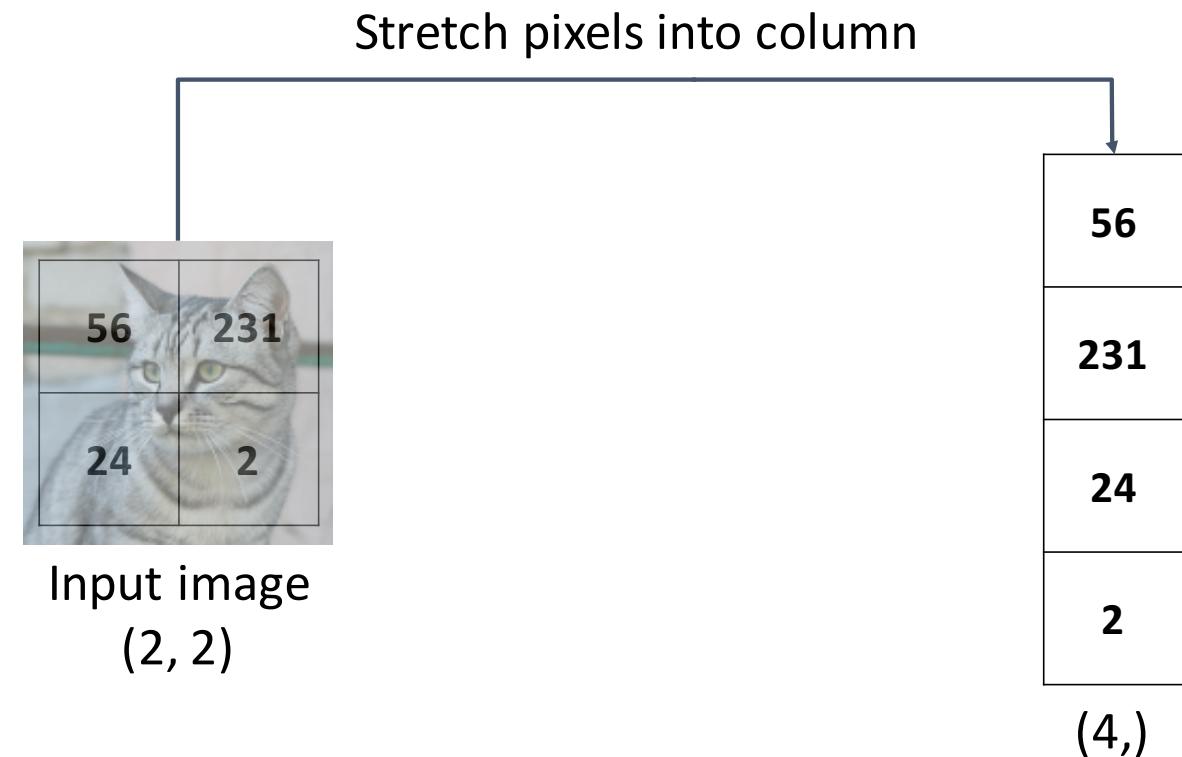
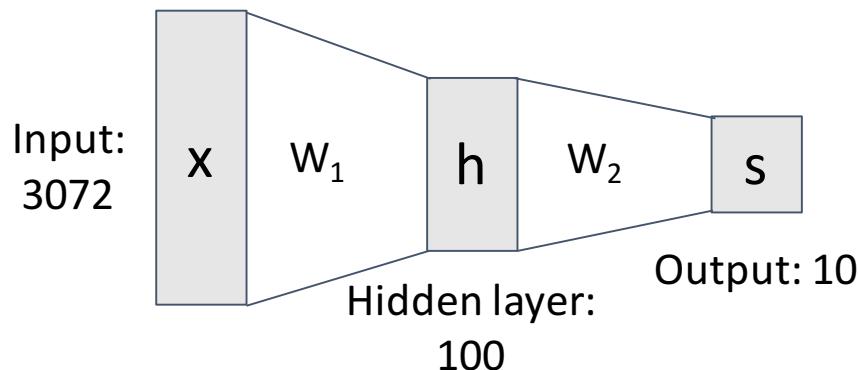


$$f(x, W) = Wx$$

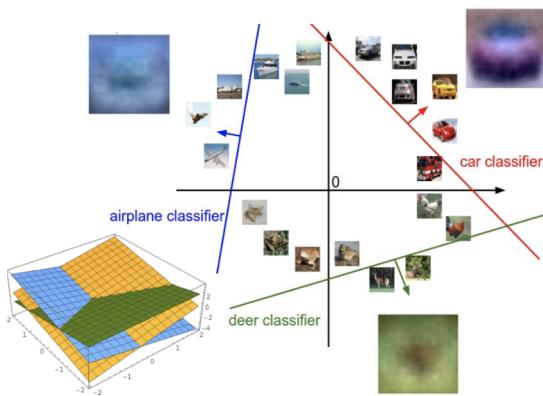


**Problem:** So far our classifiers don't respect the spatial structure of images!

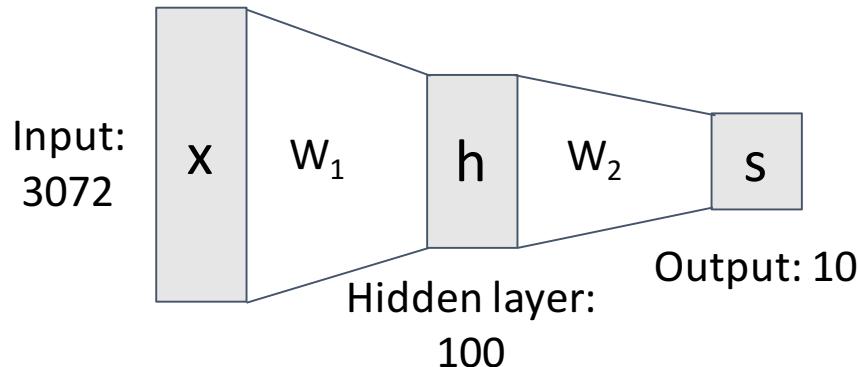
$$f = W_2 \max(0, W_1 x)$$



$$f(x, W) = Wx$$



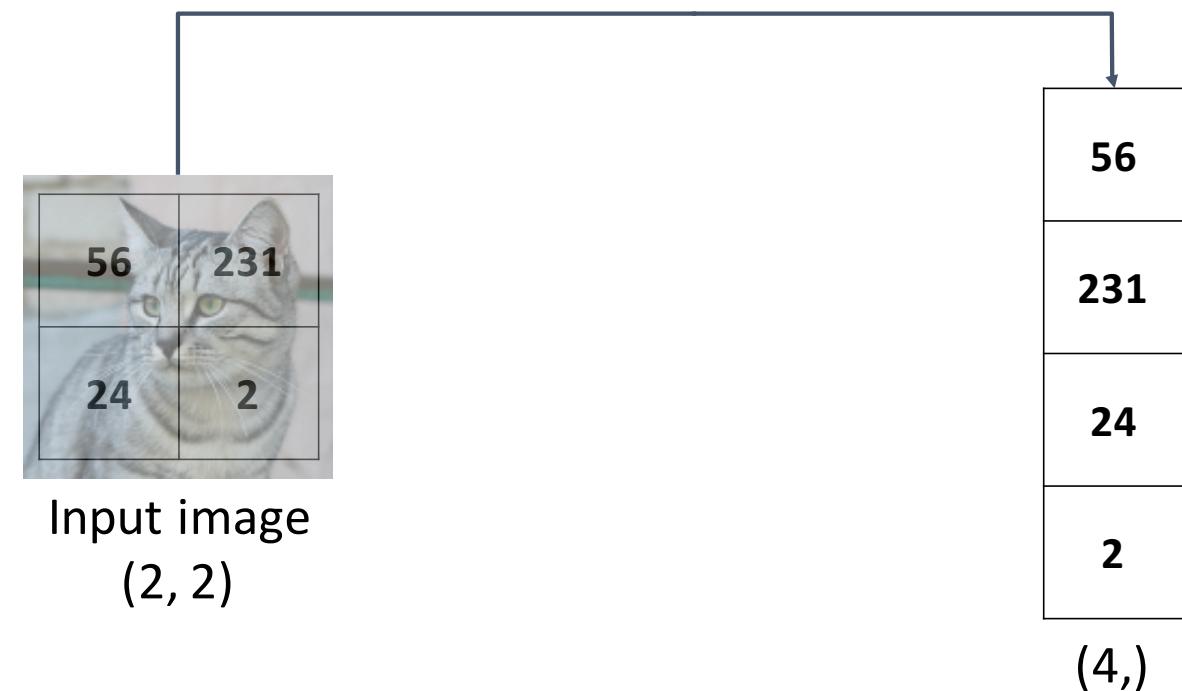
$$f = W_2 \max(0, W_1 x)$$



**Problem:** So far our classifiers don't respect the spatial structure of images!

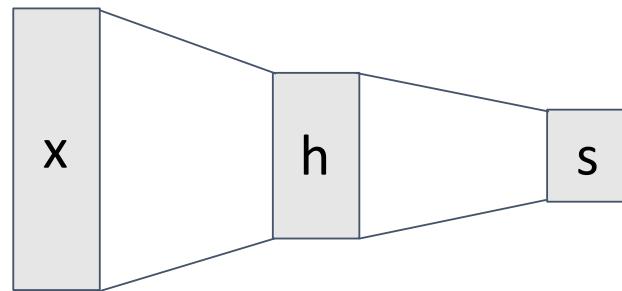
**Solution:** Define new computational nodes that operate on images!

Stretch pixels into column

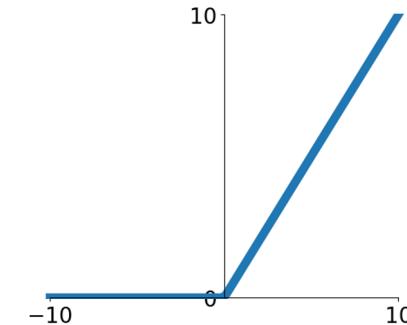


# Components of a Fully-Connected Network

## Fully-Connected Layers

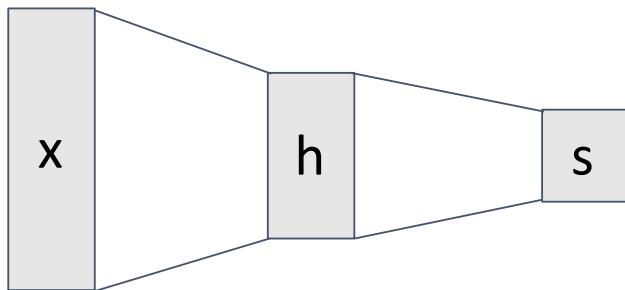


## Activation Function

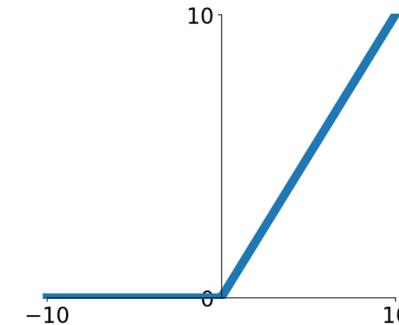


# Components of a Convolutional Network

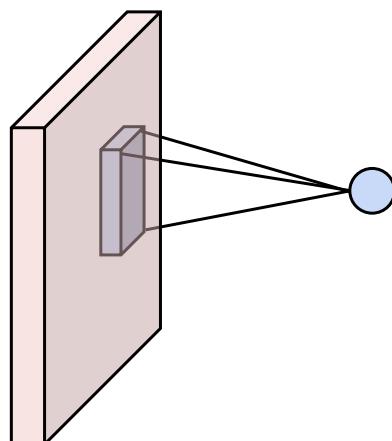
## Fully-Connected Layers



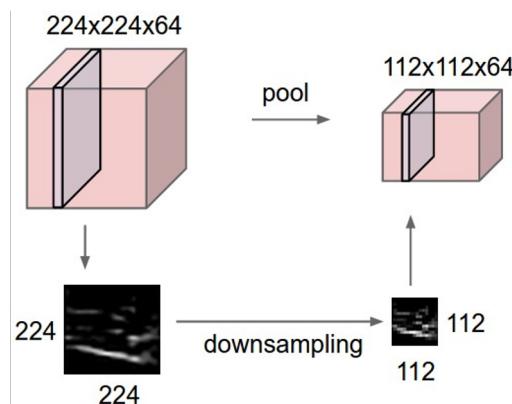
## Activation Function



## Convolution Layers



## Pooling Layers

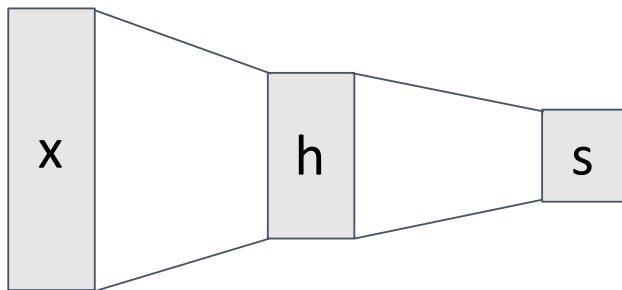


## Normalization

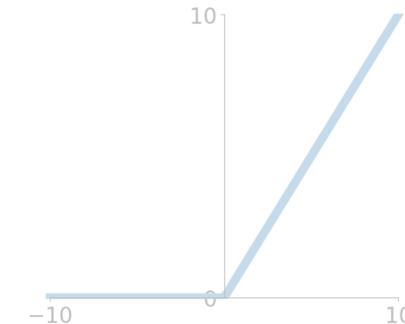
$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}}$$

# Components of a Convolutional Network

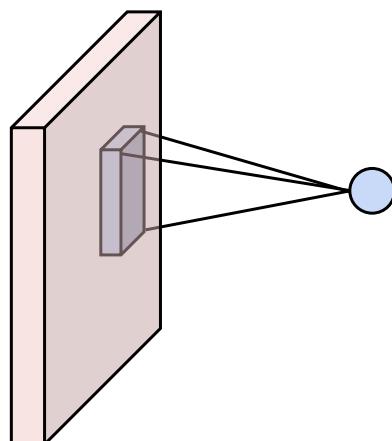
## Fully-Connected Layers



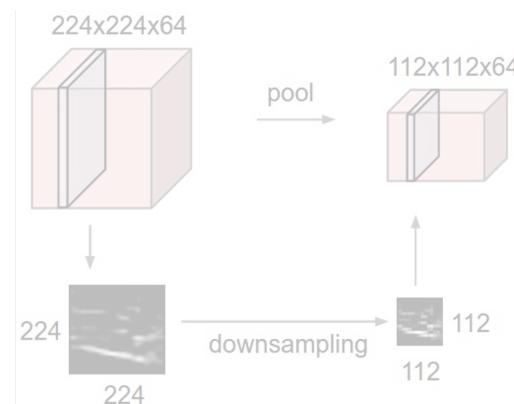
## Activation Function



## Convolution Layers



## Pooling Layers

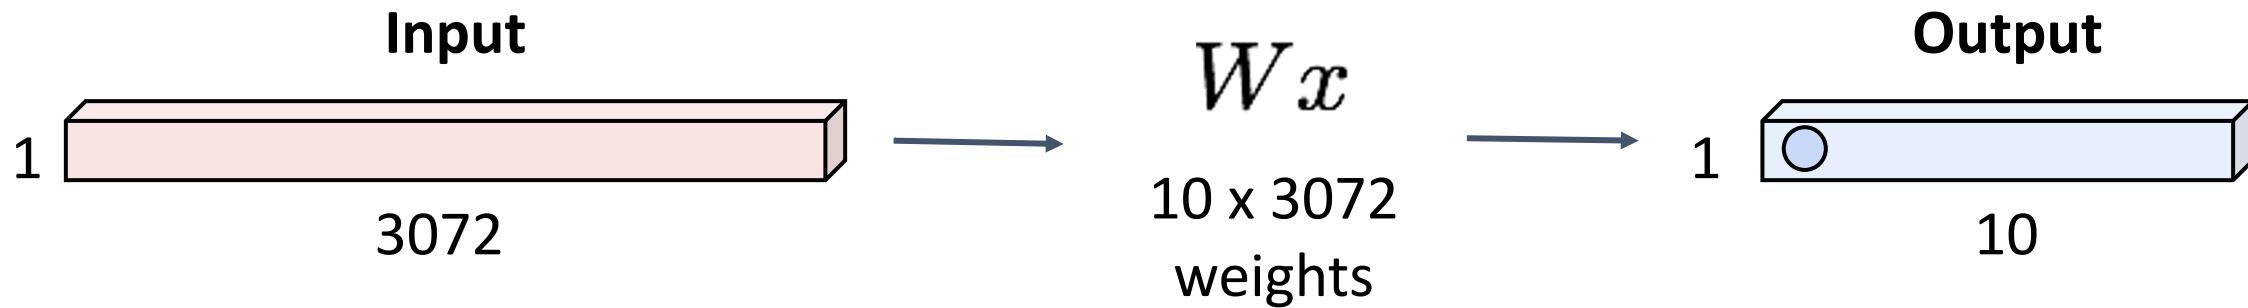


## Normalization

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}}$$

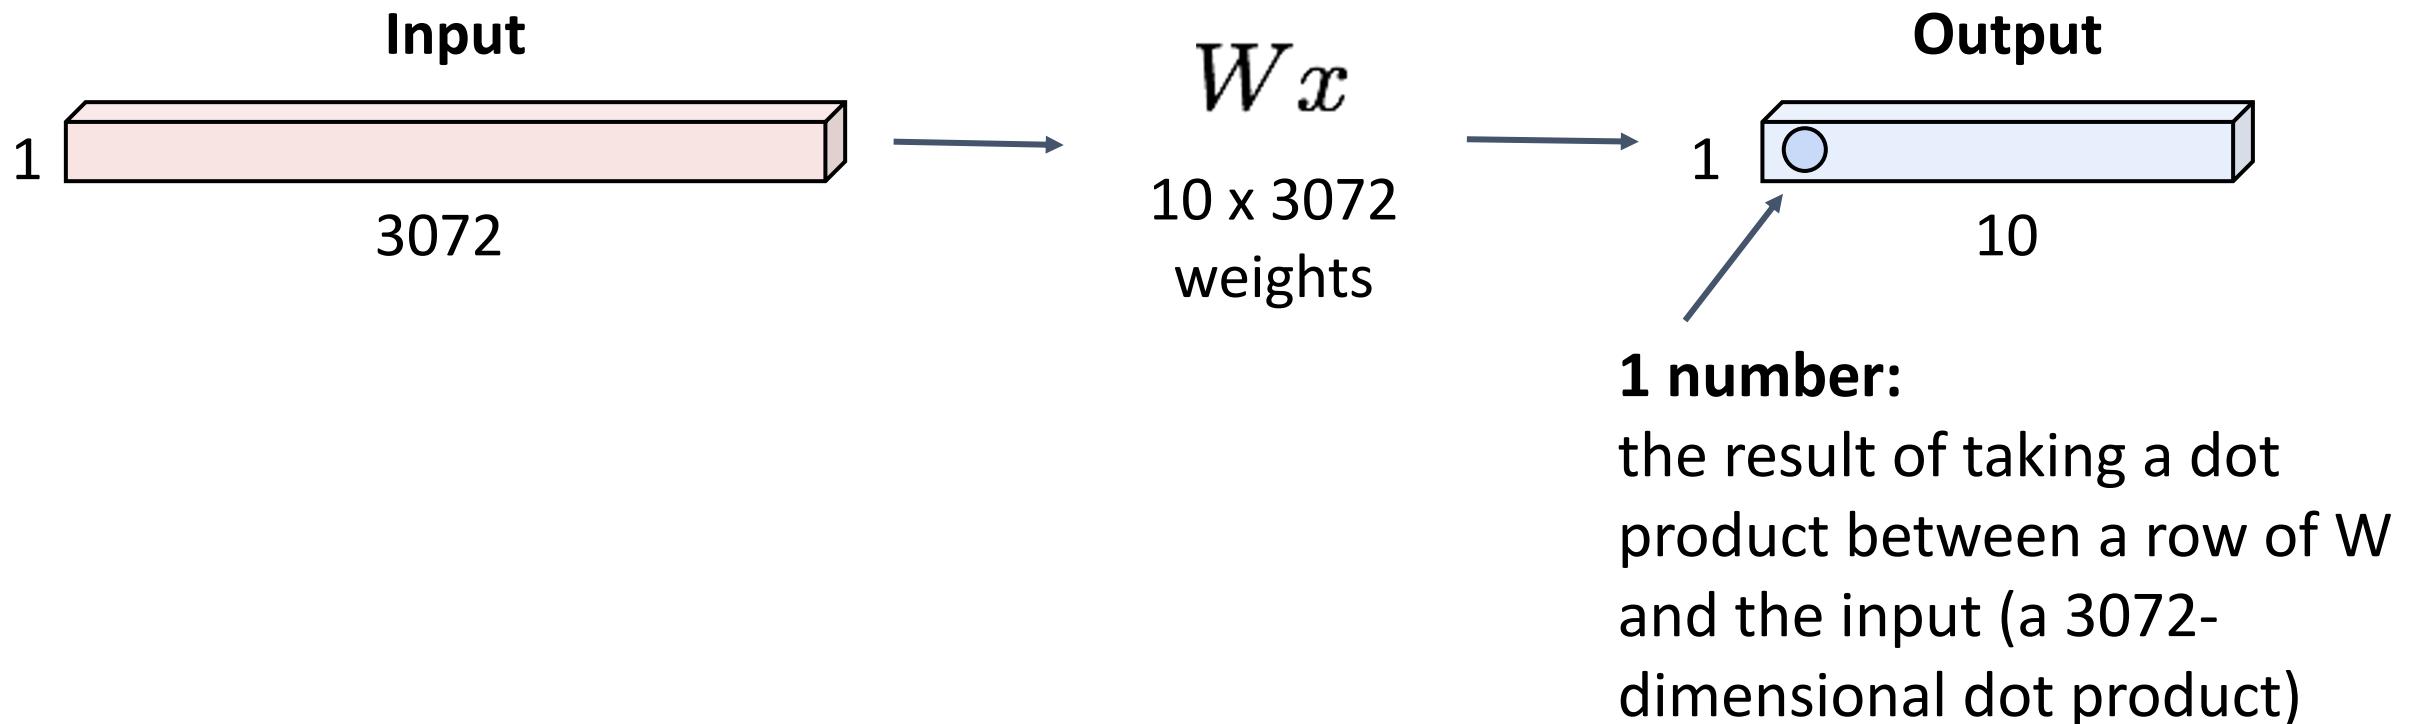
# Fully-Connected Layer

32x32x3 image -> stretch to 3072 x 1



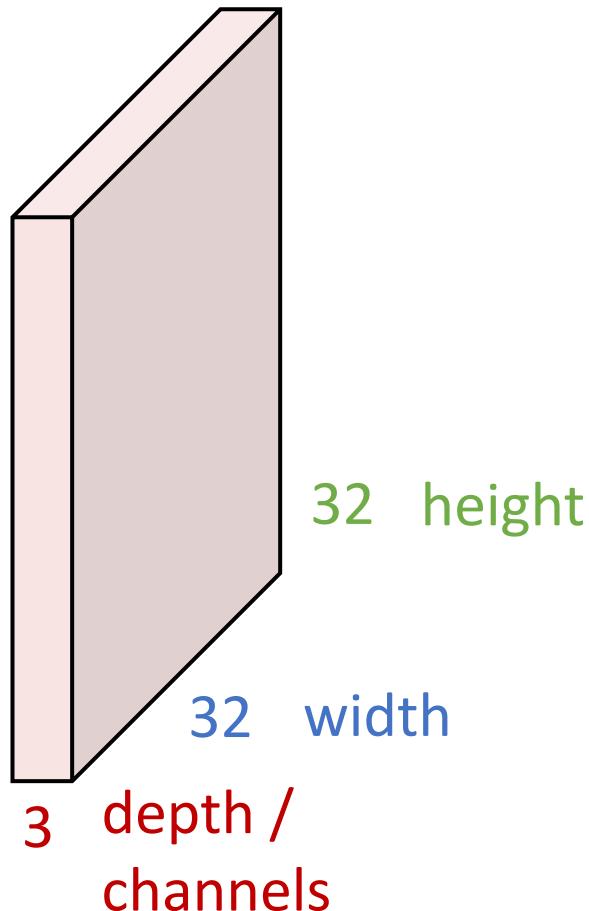
# Fully-Connected Layer

32x32x3 image -> stretch to 3072 x 1



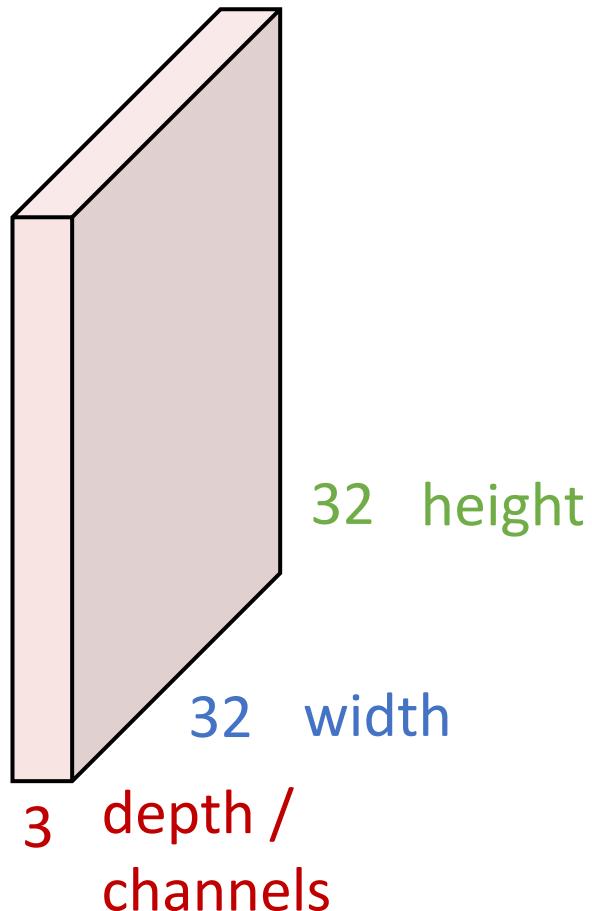
# Convolution Layer

$3 \times 32 \times 32$  image: preserve spatial structure



# Convolution Layer

$3 \times 32 \times 32$  image



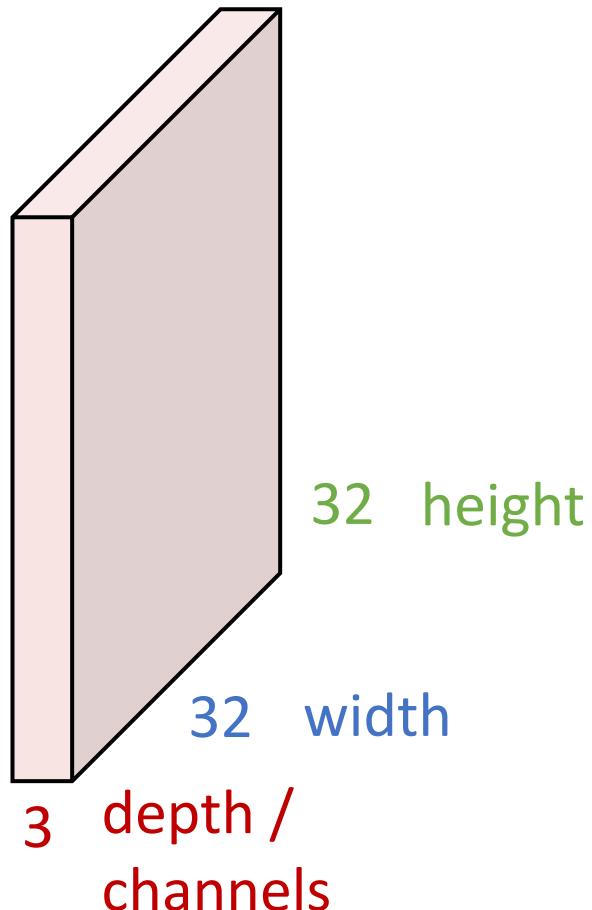
$3 \times 5 \times 5$  filter



**Convolve** the filter with the image  
i.e. “slide over the image spatially,  
computing dot products”

# Convolution Layer

$3 \times 32 \times 32$  image



$3 \times 5 \times 5$  filter

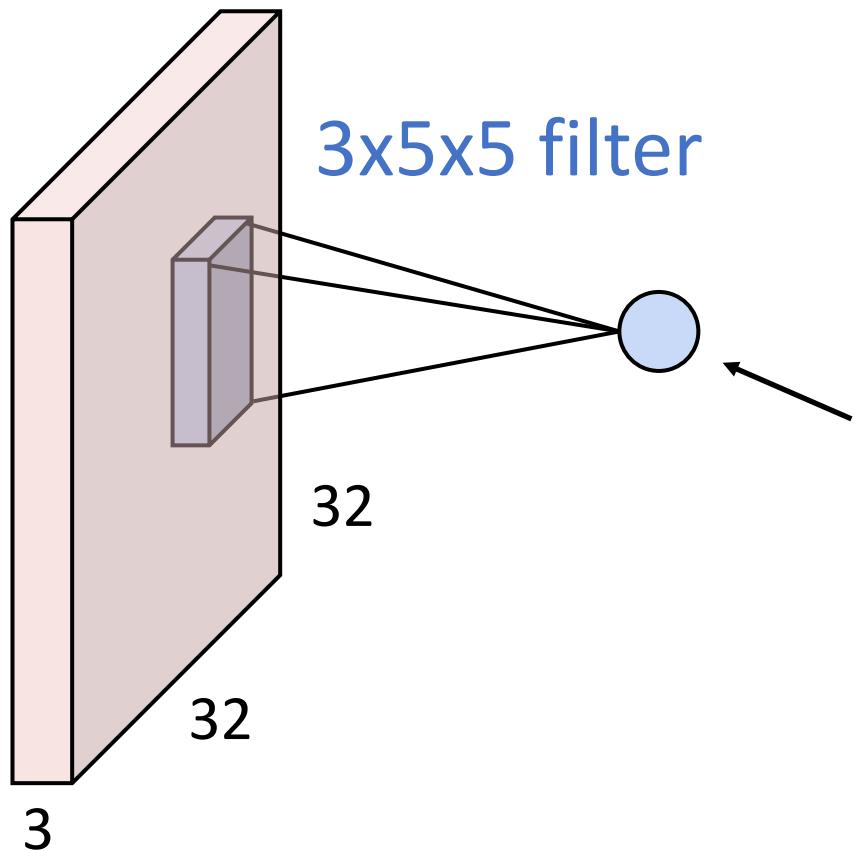


Filters always extend the full depth of the input volume

**Convolve** the filter with the image  
i.e. “slide over the image spatially,  
computing dot products”

# Convolution Layer

3x32x32 image

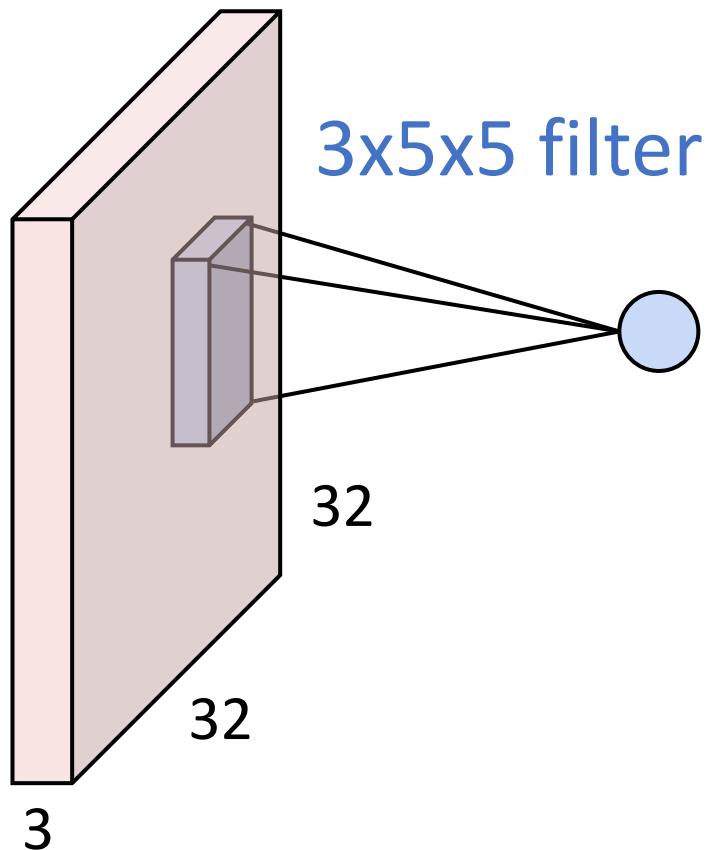


**1 number:**  
the result of taking a dot product between the filter  
and a small 3x5x5 chunk of the image  
(i.e.  $3 \times 5 \times 5 = 75$ -dimensional dot product + bias)

$$w^T x + b$$

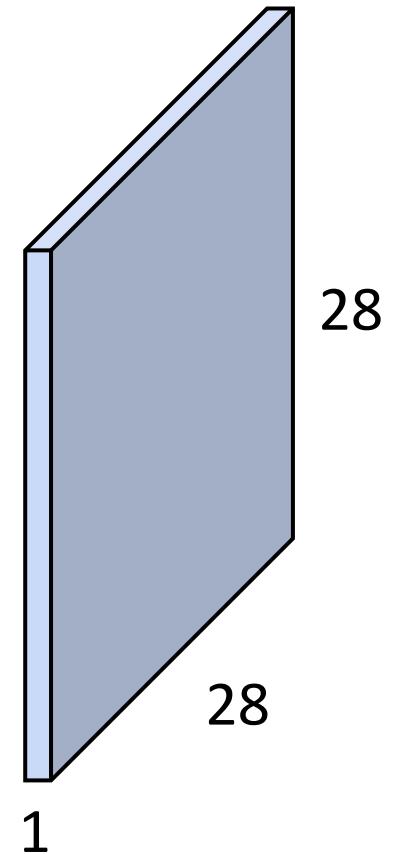
# Convolution Layer

3x32x32 image



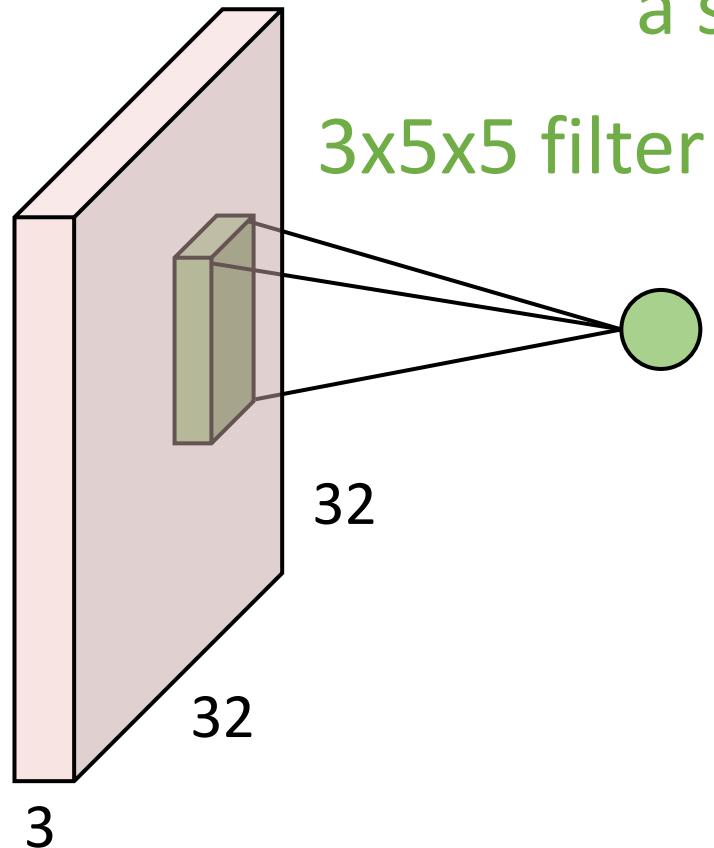
convolve (slide) over  
all spatial locations

1x28x28  
activation map



# Convolution Layer

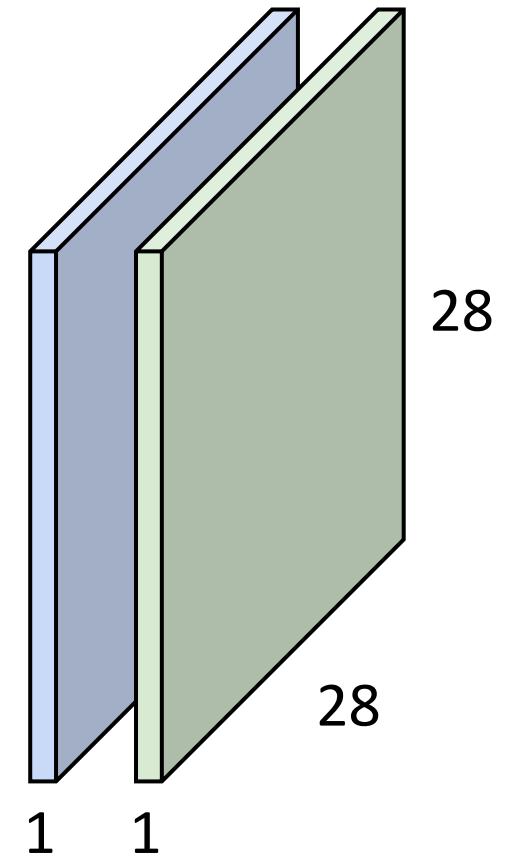
3x32x32 image



Consider repeating with  
a second (green) filter:

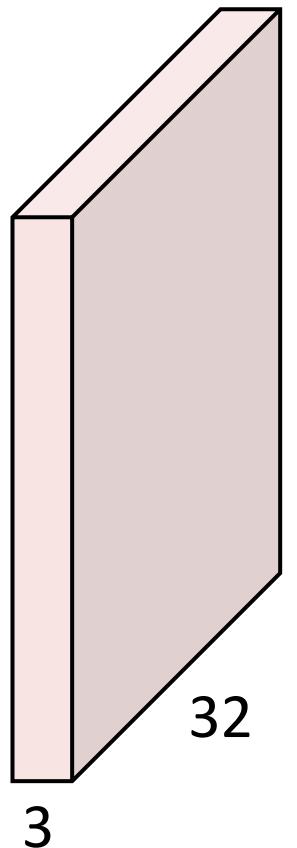
convolve (slide) over  
all spatial locations

two 1x28x28  
activation map

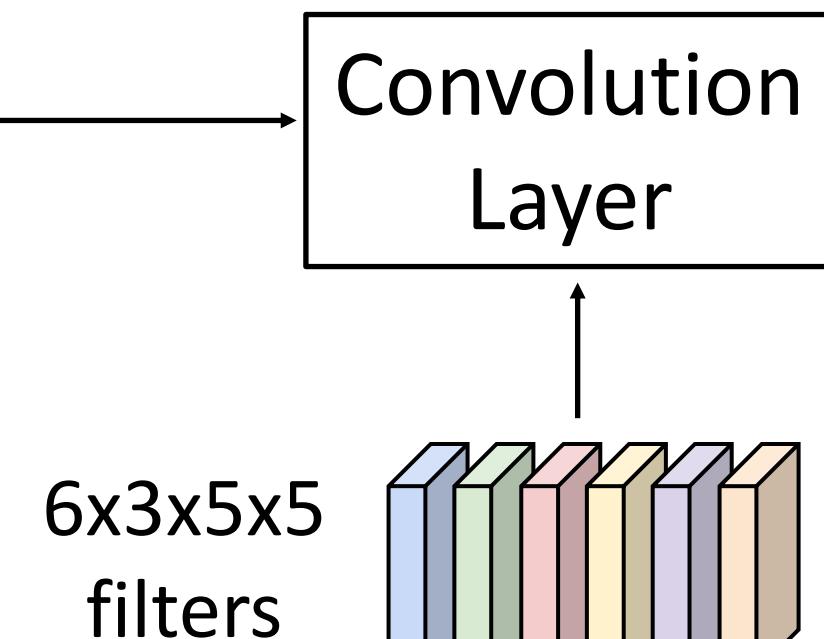


# Convolution Layer

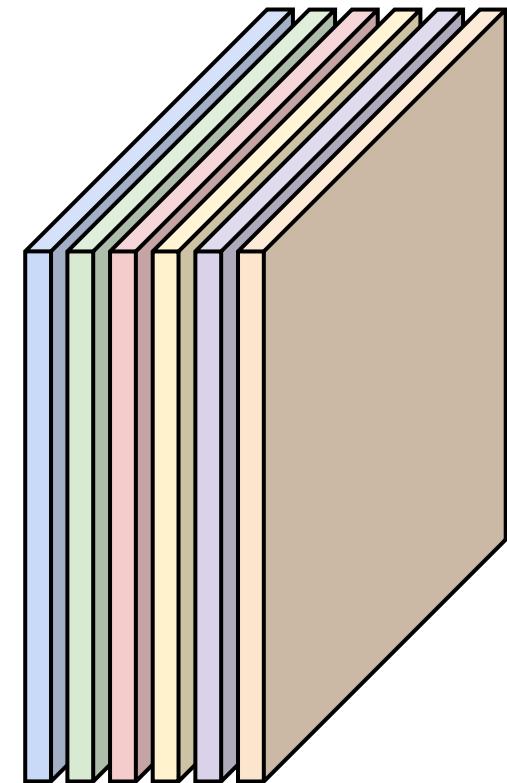
3x32x32 image



Consider 6 filters,  
each 3x5x5



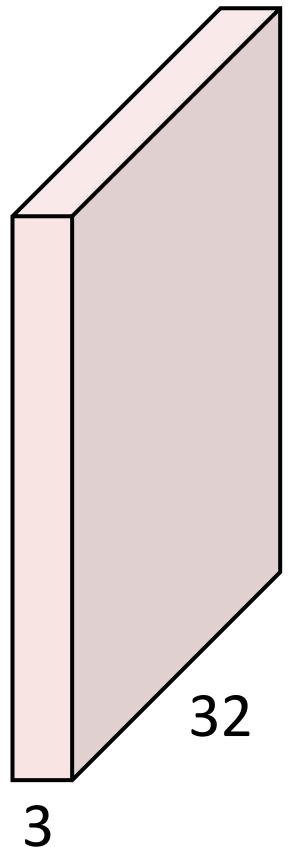
6 activation maps,  
each 1x28x28



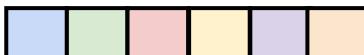
Stack activations to get a  
6x28x28 output image!

# Convolution Layer

3x32x32 image

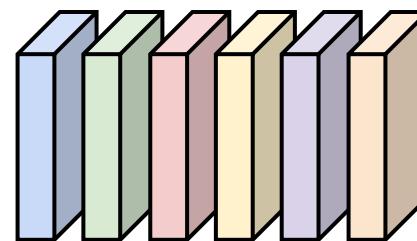


Also 6-dim bias vector:

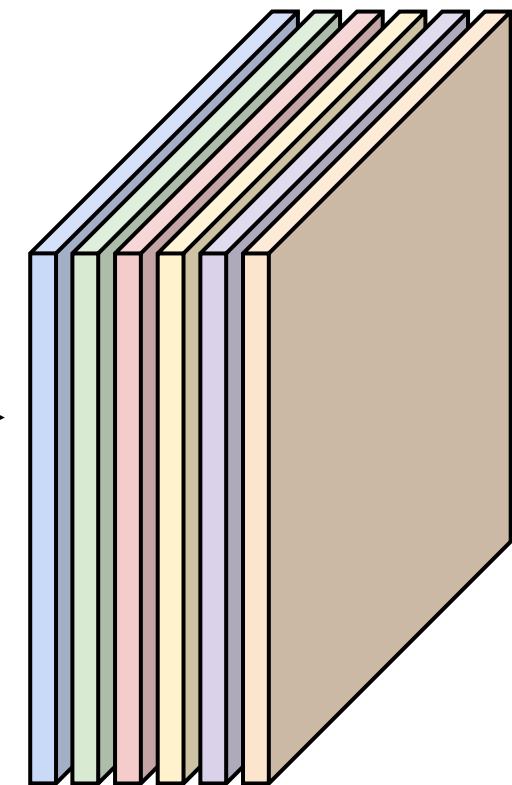


Convolution  
Layer

6x3x5x5  
filters



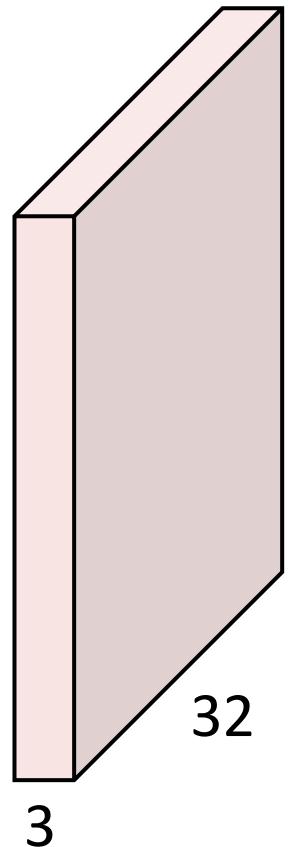
6 activation maps,  
each 1x28x28



Stack activations to get a  
6x28x28 output image!

# Convolution Layer

3x32x32 image

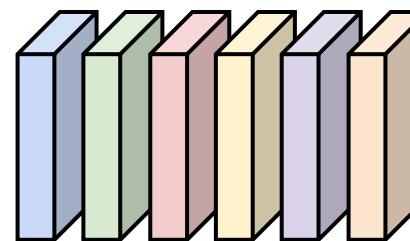


Also 6-dim bias vector:

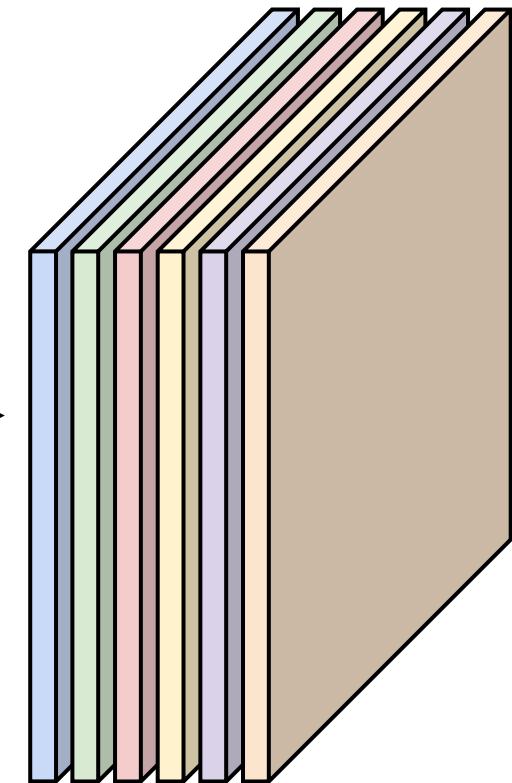


Convolution  
Layer

6x3x5x5  
filters



28x28 grid, at each  
point a 6-dim vector

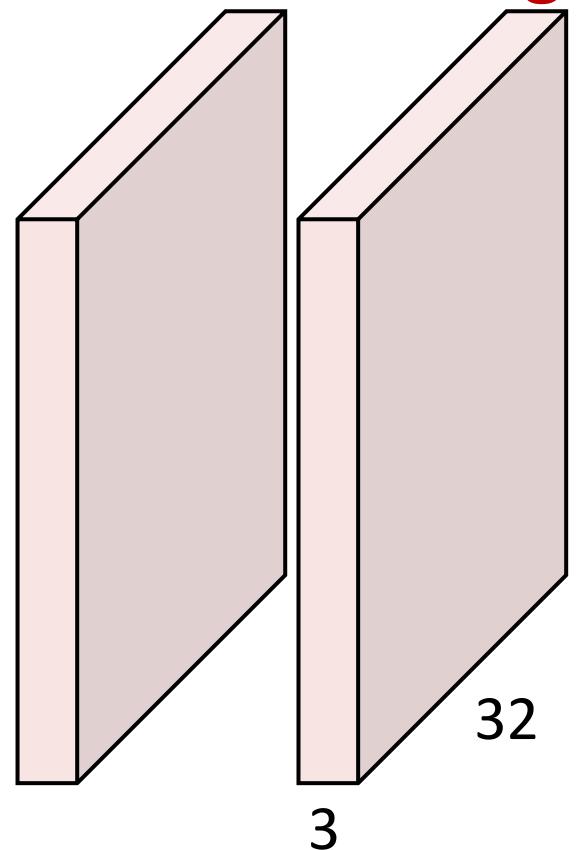


Stack activations to get a  
6x28x28 output image!

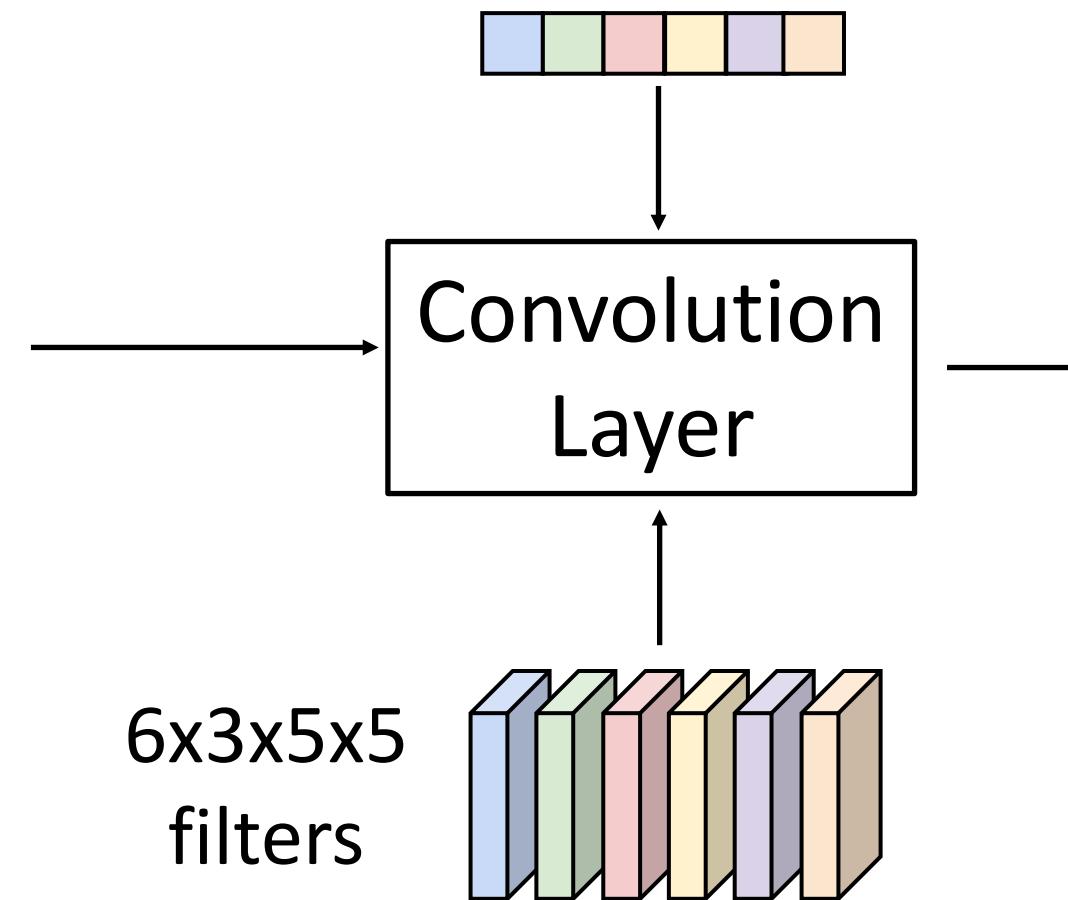
# Convolution Layer

$2 \times 3 \times 32 \times 32$

Batch of images

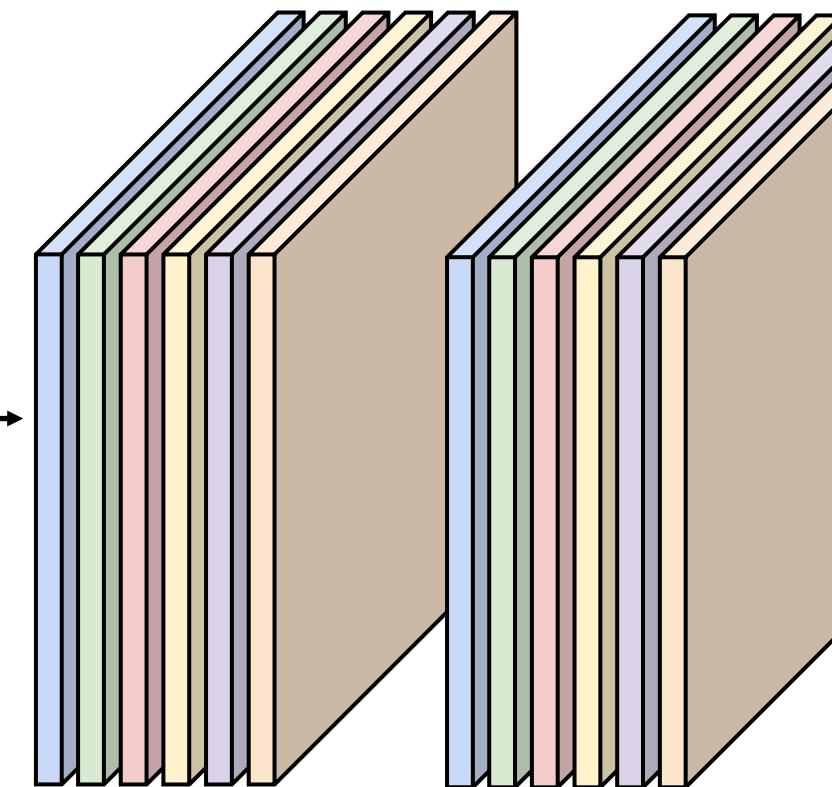


Also 6-dim bias vector:



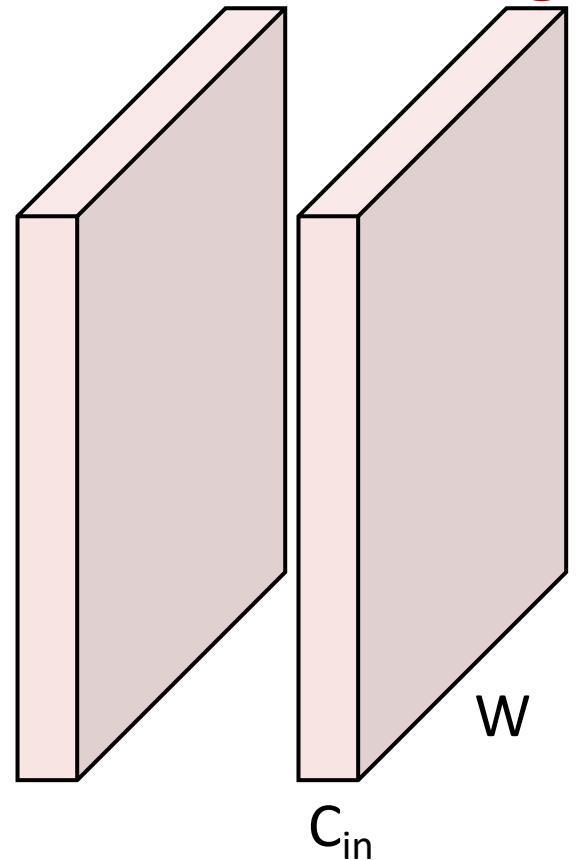
$2 \times 6 \times 28 \times 28$

Batch of outputs



# Convolution Layer

$N \times C_{in} \times H \times W$   
Batch of images

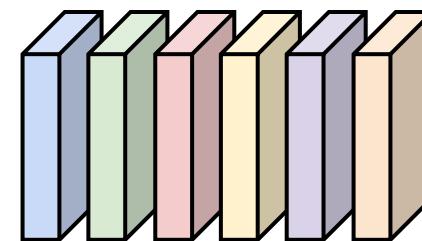


Also  $C_{out}$ -dim bias vector:

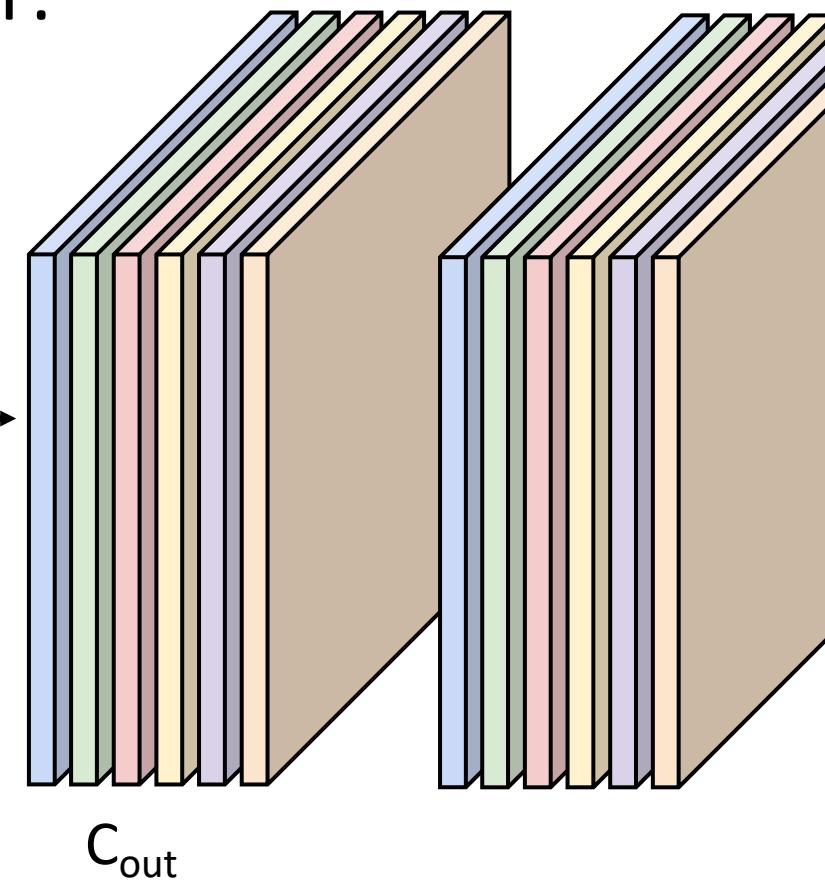


Convolution  
Layer

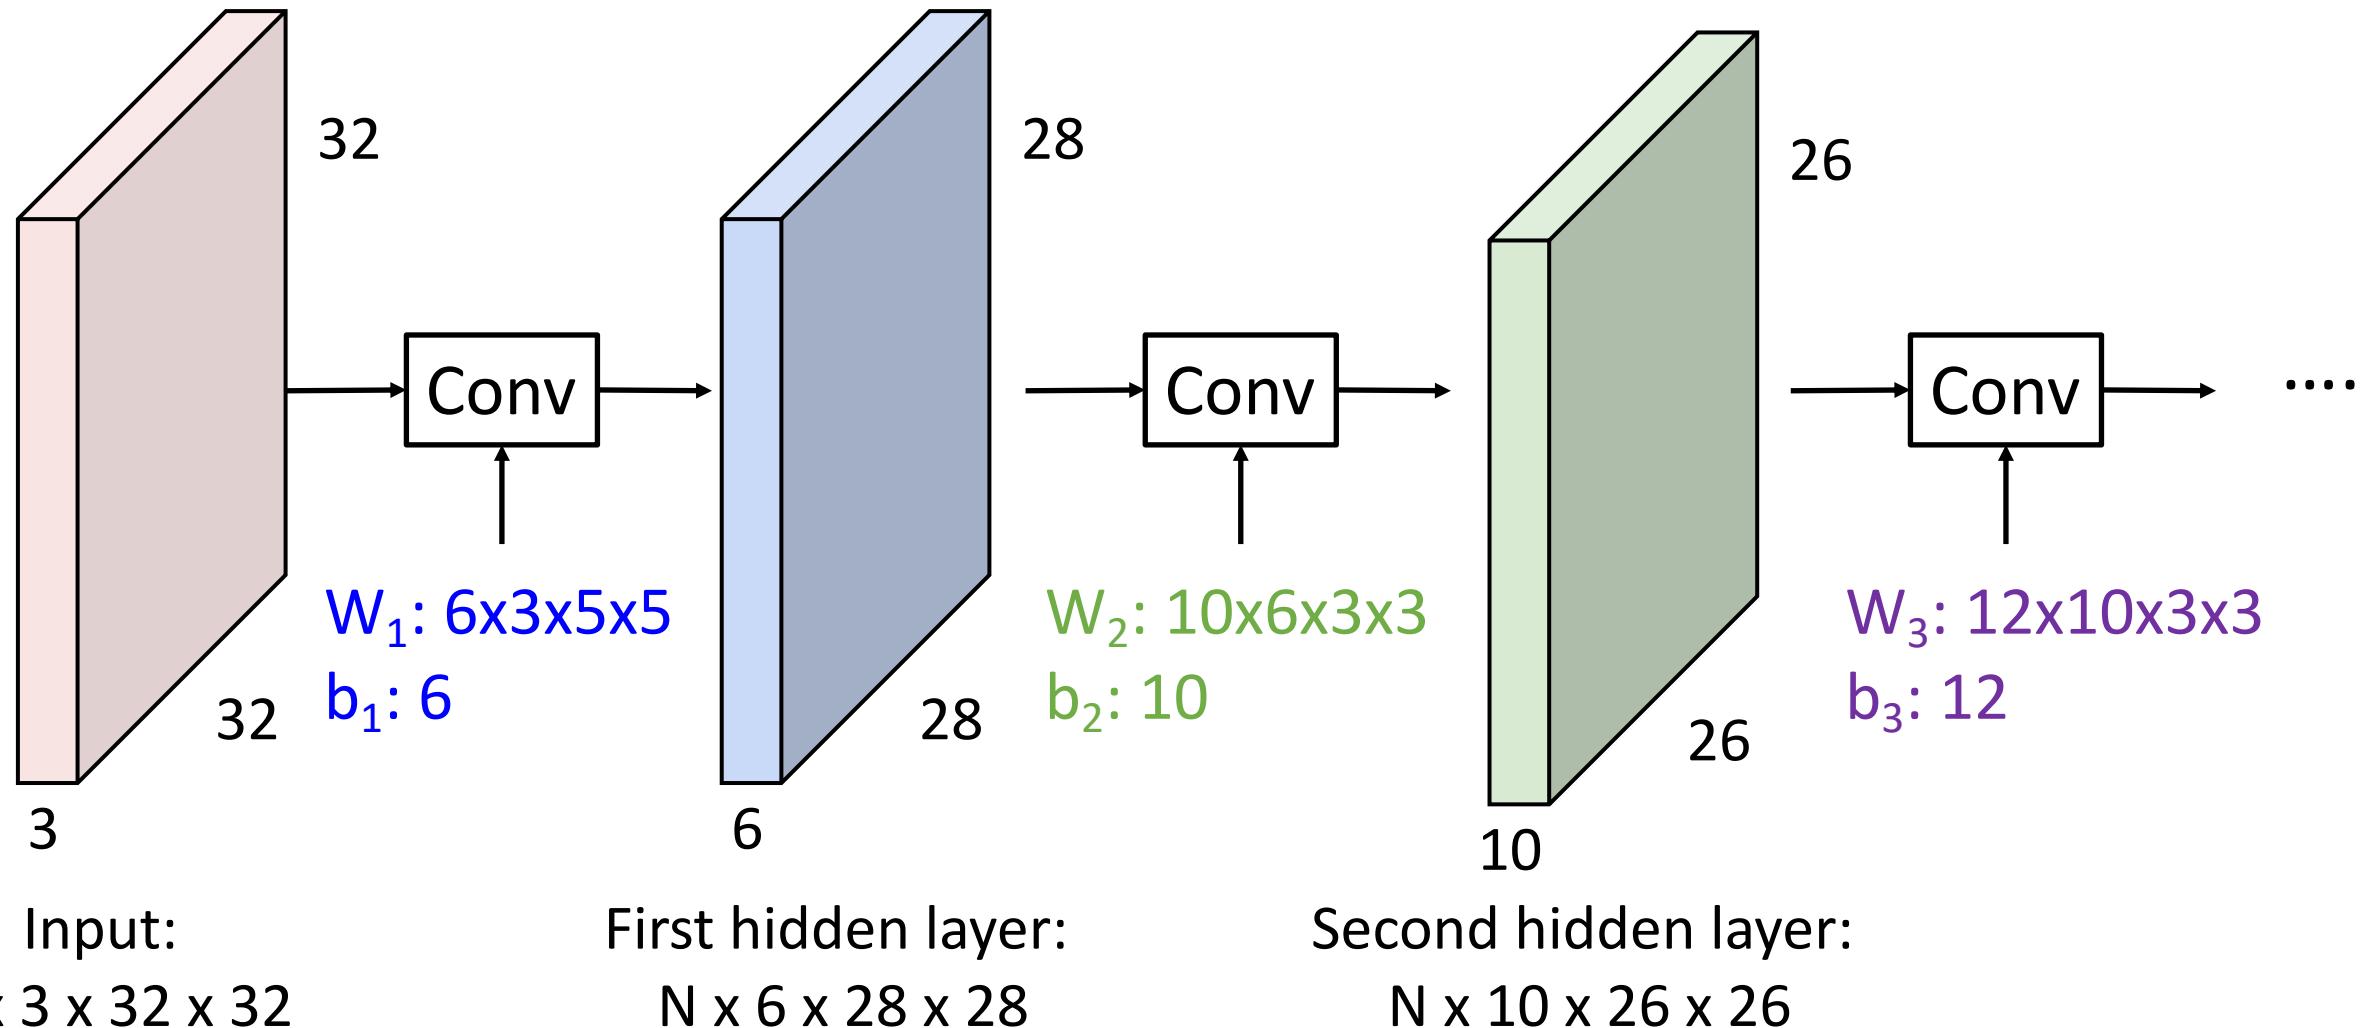
$C_{out} \times C_{in} \times K_w \times K_h$   
filters



$N \times C_{out} \times H' \times W'$   
Batch of outputs

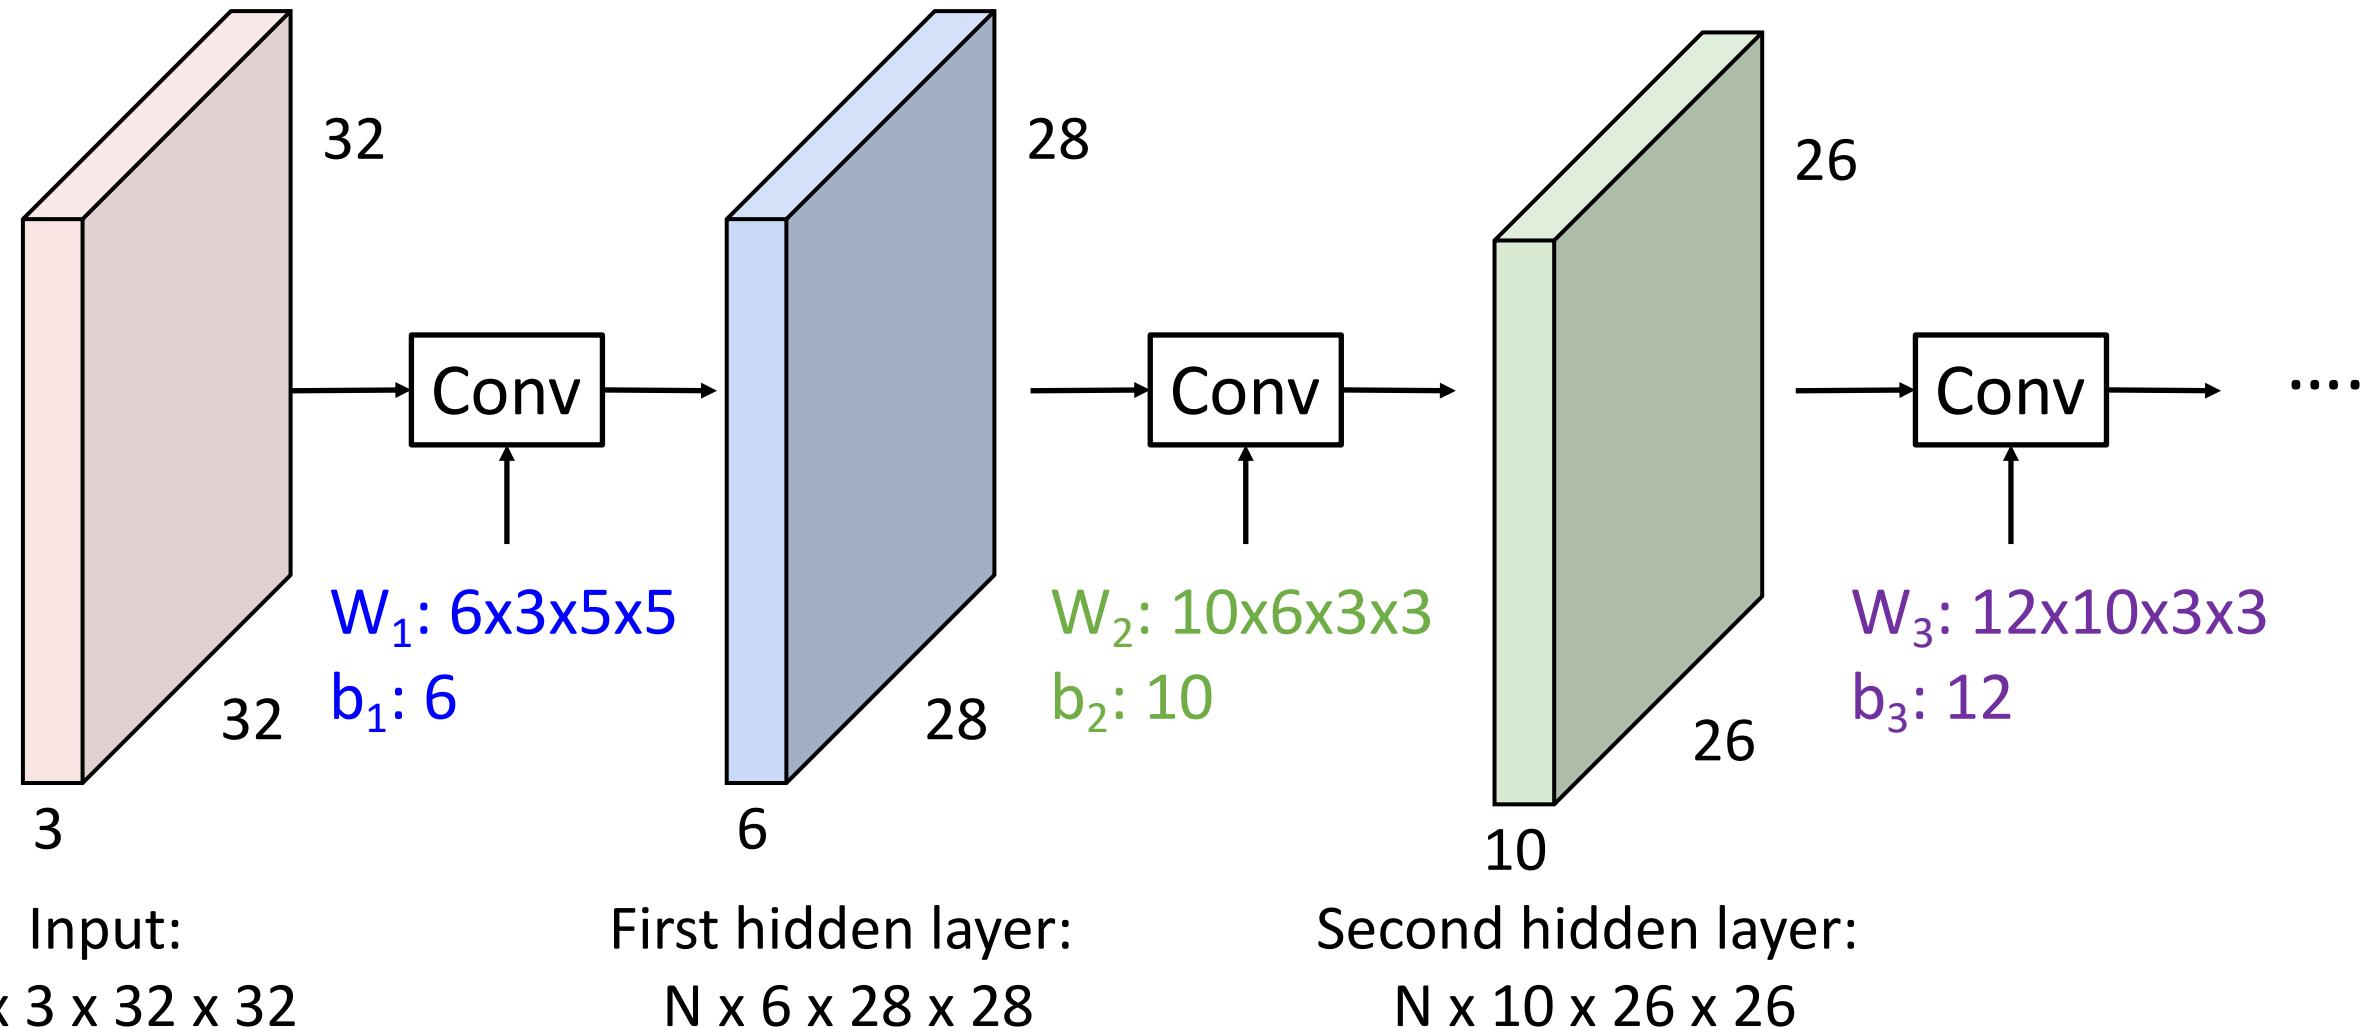


# Stacking Convolutions

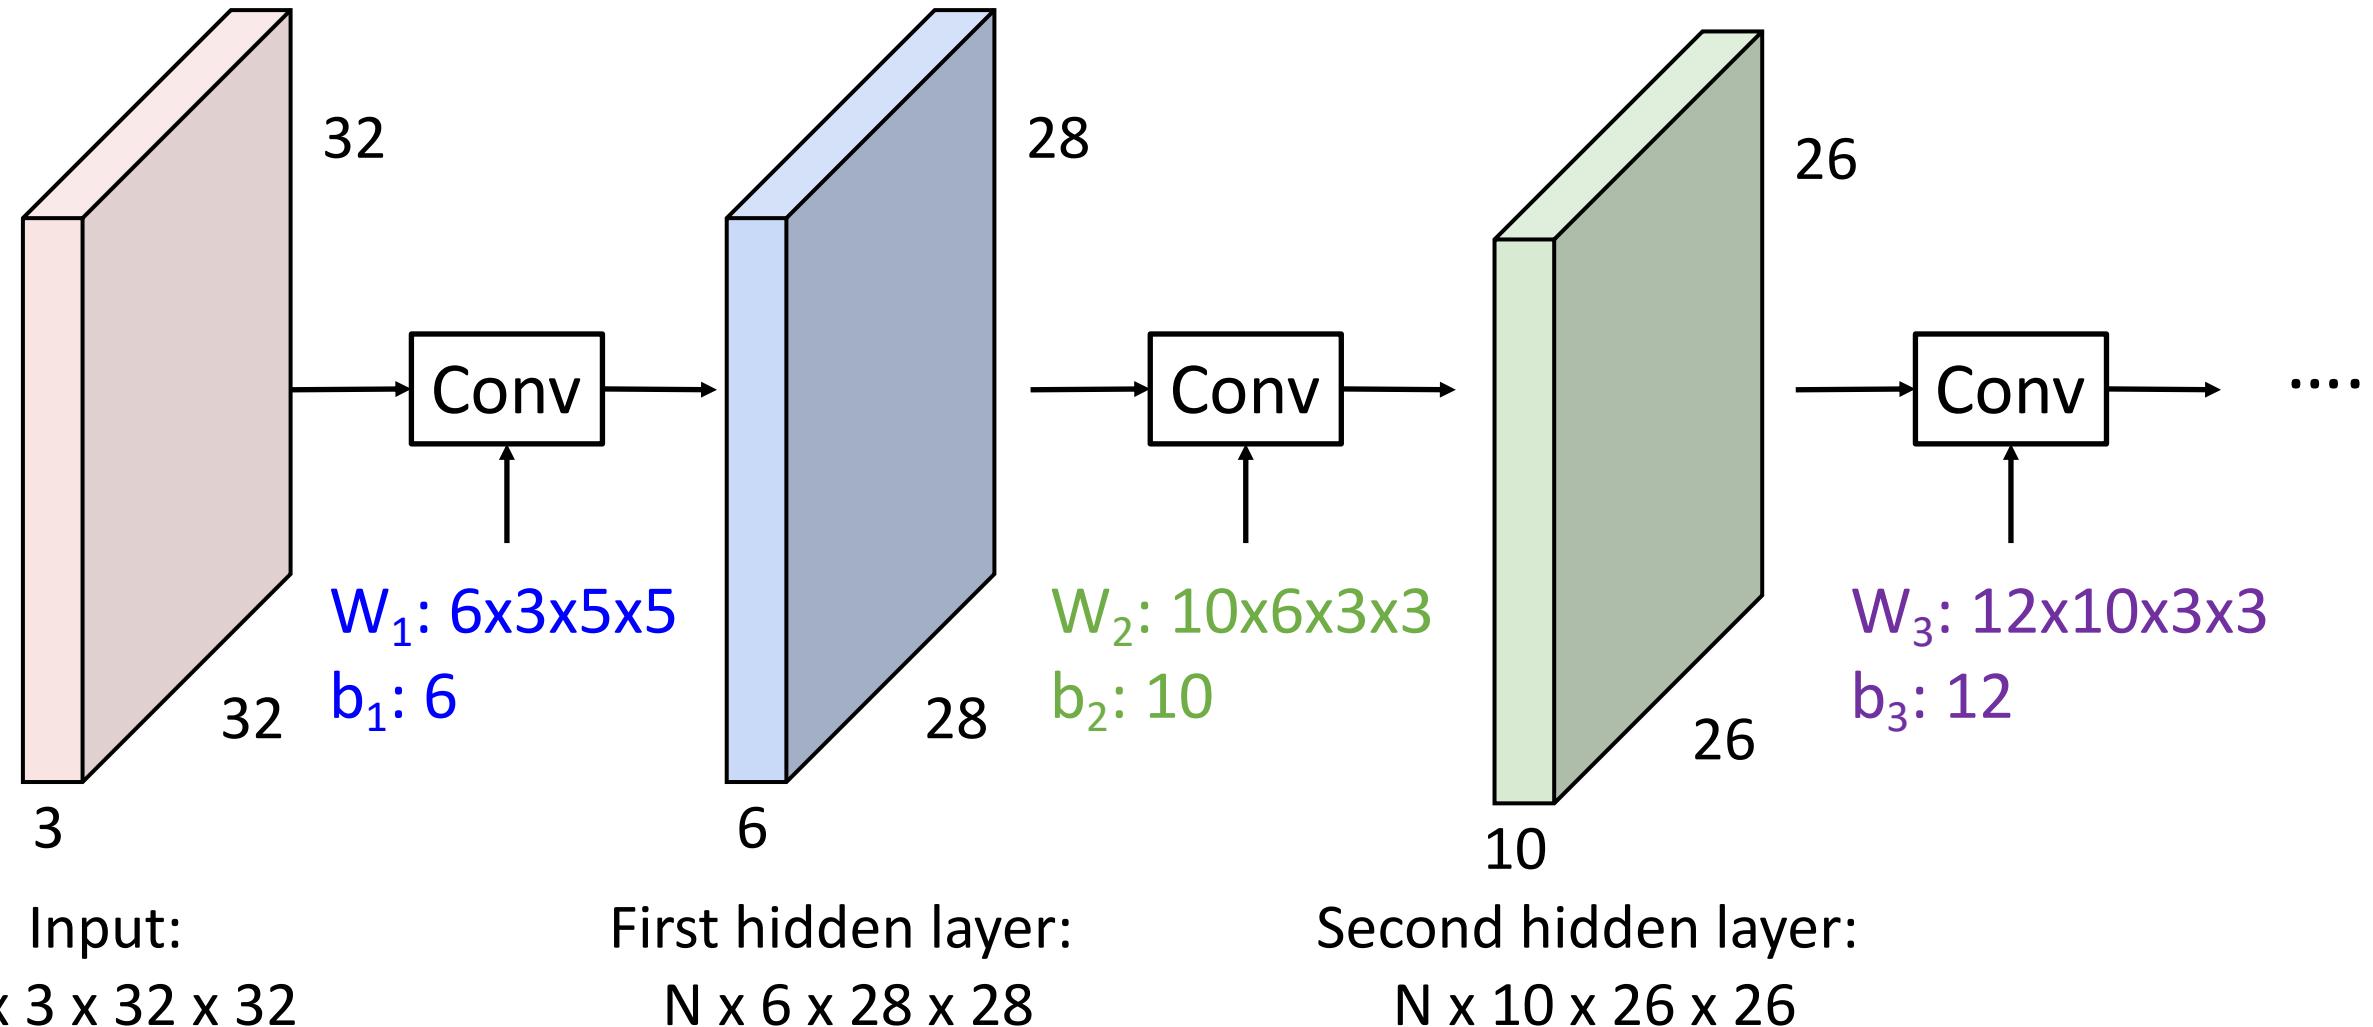


# Stacking Convolutions

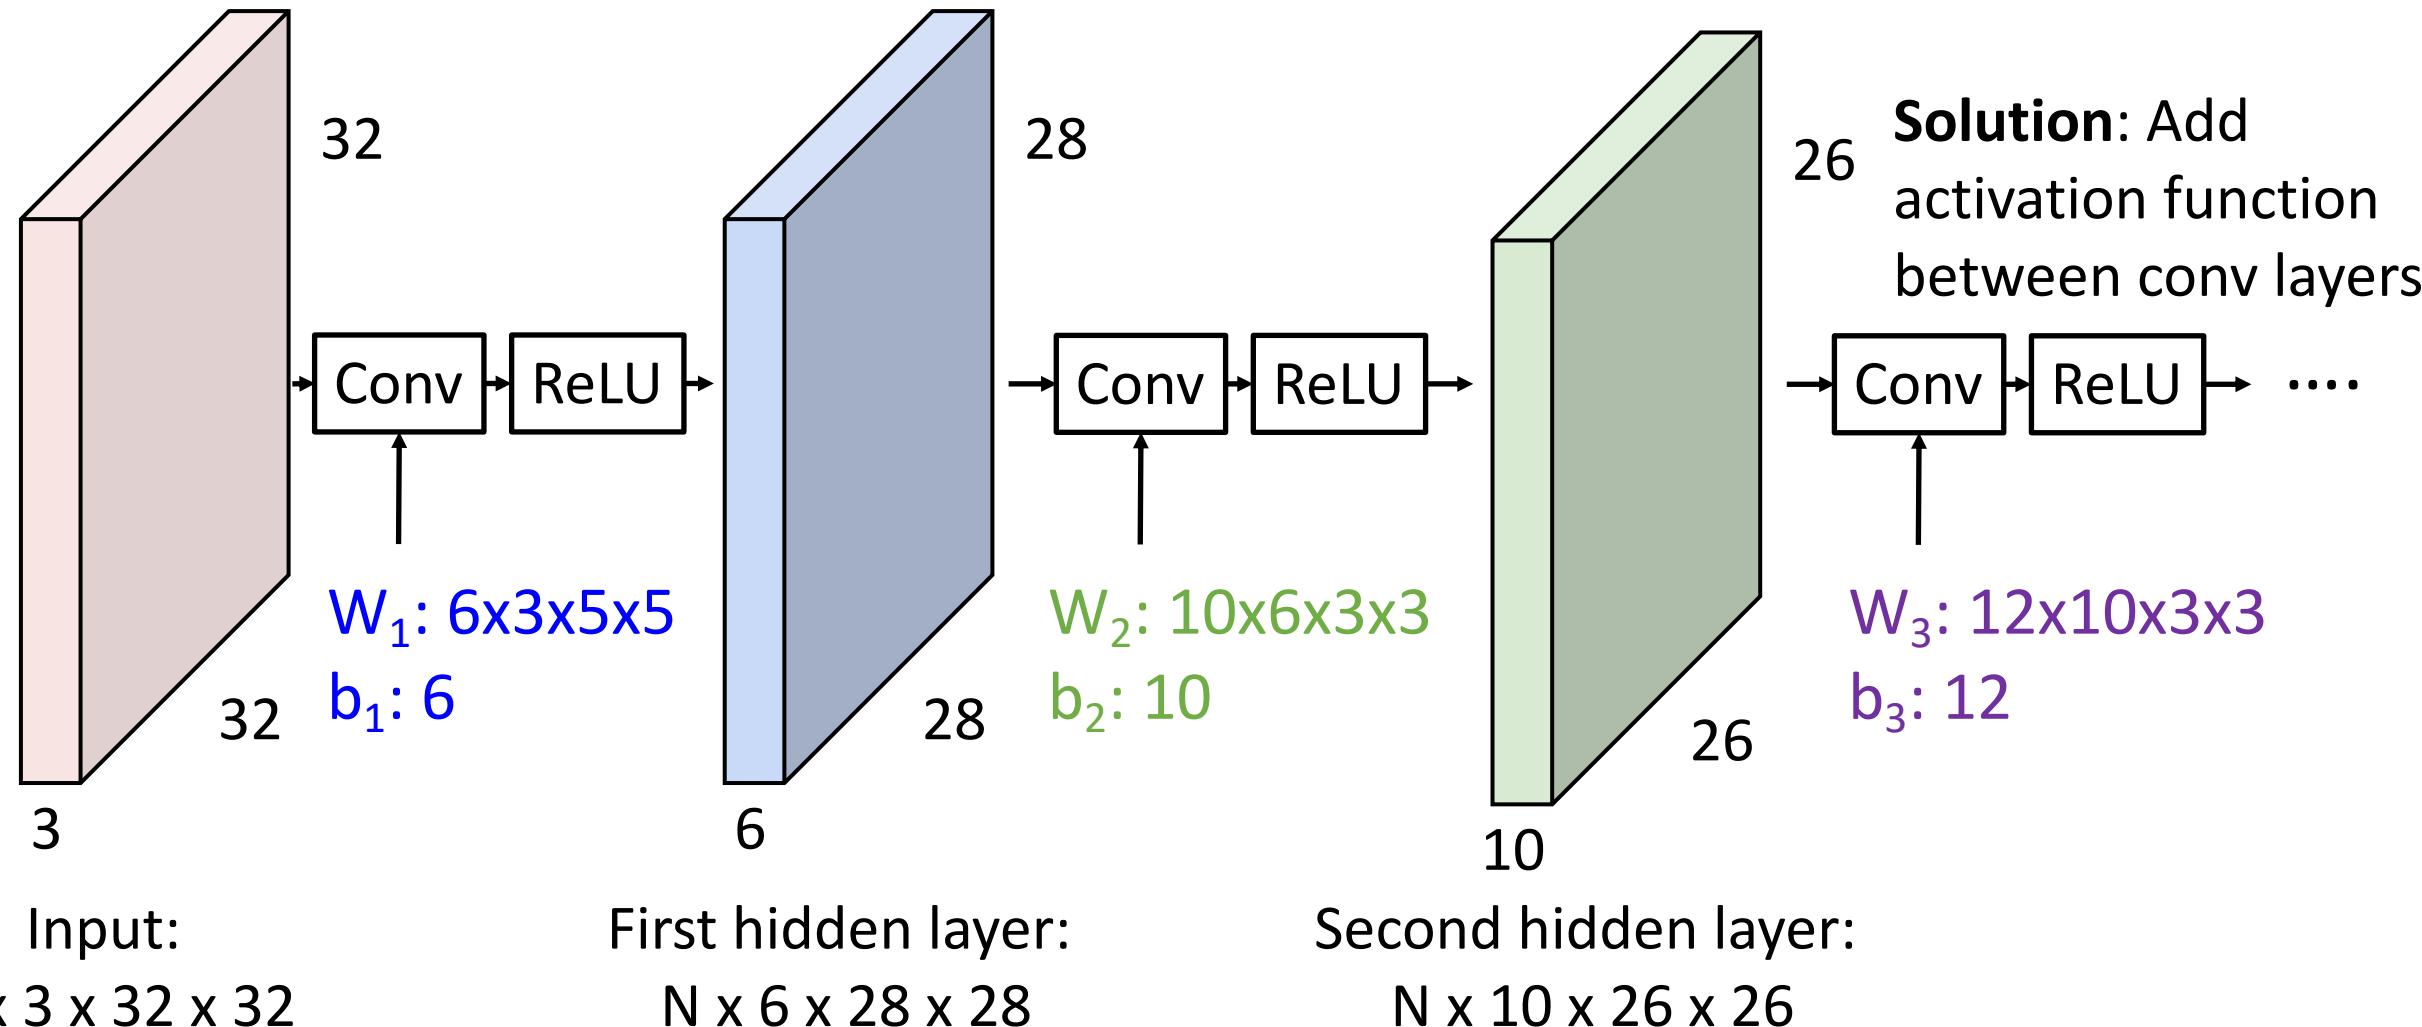
Q: What happens if we stack two convolution layers?



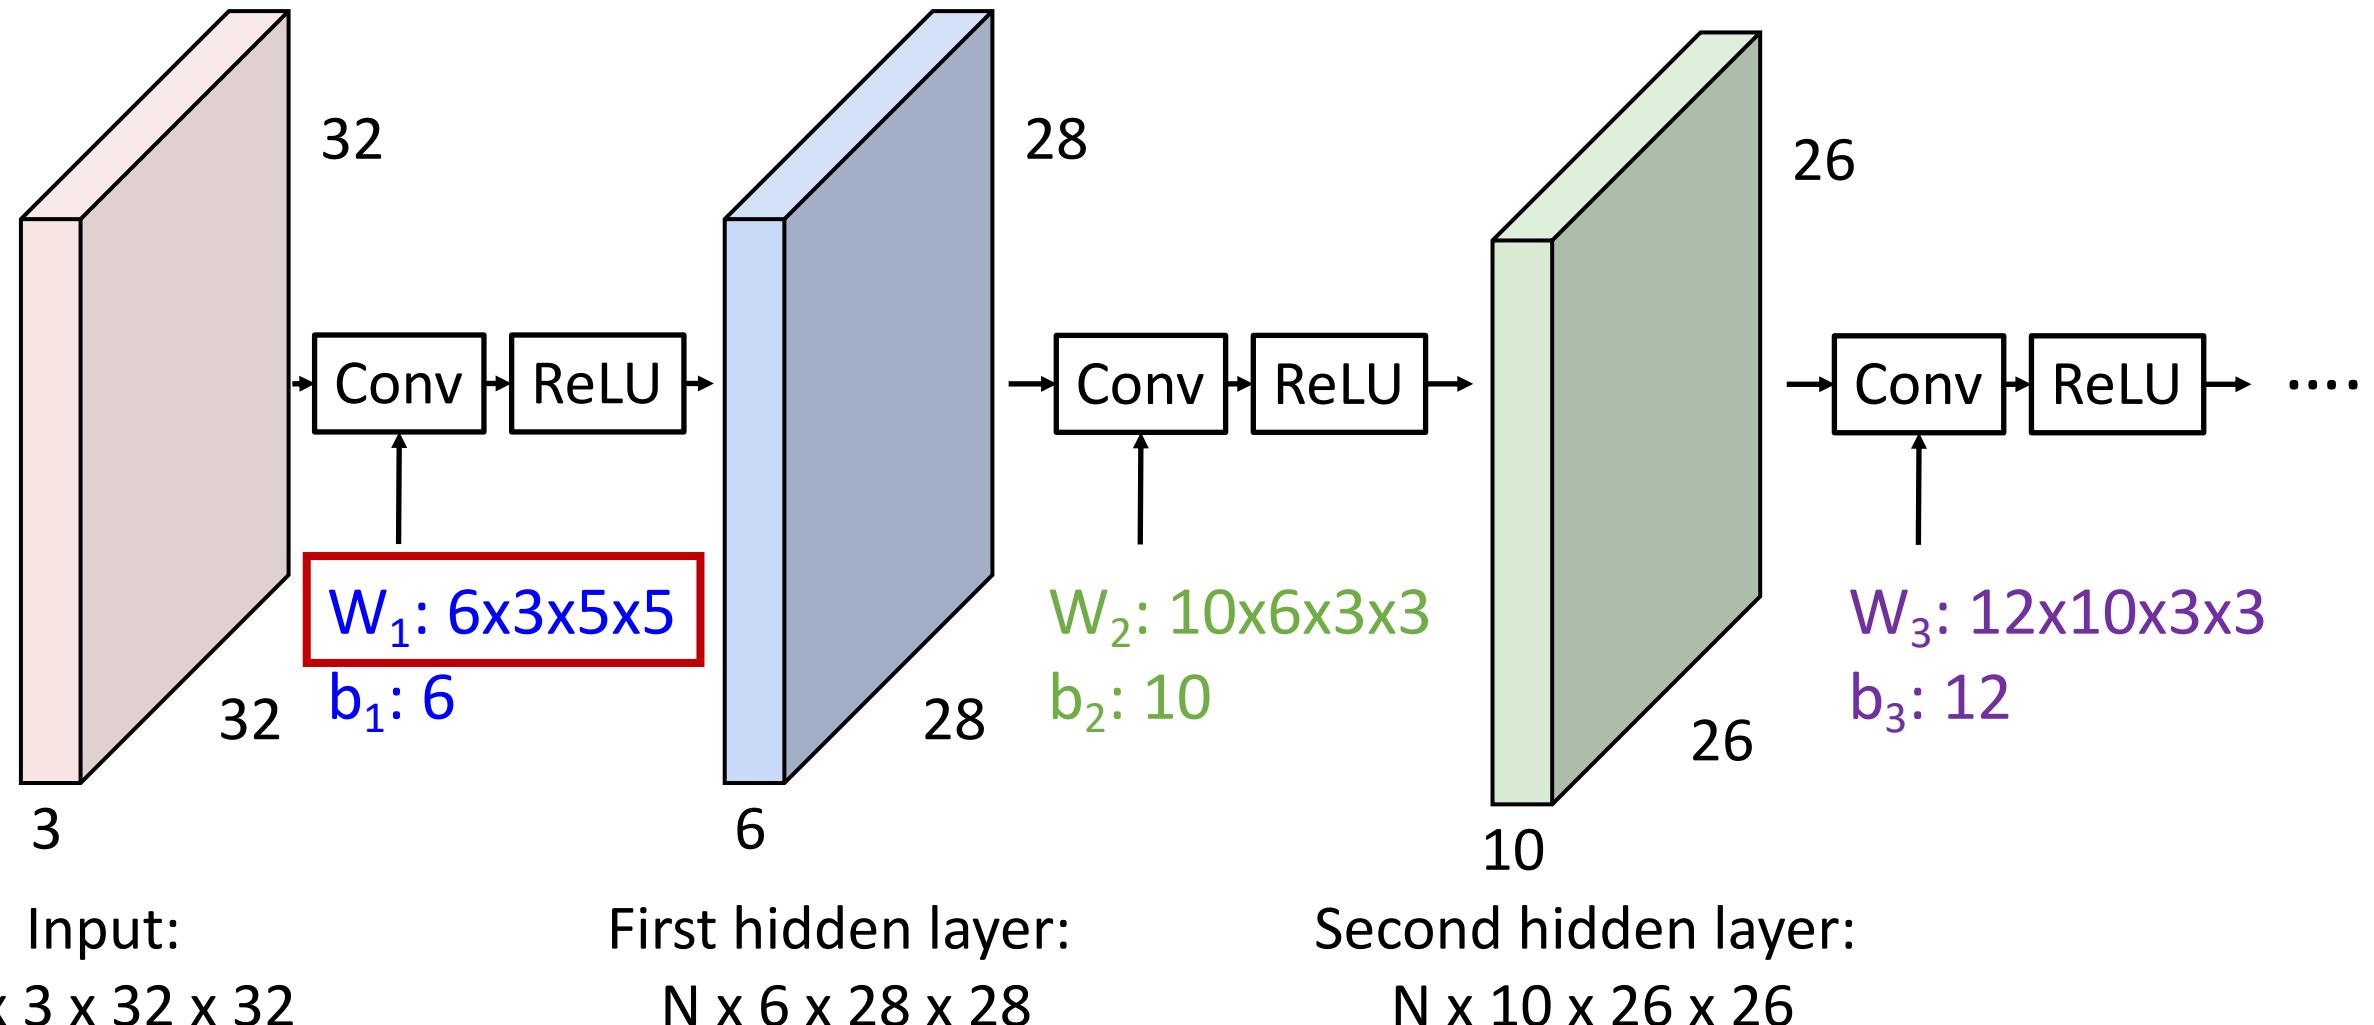
# Stacking Convolutions



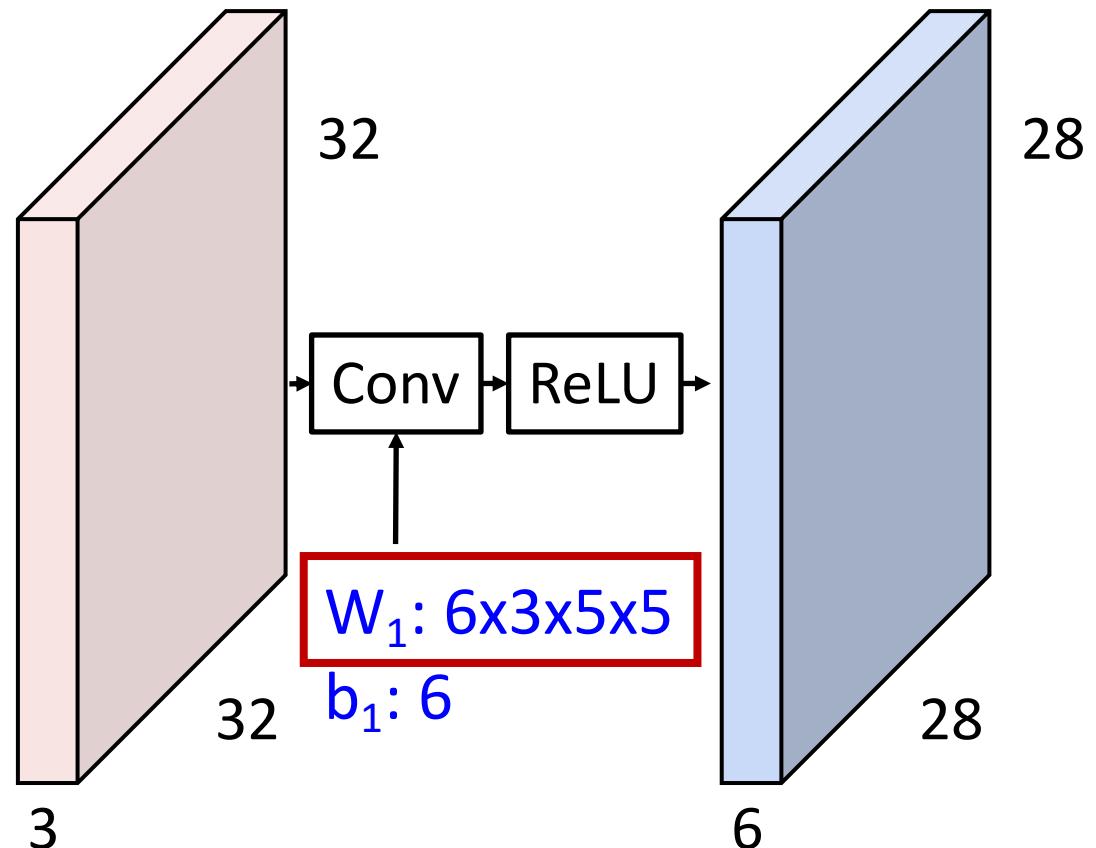
# Stacking Convolutions



# What do convolutional filters learn?



# What do convolutional filters learn?



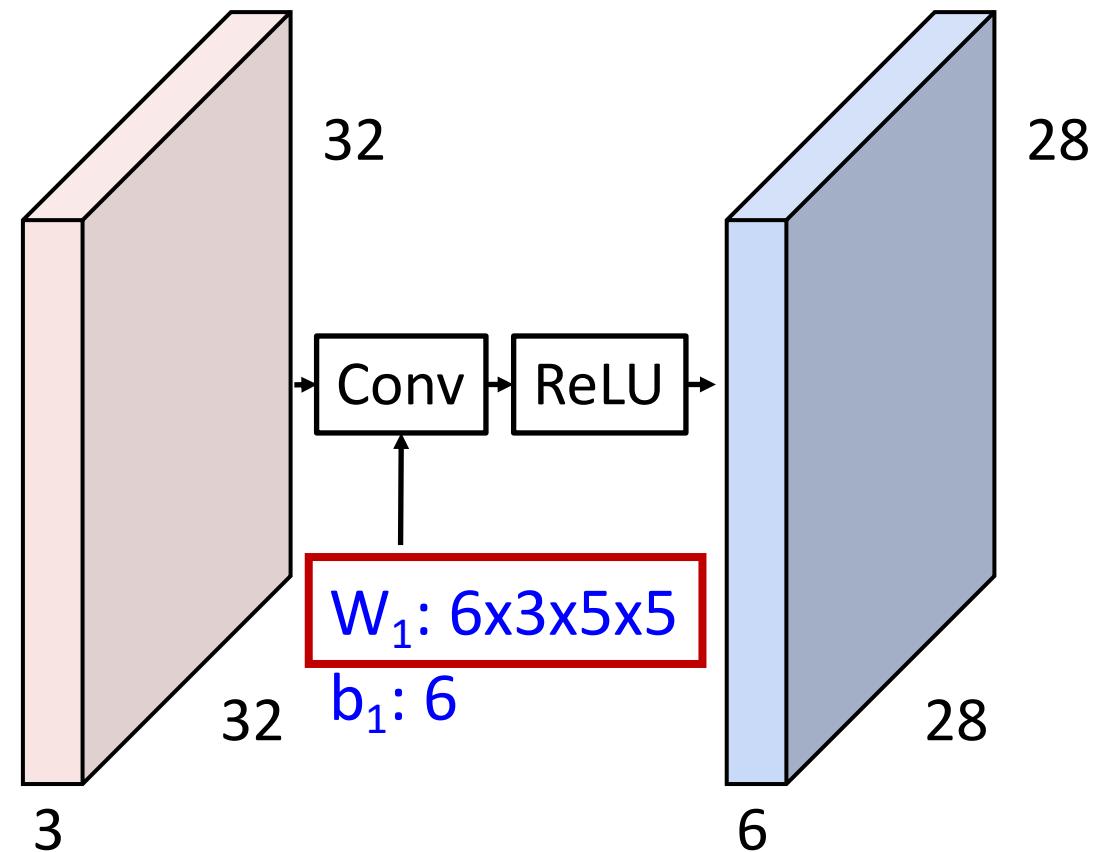
Input:  
 $N \times 3 \times 32 \times 32$

First hidden layer:  
 $N \times 6 \times 28 \times 28$

Linear classifier: One template per class



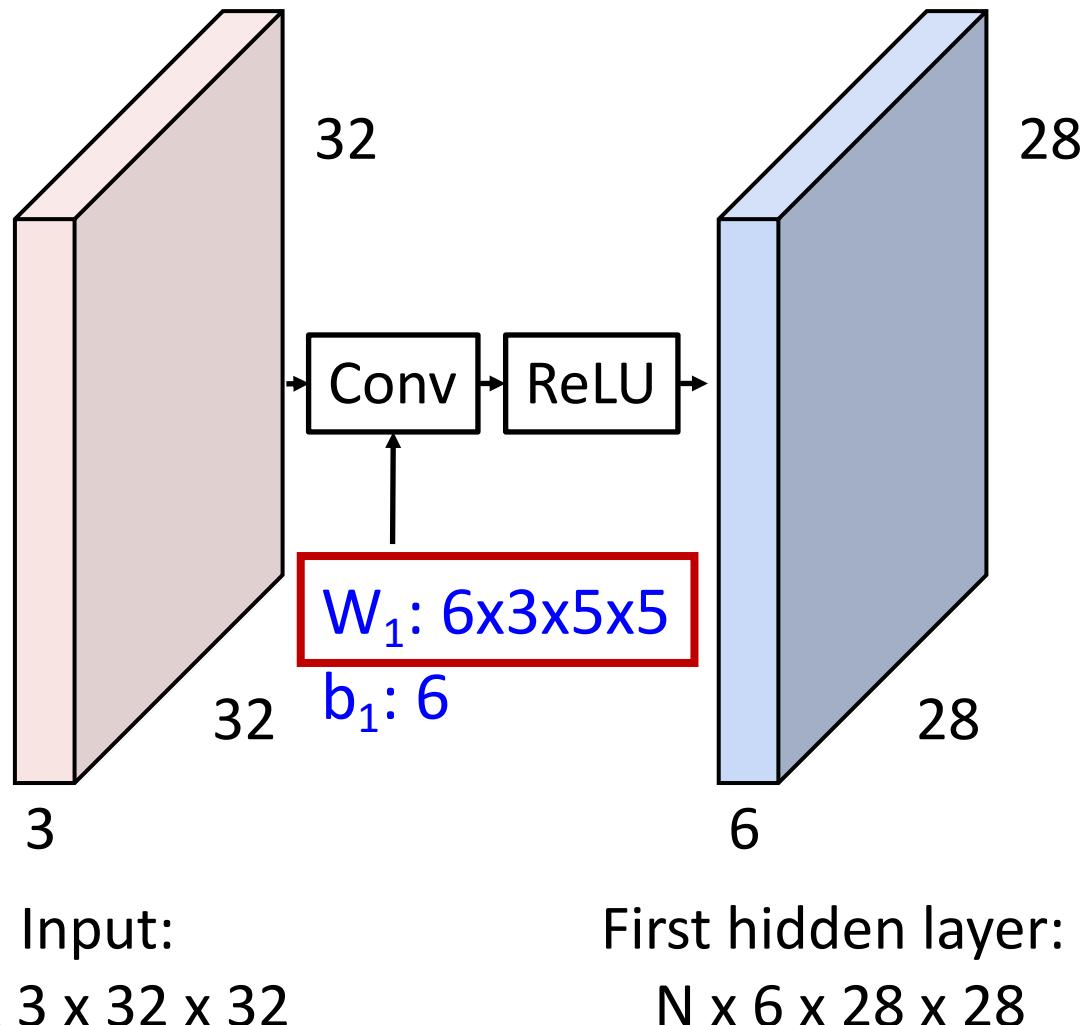
# What do convolutional filters learn?



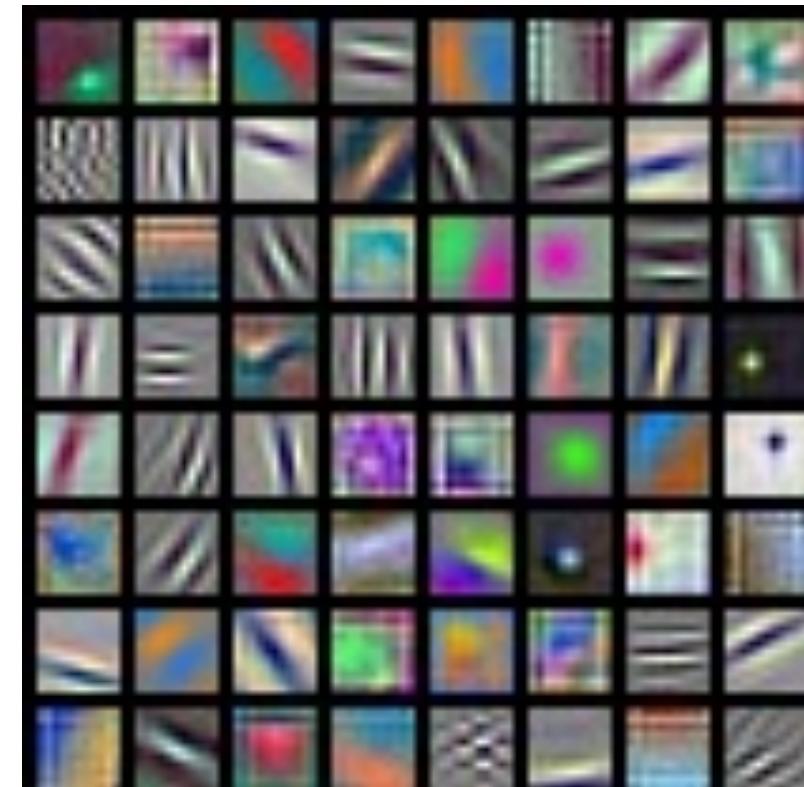
MLP: Bank of whole-image templates



# What do convolutional filters learn?

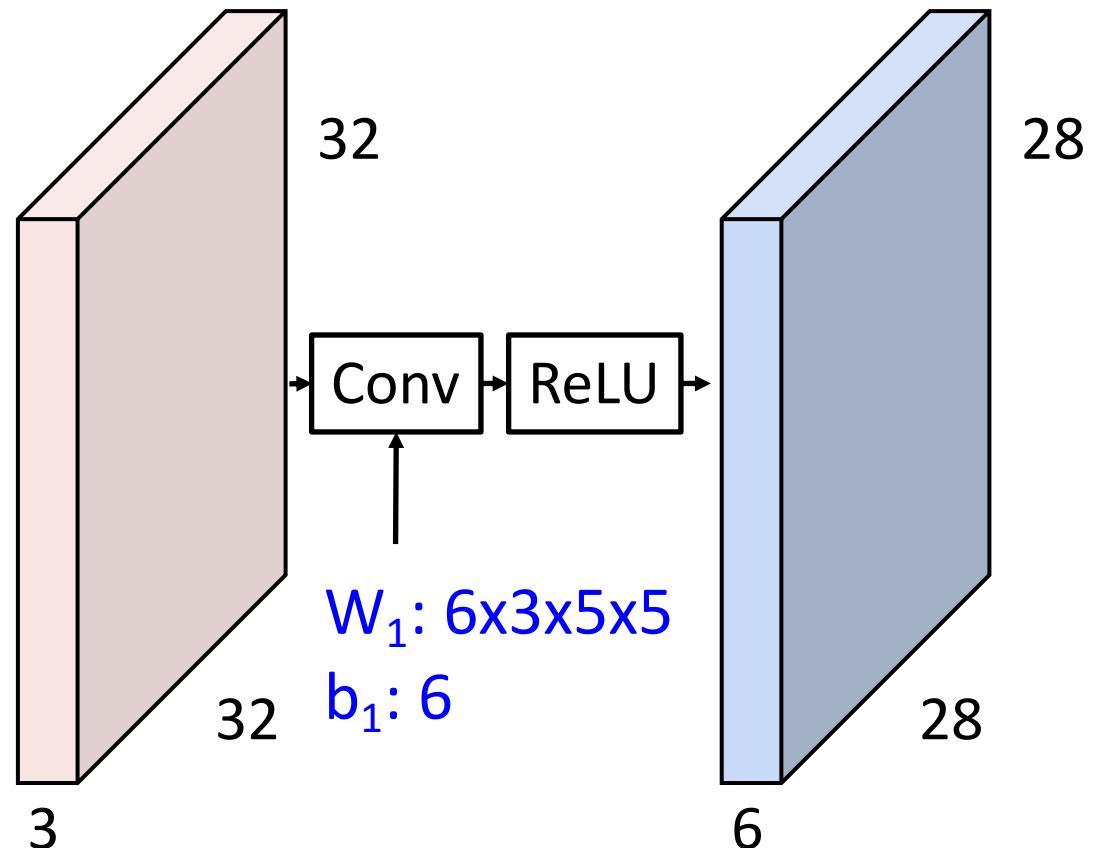


First-layer conv filters: local image templates  
(Often learns oriented edges, opposing colors)



AlexNet: 64 filters, each  $3 \times 11 \times 11$

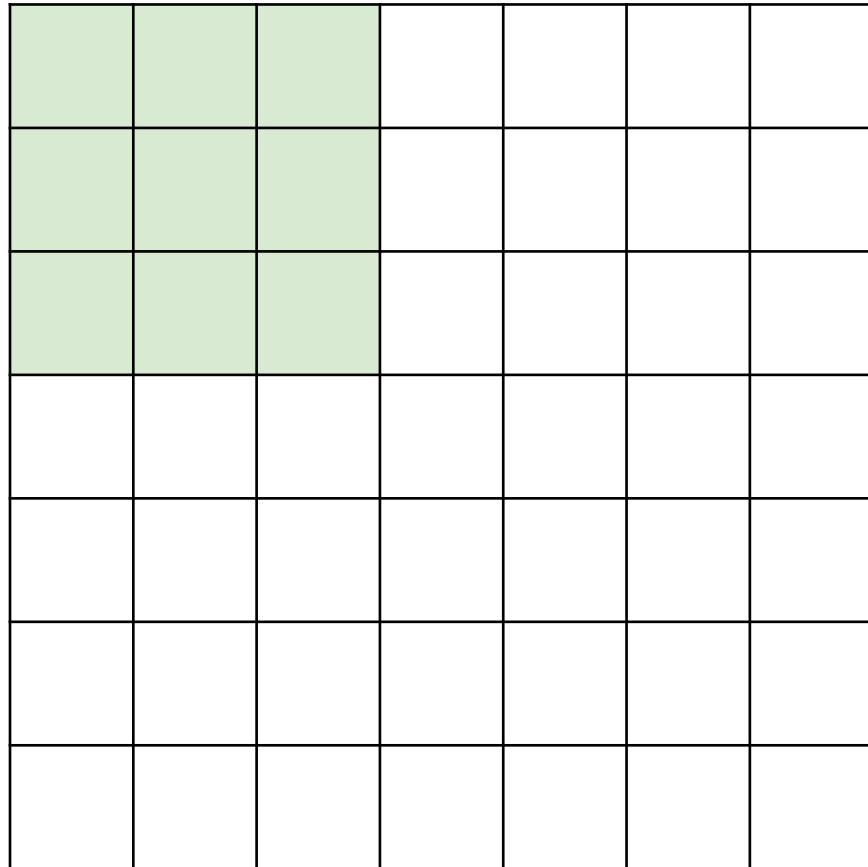
# A closer look at spatial dimensions



Input:  
 $N \times 3 \times 32 \times 32$

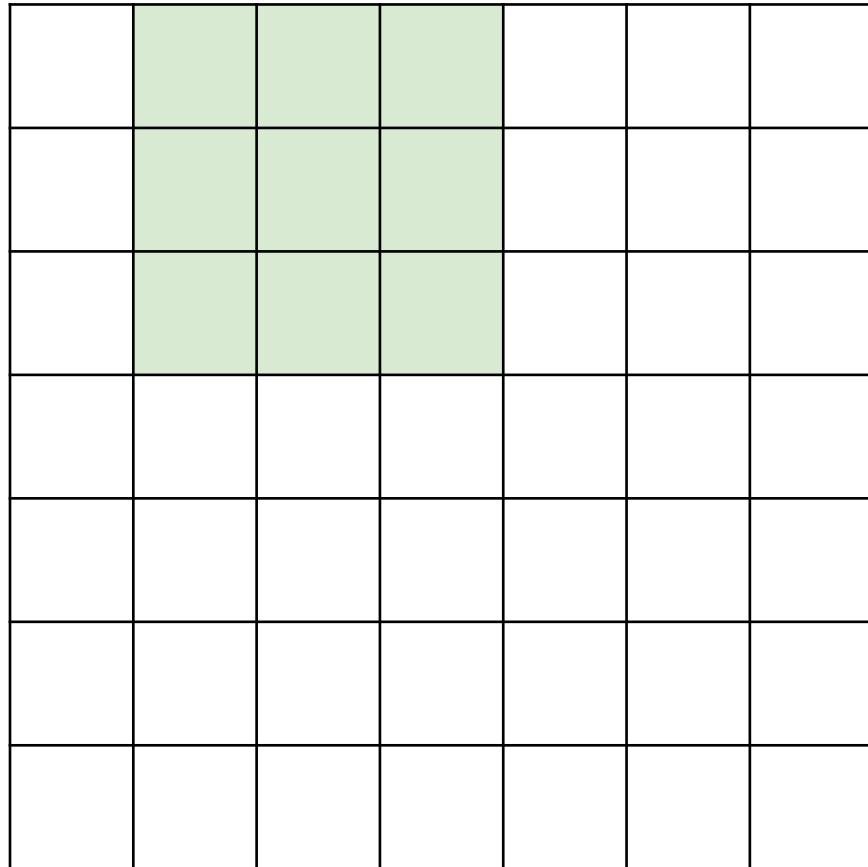
First hidden layer:  
 $N \times 6 \times 28 \times 28$

# A closer look at spatial dimensions



Input: 7x7  
Filter: 3x3

# A closer look at spatial dimensions

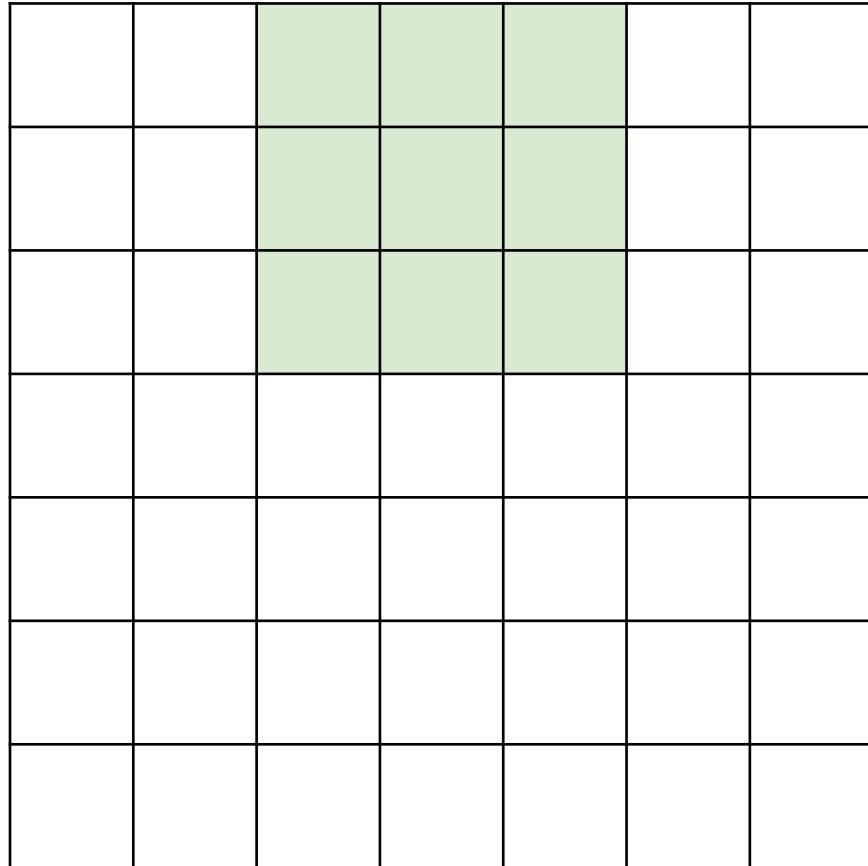


7

Input: 7x7  
Filter: 3x3

7

# A closer look at spatial dimensions

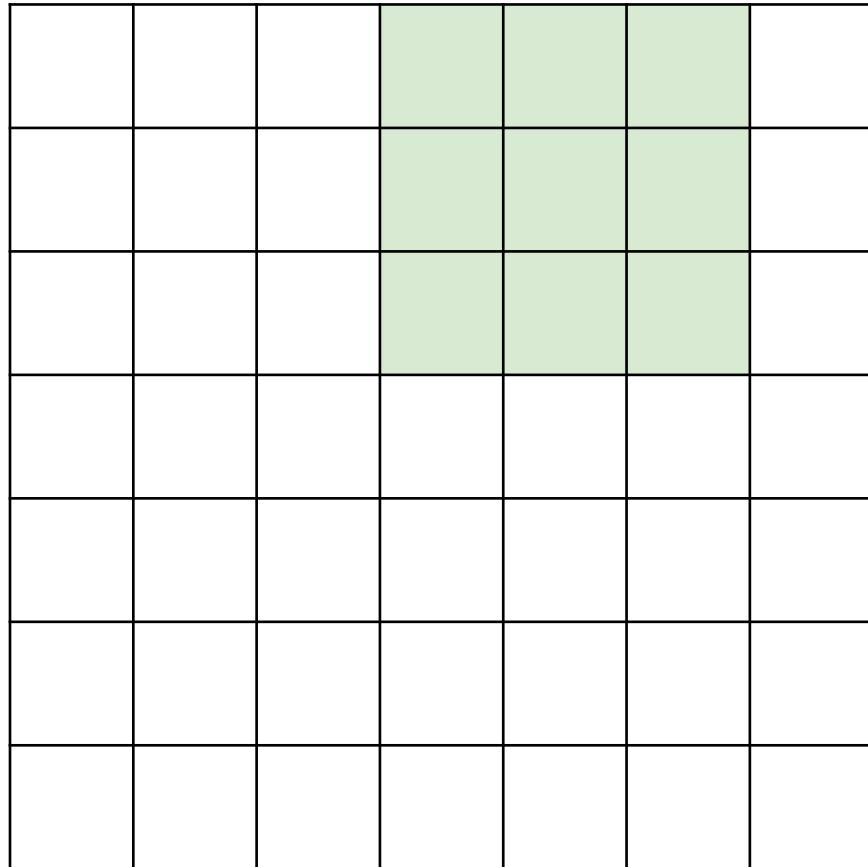


7

Input: 7x7  
Filter: 3x3

7

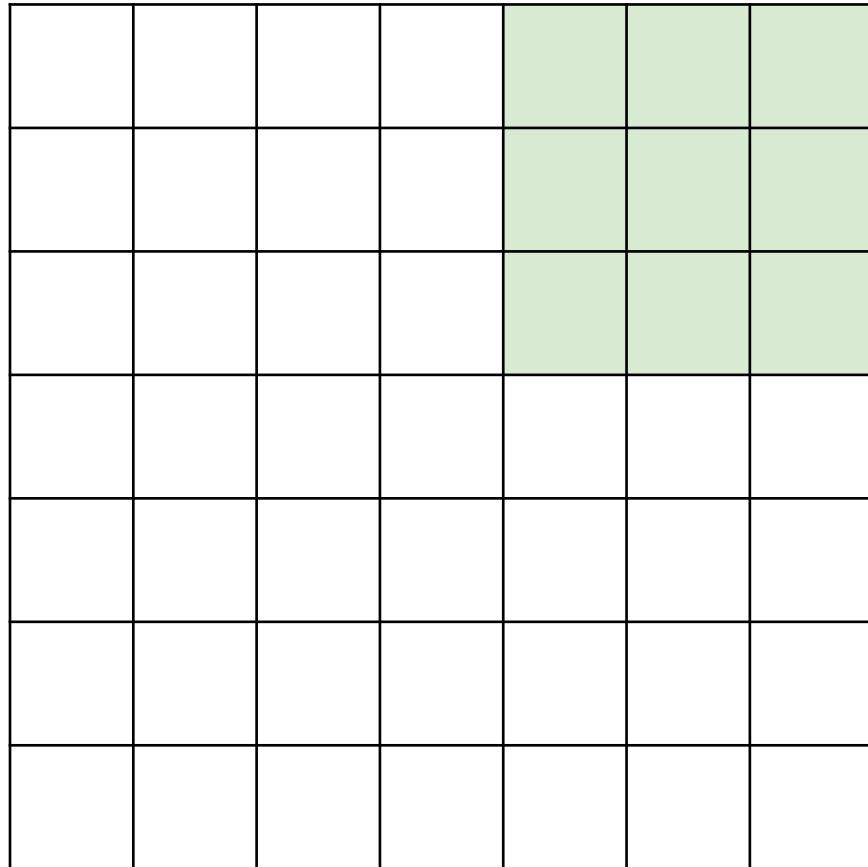
# A closer look at spatial dimensions



7

Input: 7x7  
Filter: 3x3

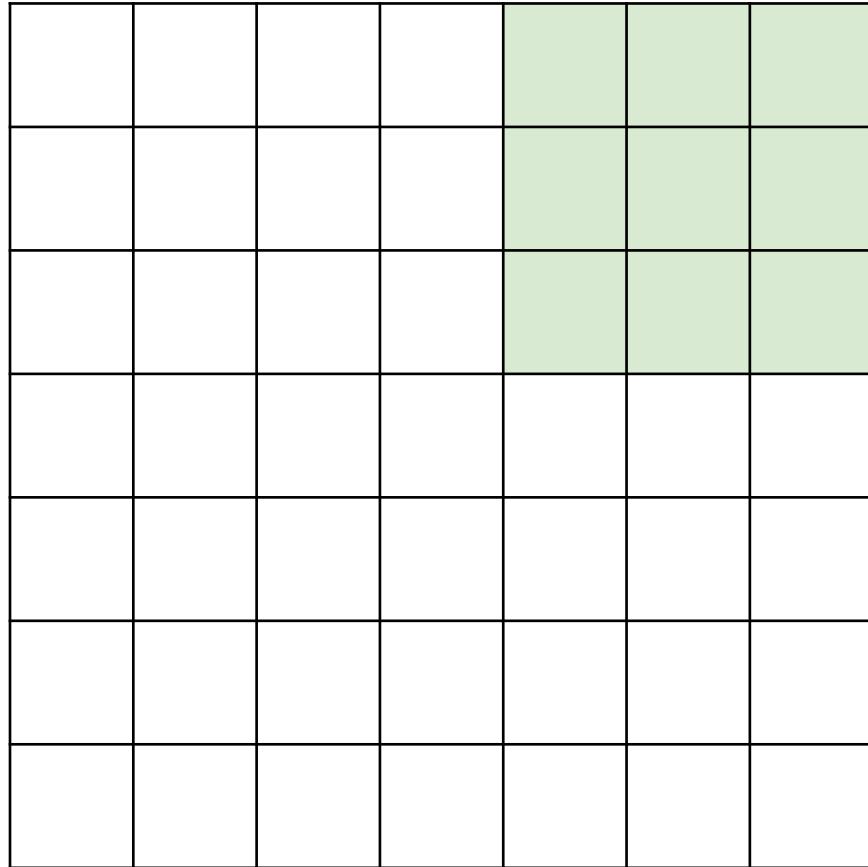
# A closer look at spatial dimensions



7

Input: 7x7  
Filter: 3x3  
Output: 5x5

# A closer look at spatial dimensions



7

Input: 7x7

Filter: 3x3

Output: 5x5

In general:

Input:  $W$

Filter:  $K$

Output:  $W - K + 1$

Problem: Feature  
maps “shrink”  
with each layer!

# A closer look at spatial dimensions

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Input: 7x7

Filter: 3x3

Output: 5x5

In general:

Input:  $W$

Filter:  $K$

Output:  $W - K + 1$

Problem: Feature  
maps “shrink”  
with each layer!

Solution: padding

Add zeros around the input

# A closer look at spatial dimensions

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Input: 7x7

Filter: 3x3

Output: 5x5

In general:

Input:  $W$

Filter:  $K$

Padding:  $P$

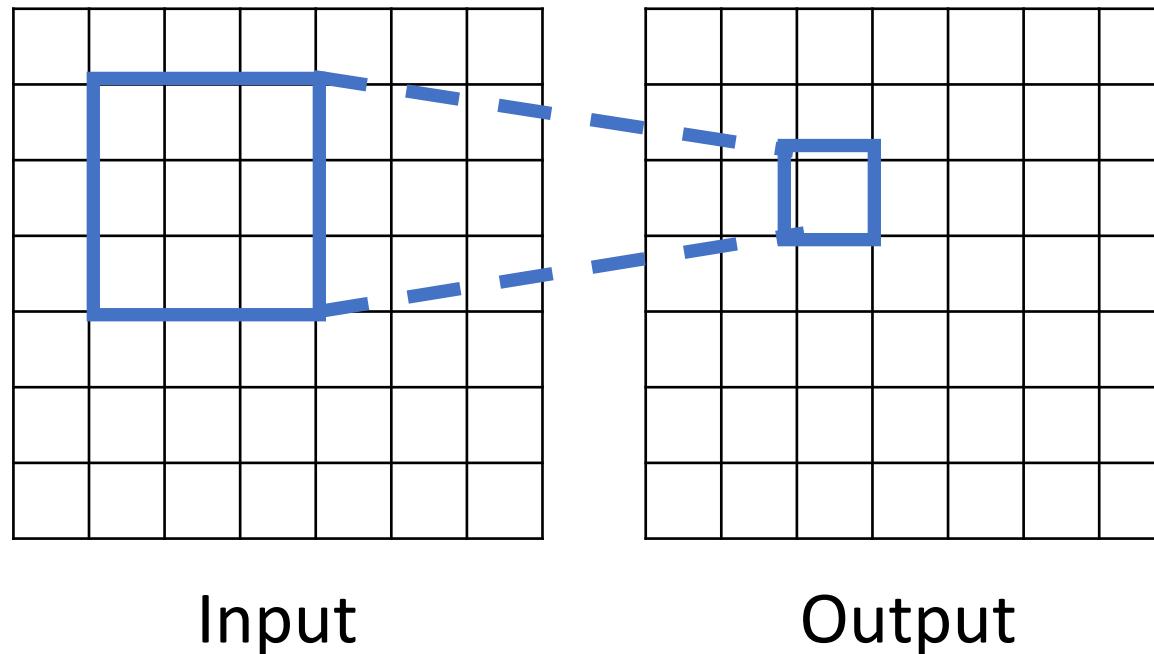
Output:  $W - K + 1 + 2P$

Very common:

Set  $P = (K - 1) / 2$  to  
make output have  
same size as input!

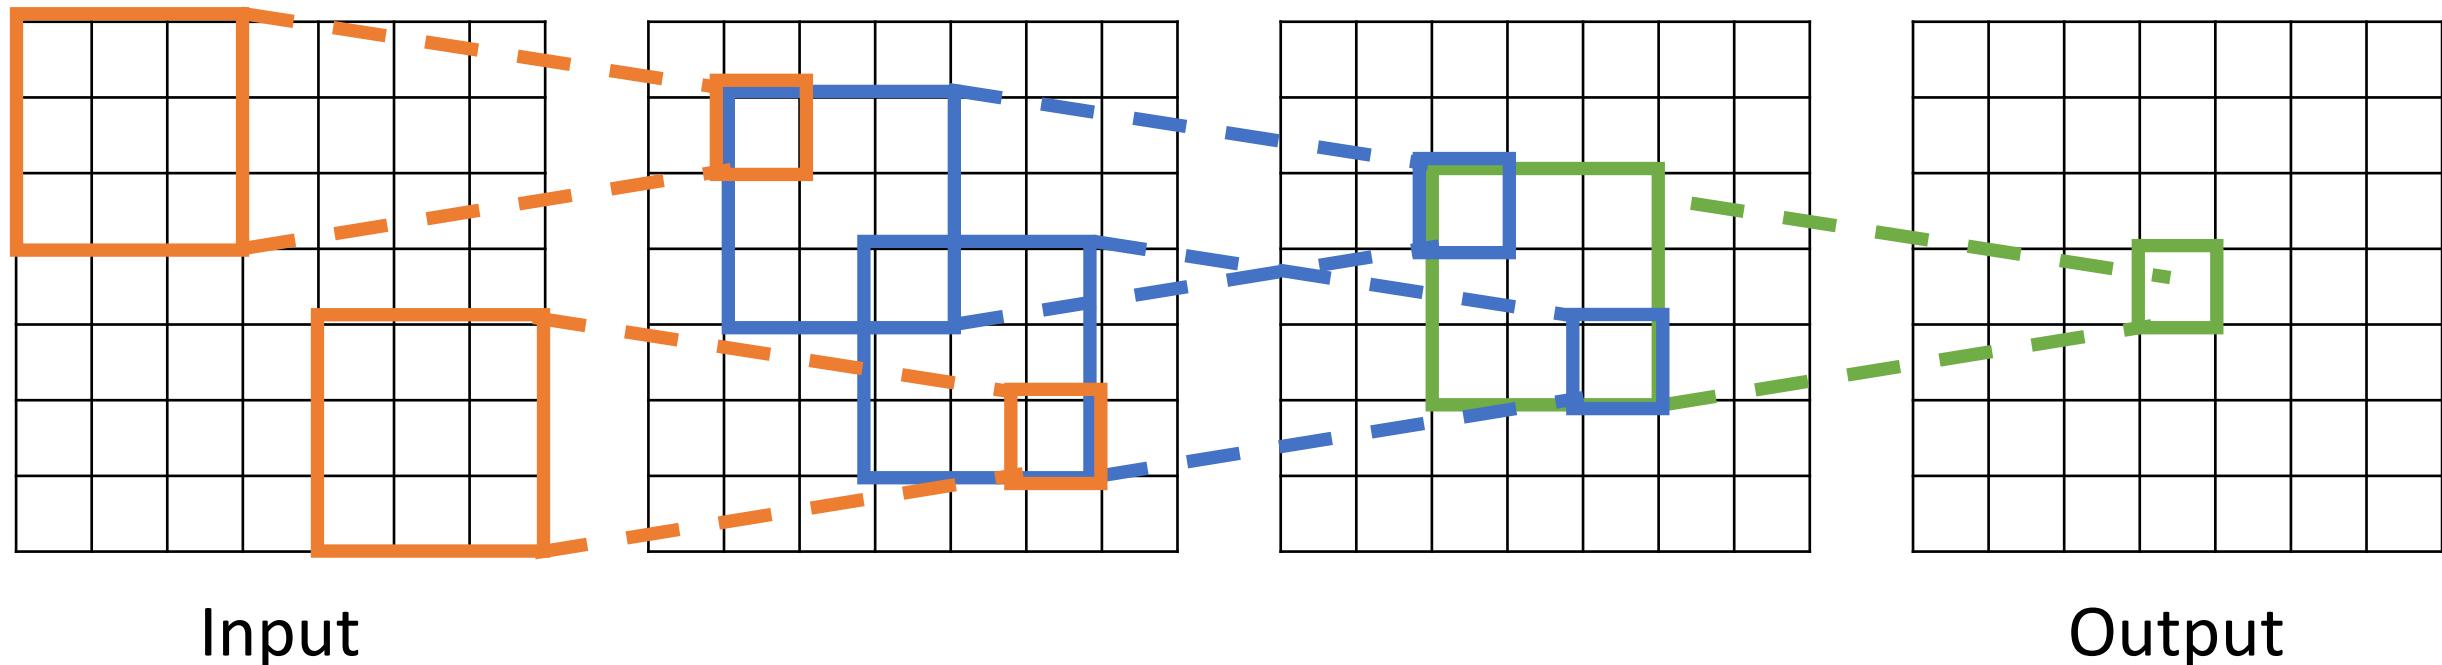
# Receptive Fields

For convolution with kernel size K, each element in the output depends on a  $K \times K$  **receptive field** in the input



# Receptive Fields

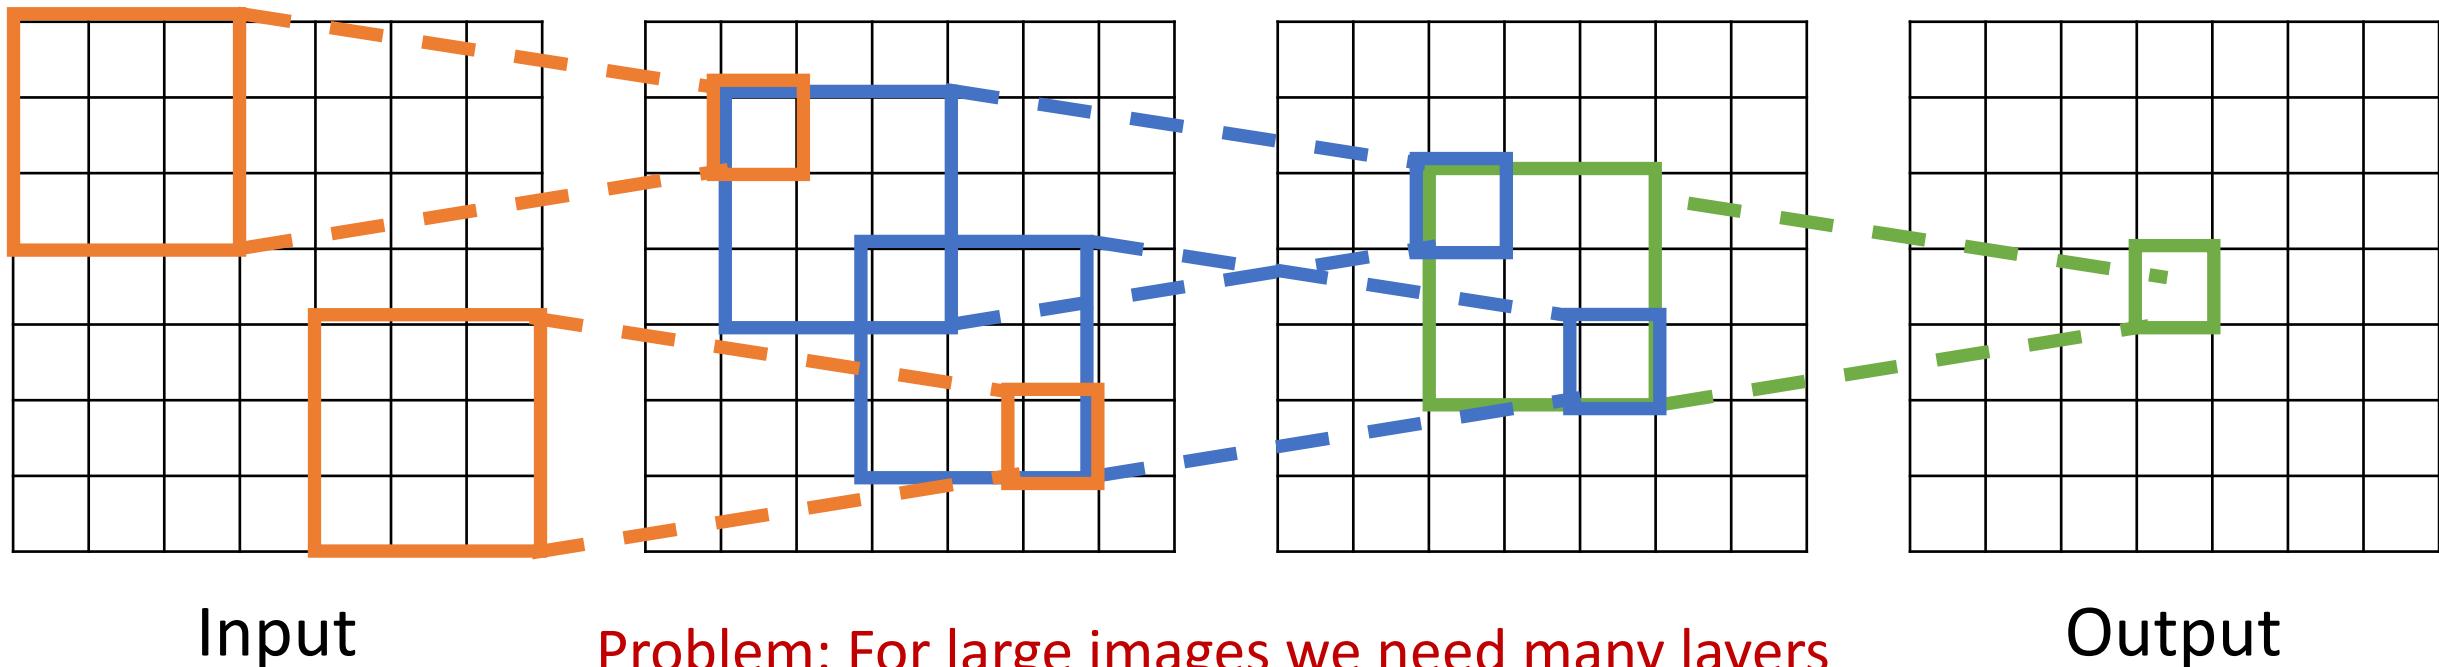
Each successive convolution adds  $K - 1$  to the receptive field size  
With  $L$  layers the receptive field size is  $1 + L * (K - 1)$



Be careful – “receptive field in the input” vs “receptive field in the previous layer”  
Hopefully clear from context!

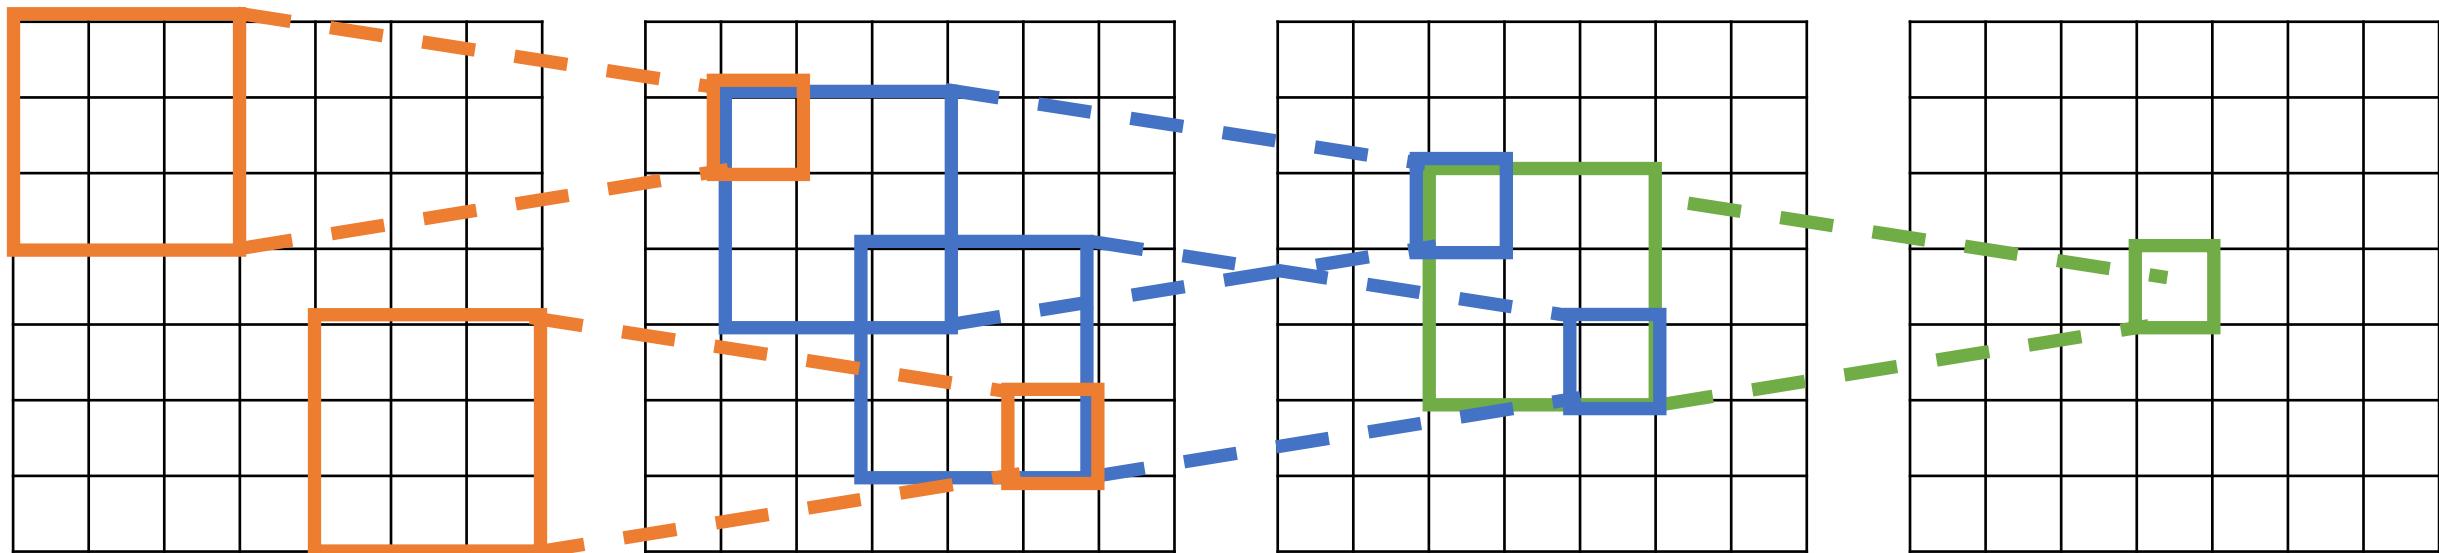
# Receptive Fields

Each successive convolution adds  $K - 1$  to the receptive field size  
With  $L$  layers the receptive field size is  $1 + L * (K - 1)$



# Receptive Fields

Each successive convolution adds  $K - 1$  to the receptive field size  
With  $L$  layers the receptive field size is  $1 + L * (K - 1)$



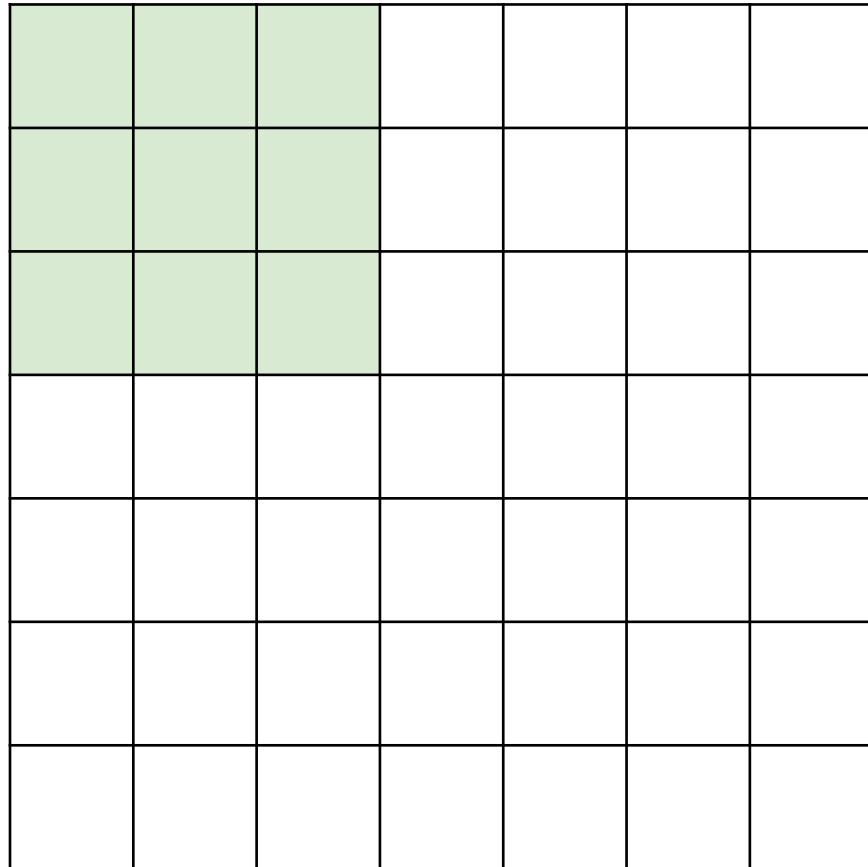
Input

Problem: For large images we need many layers  
for each output to “see” the whole image

Solution: Downsample inside the network

Output

# Strided Convolution

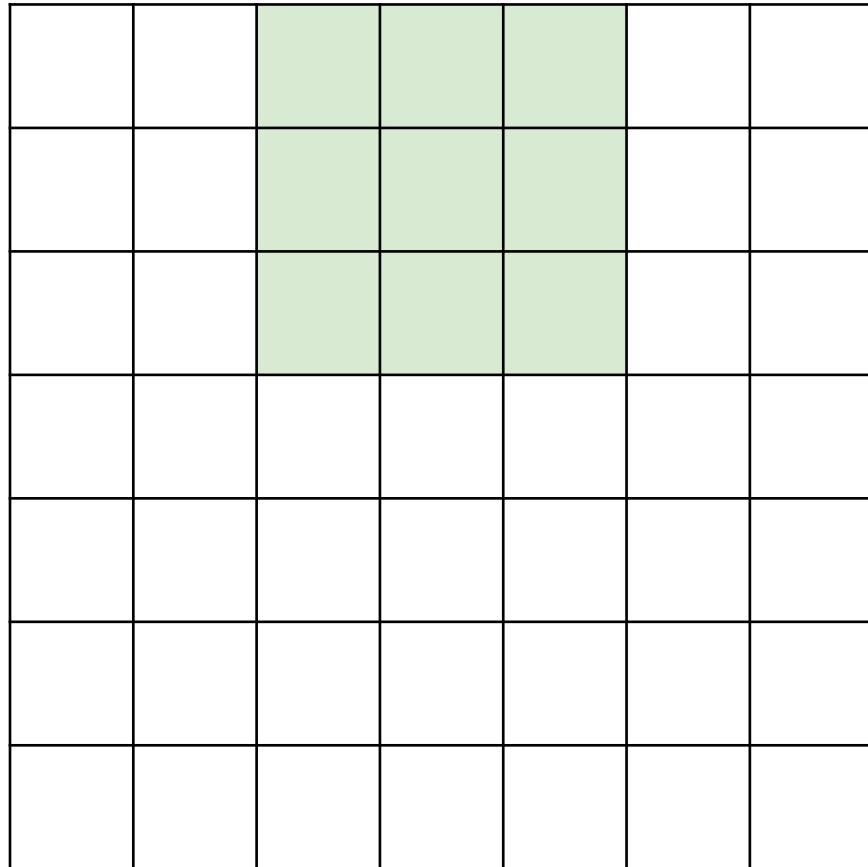


Input: 7x7

Filter: 3x3

Stride: 2

# Strided Convolution

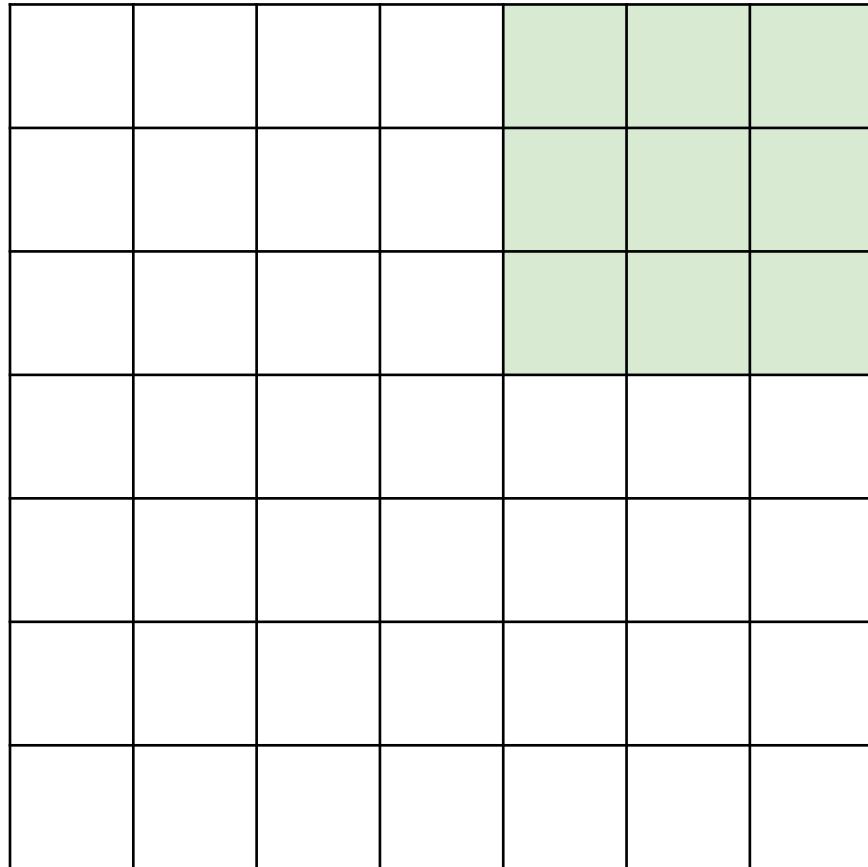


Input: 7x7

Filter: 3x3

Stride: 2

# Strided Convolution



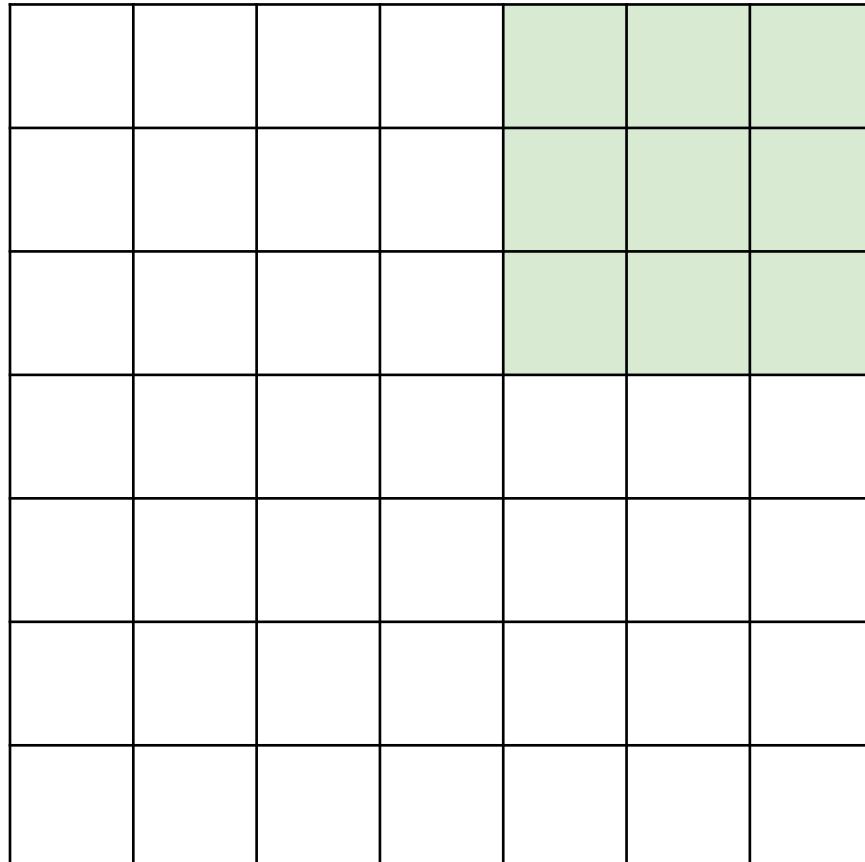
Input: 7x7

Filter: 3x3

Stride: 2

Output: 3x3

# Strided Convolution



Input: 7x7

Filter: 3x3

Stride: 2

Output: 3x3

In general:

Input:  $W$

Filter:  $K$

Padding:  $P$

Stride:  $S$

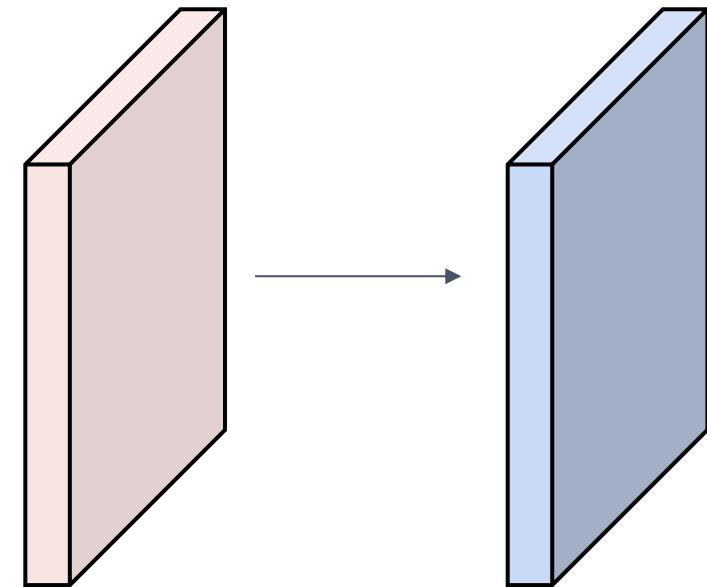
Output:  $(W - K + 2P) / S + 1$

# Convolution Example

Input volume:  $3 \times 32 \times 32$

10 5x5 filters with stride 1, pad 2

Output volume size: ?



# Convolution Example

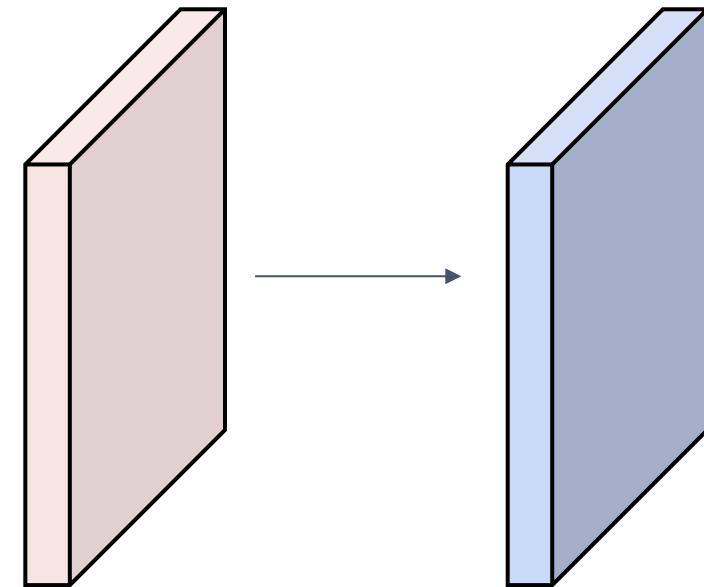
Input volume:  $3 \times 32 \times 32$

**10** **5x5** filters with stride **1**, pad **2**

Output volume size:

$(32+2*2-5)/1+1 = 32$  spatially, so

**10**  $\times 32 \times 32$



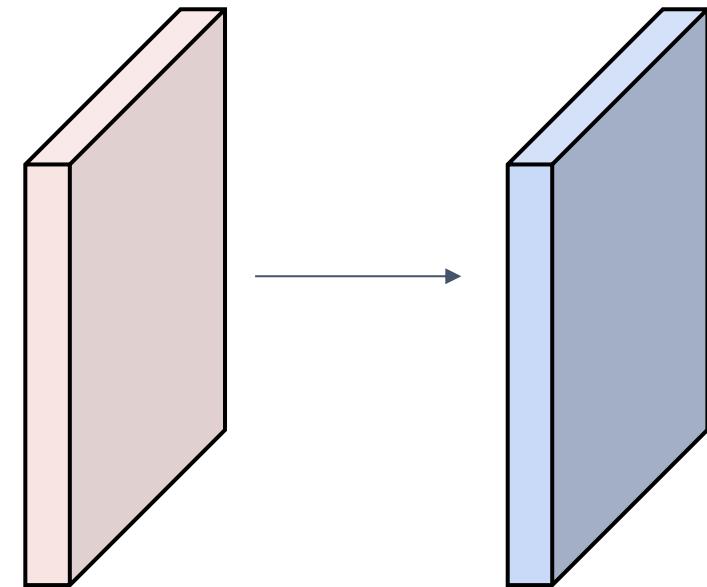
# Convolution Example

Input volume:  $3 \times 32 \times 32$

10 5x5 filters with stride 1, pad 2

Output volume size:  $10 \times 32 \times 32$

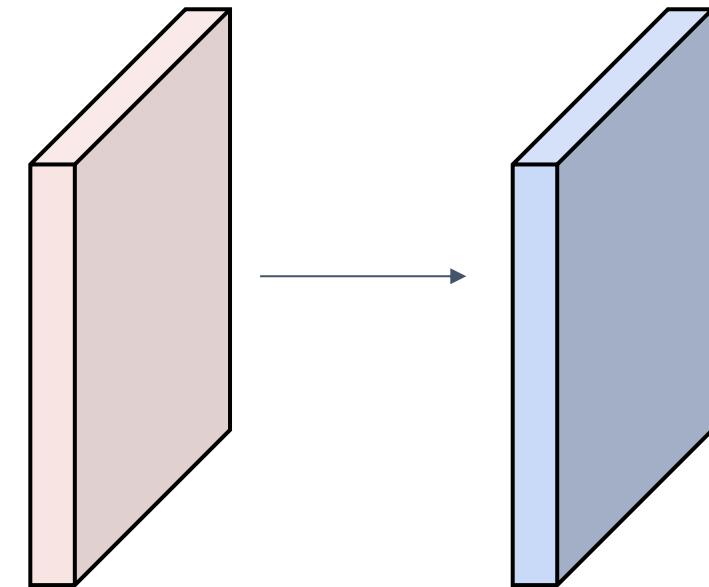
Number of learnable parameters: ?



# Convolution Example

Input volume: **3** x 32 x 32

**10** **5x5** filters with stride 1, pad 2



Output volume size:  $10 \times 32 \times 32$

Number of learnable parameters: **760**

Parameters per filter:  $3 * 5 * 5 + 1$  (for bias) = **76**

**10** filters, so total is **10 \* 76 = 760**

# Convolution Example

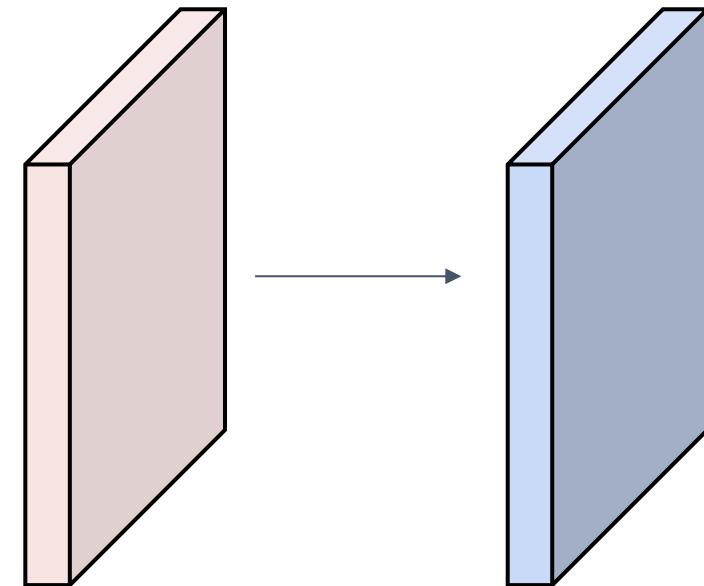
Input volume:  $3 \times 32 \times 32$

10 5x5 filters with stride 1, pad 2

Output volume size:  $10 \times 32 \times 32$

Number of learnable parameters: 760

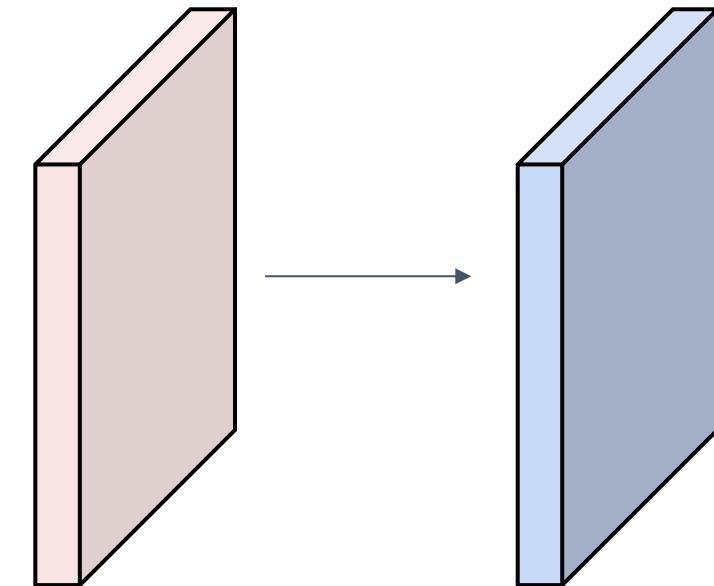
Number of multiply-add operations: ?



# Convolution Example

Input volume: **3 x 32 x 32**

10 **5x5** filters with stride 1, pad 2



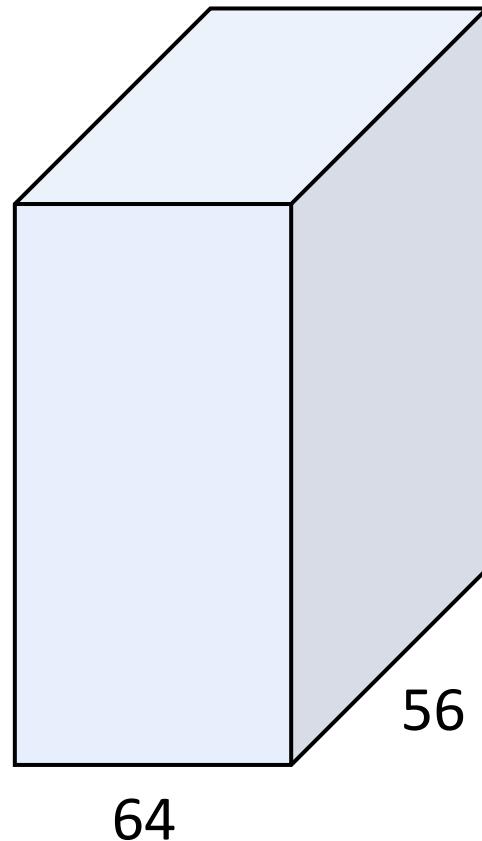
Output volume size: **10 x 32 x 32**

Number of learnable parameters: 760

Number of multiply-add operations: **768,000**

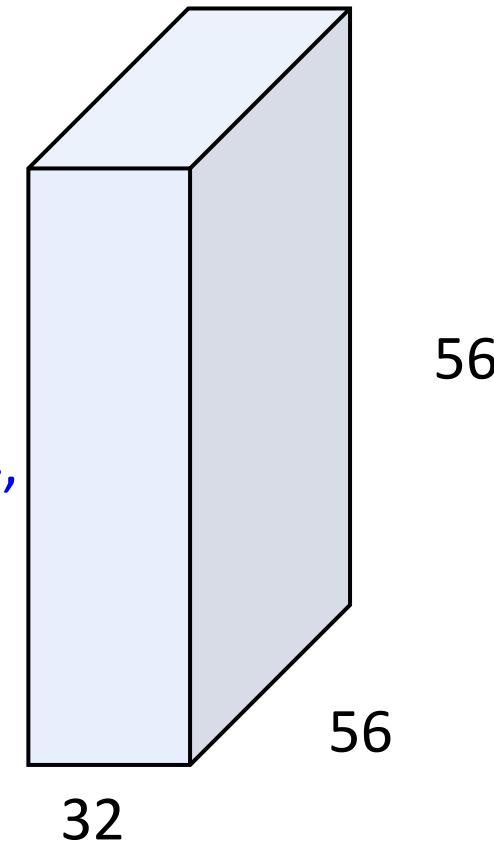
**$10*32*32$**  = 10,240 outputs; each output is the inner product of two **3x5x5** tensors (75 elems); total =  $75*10240 = 768K$

# Example: 1x1 Convolution

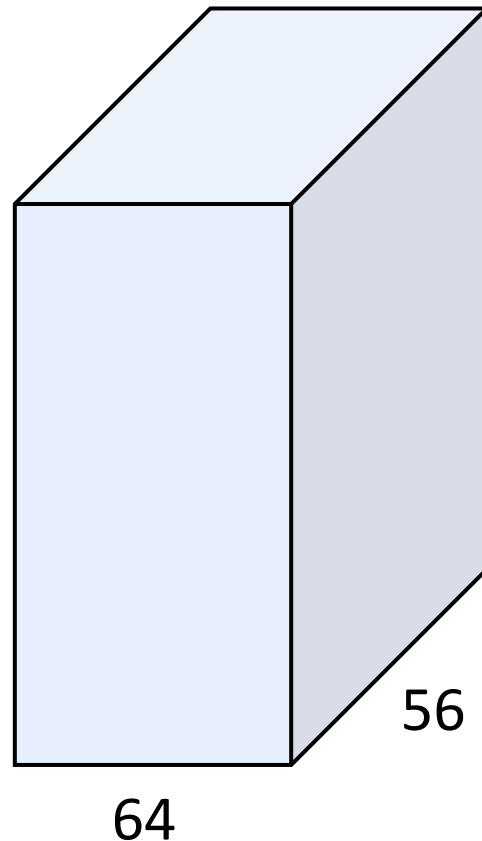


1x1 CONV  
with 32 filters

(each filter has size  $1 \times 1 \times 64$ ,  
and performs a 64-  
dimensional dot product)



# Example: 1x1 Convolution

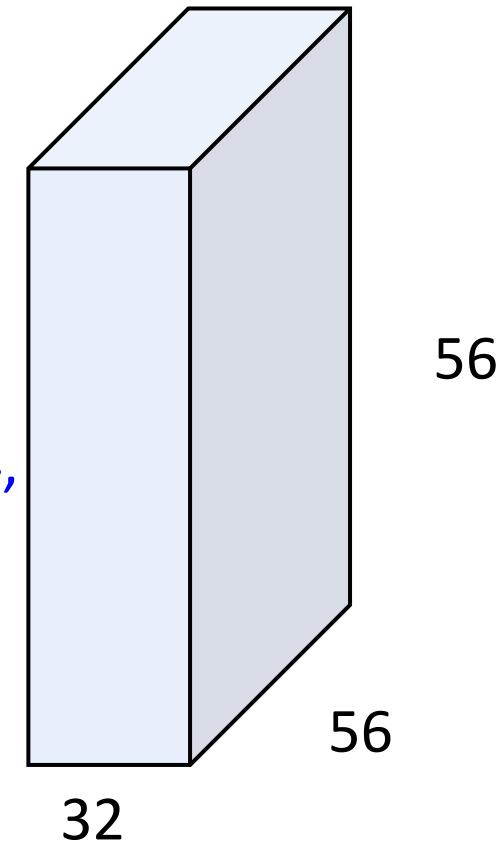


1x1 CONV  
with 32 filters

→

(each filter has size  $1 \times 1 \times 64$ ,  
and performs a 64-dimensional dot product)

Stacking 1x1 conv layers  
gives MLP operating on  
each input position



Lin et al, "Network in Network", ICLR 2014

# Convolution Summary

**Input:**  $C_{in} \times H \times W$

**Hyperparameters:**

- **Kernel size:**  $K_H \times K_W$
- **Number filters:**  $C_{out}$
- **Padding:**  $P$
- **Stride:**  $S$

**Weight matrix:**  $C_{out} \times C_{in} \times K_H \times K_W$

giving  $C_{out}$  filters of size  $C_{in} \times K_H \times K_W$

**Bias vector:**  $C_{out}$

**Output size:**  $C_{out} \times H' \times W'$  where:

- $H' = (H - K + 2P) / S + 1$
- $W' = (W - K + 2P) / S + 1$

# Convolution Summary

**Input:**  $C_{in} \times H \times W$

**Hyperparameters:**

- **Kernel size:**  $K_H \times K_W$
- **Number filters:**  $C_{out}$
- **Padding:**  $P$
- **Stride:**  $S$

**Weight matrix:**  $C_{out} \times C_{in} \times K_H \times K_W$   
giving  $C_{out}$  filters of size  $C_{in} \times K_H \times K_W$

**Bias vector:**  $C_{out}$

**Output size:**  $C_{out} \times H' \times W'$  where:

- $H' = (H - K + 2P) / S + 1$
- $W' = (W - K + 2P) / S + 1$

**Common settings:**

$K_H = K_W$  (Small square filters)

$P = (K - 1) / 2$  ("Same" padding)

$C_{in}, C_{out} = 32, 64, 128, 256$  (powers of 2)

$K = 3, P = 1, S = 1$  (3x3 conv)

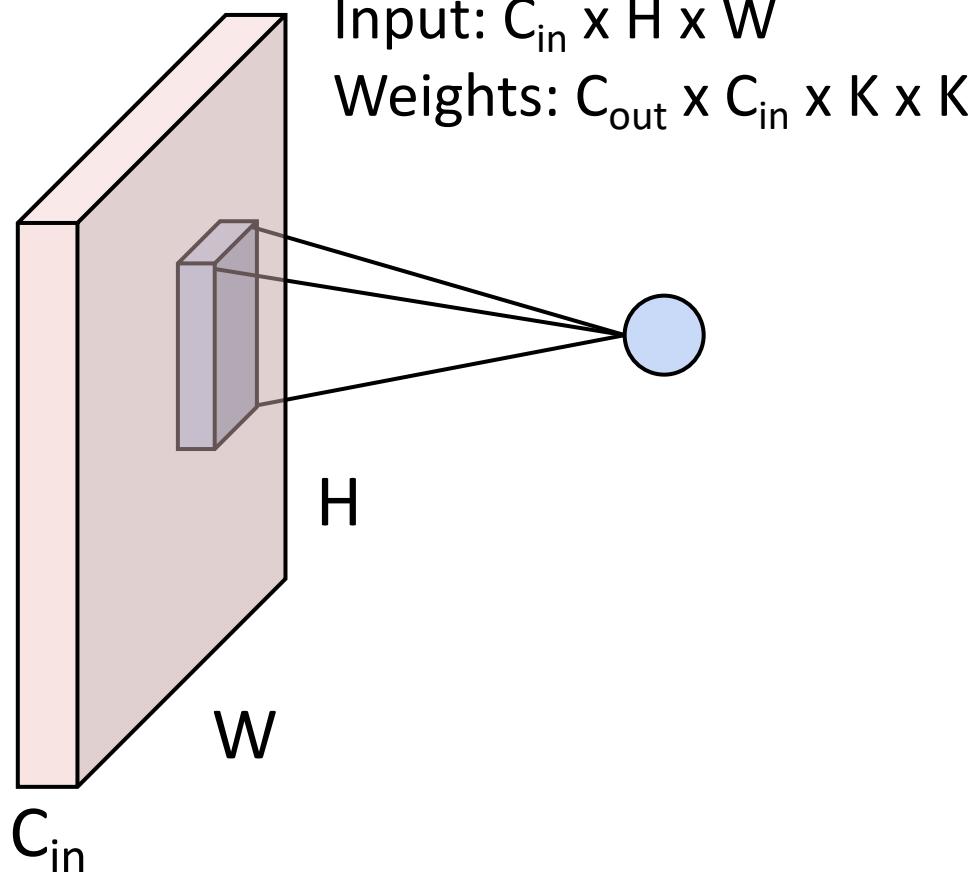
$K = 5, P = 2, S = 1$  (5x5 conv)

$K = 1, P = 0, S = 1$  (1x1 conv)

$K = 3, P = 1, S = 2$  (Downsample by 2)

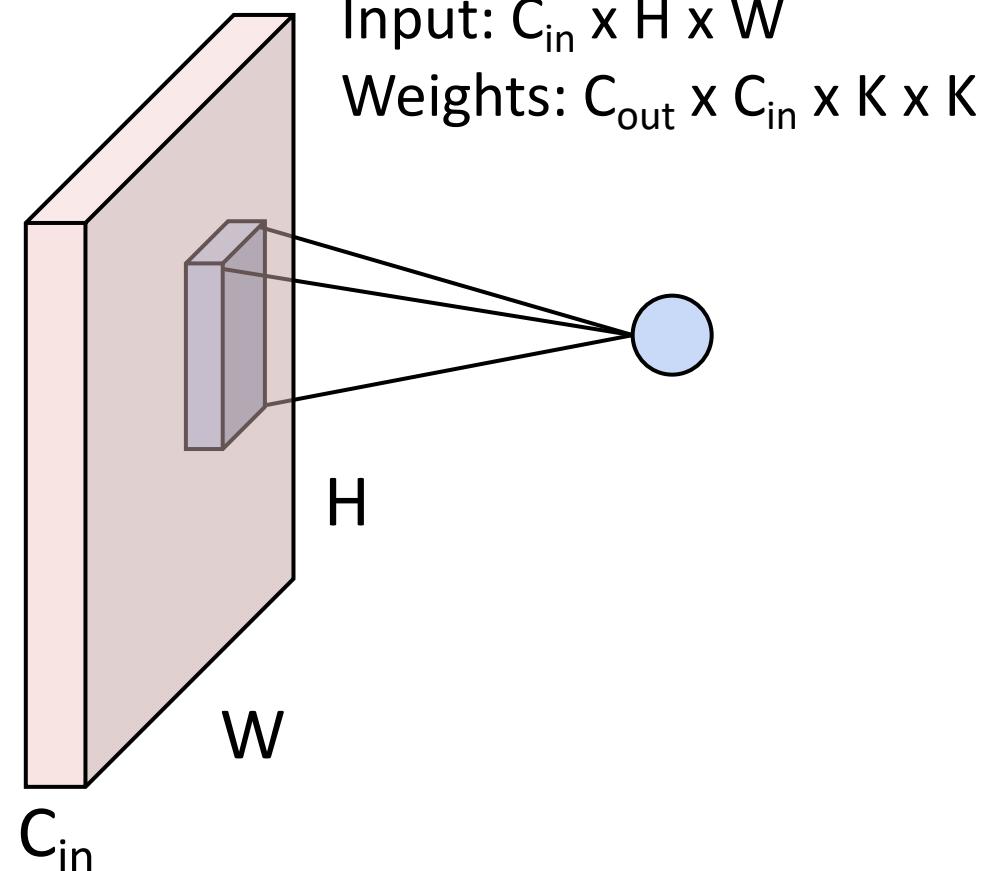
# Other types of convolution

So far: 2D Convolution



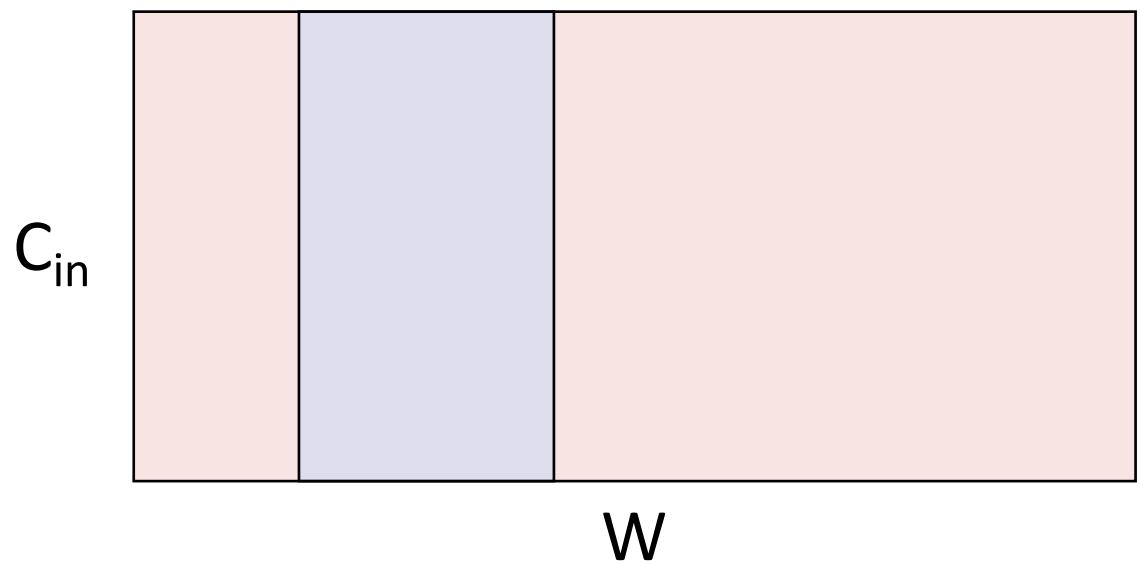
# Other types of convolution

So far: 2D Convolution



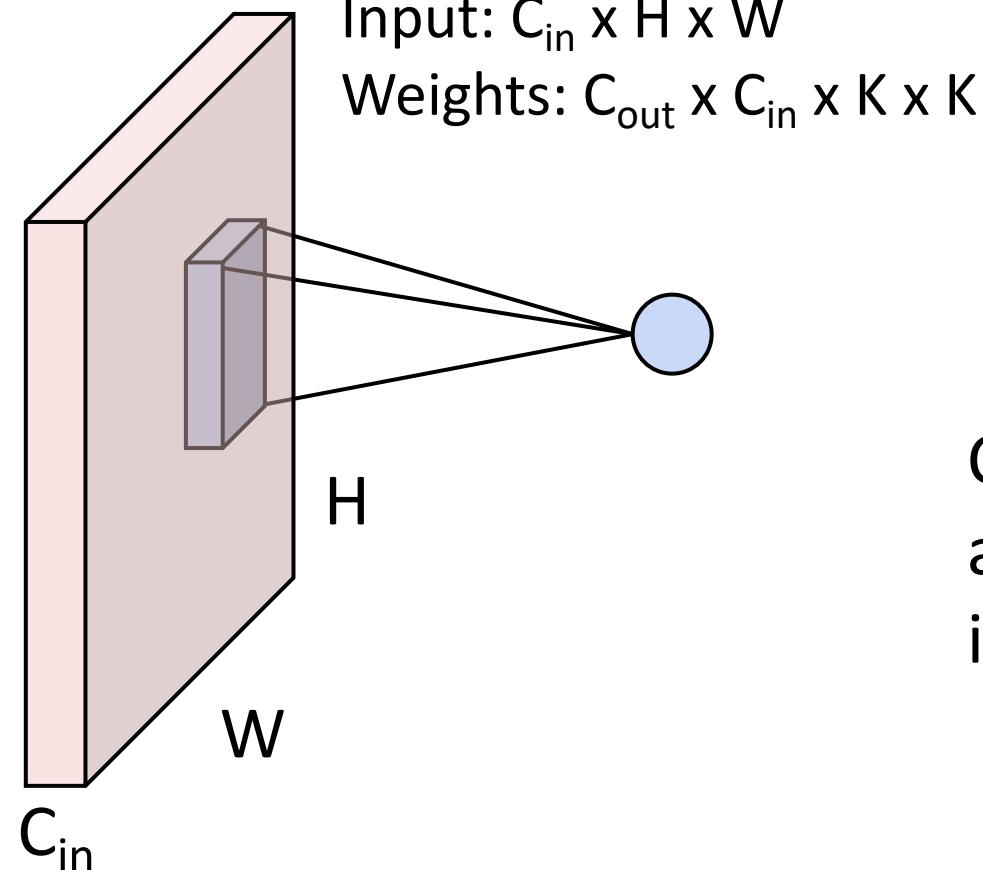
1D Convolution

Input:  $C_{in} \times W$   
Weights:  $C_{out} \times C_{in} \times K$



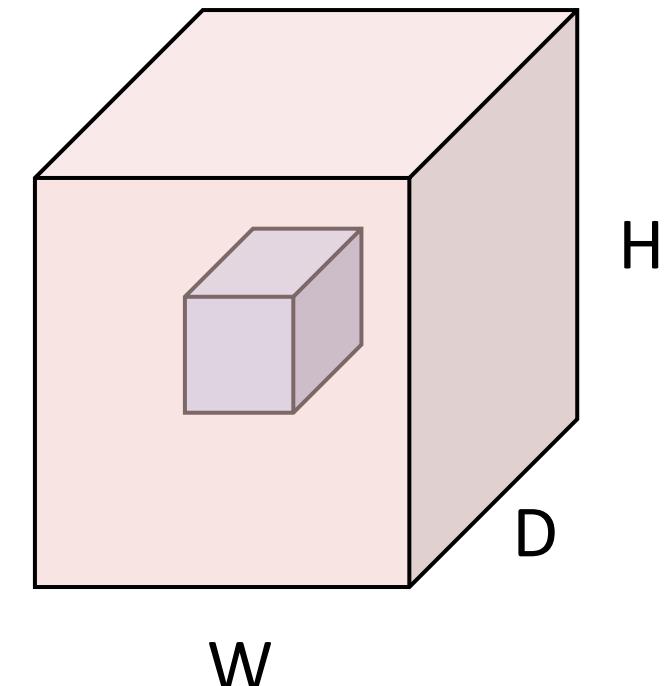
# Other types of convolution

So far: 2D Convolution



3D Convolution

Input:  $C_{in} \times H \times W \times D$   
Weights:  $C_{out} \times C_{in} \times K \times K \times K$



$C_{in}$ -dim vector  
at each point  
in the volume

# PyTorch Convolution Layer

## Conv2d

---

**CLASS** `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros')`

[SOURCE]

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size  $(N, C_{\text{in}}, H, W)$  and output  $(N, C_{\text{out}}, H_{\text{out}}, W_{\text{out}})$  can be precisely described as:

$$\text{out}(N_i, C_{\text{out}_j}) = \text{bias}(C_{\text{out}_j}) + \sum_{k=0}^{C_{\text{in}}-1} \text{weight}(C_{\text{out}_j}, k) \star \text{input}(N_i, k)$$

# PyTorch Convolution Layers

## Conv2d

---

**CLASS** `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros')`

[SOURCE] 

## Conv1d

---

**CLASS** `torch.nn.Conv1d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros')`

[SOURCE] 

## Conv3d

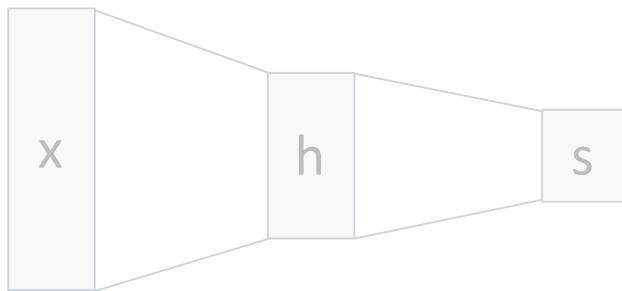
---

**CLASS** `torch.nn.Conv3d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros')`

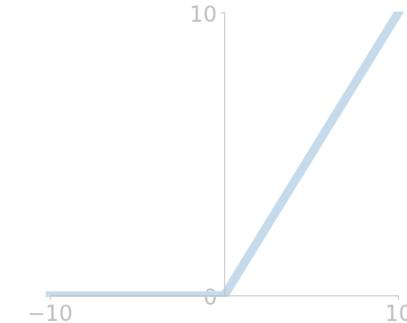
[SOURCE] 

# Components of a Convolutional Network

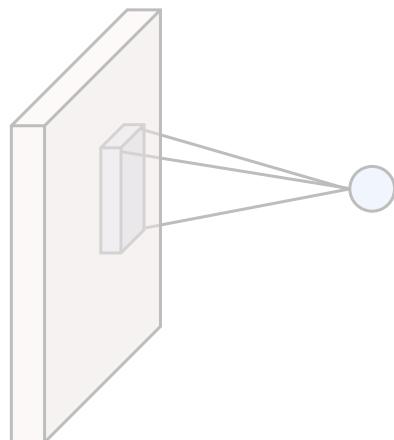
Fully-Connected Layers



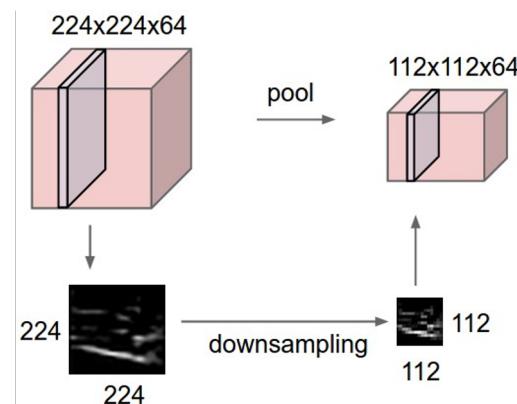
Activation Function



Convolution Layers



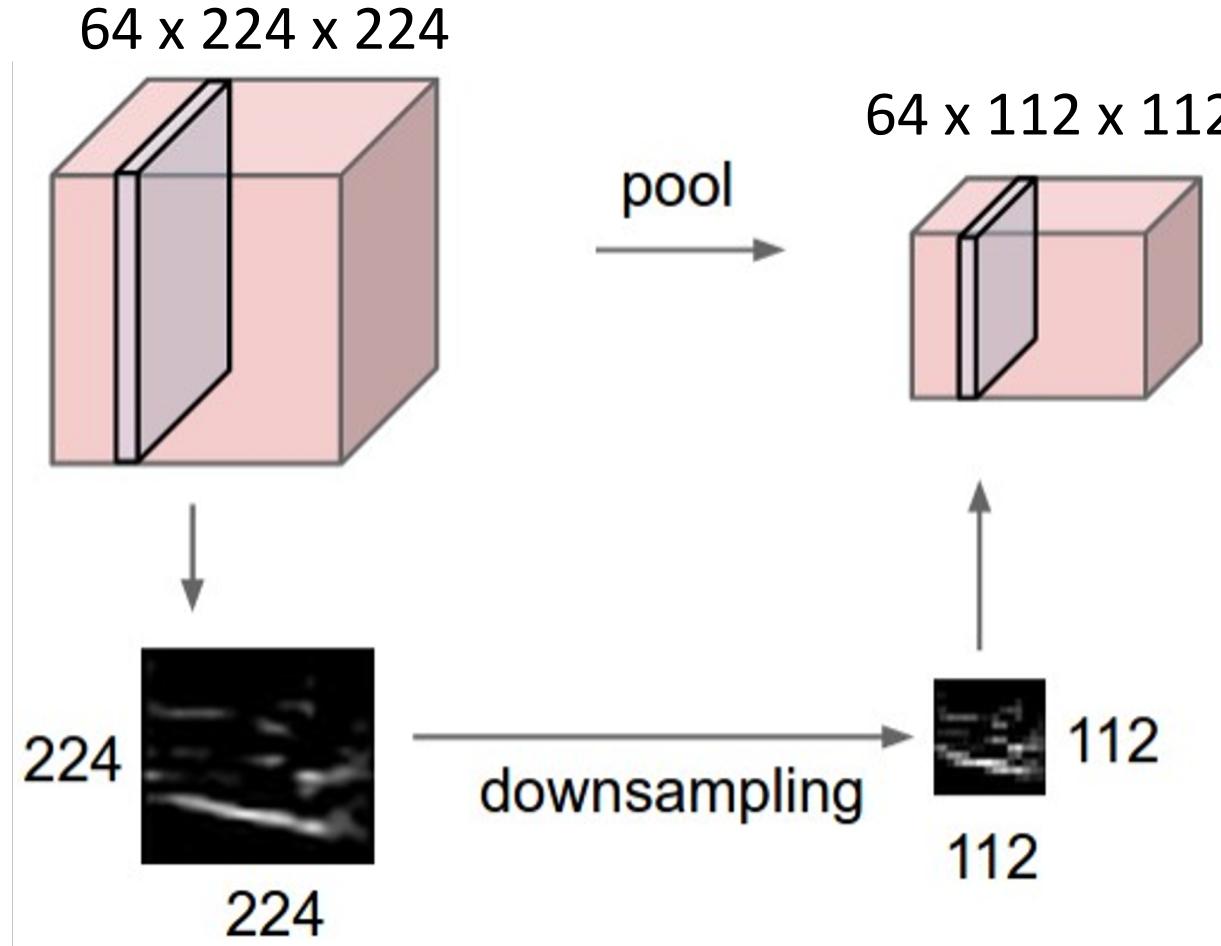
Pooling Layers



Normalization

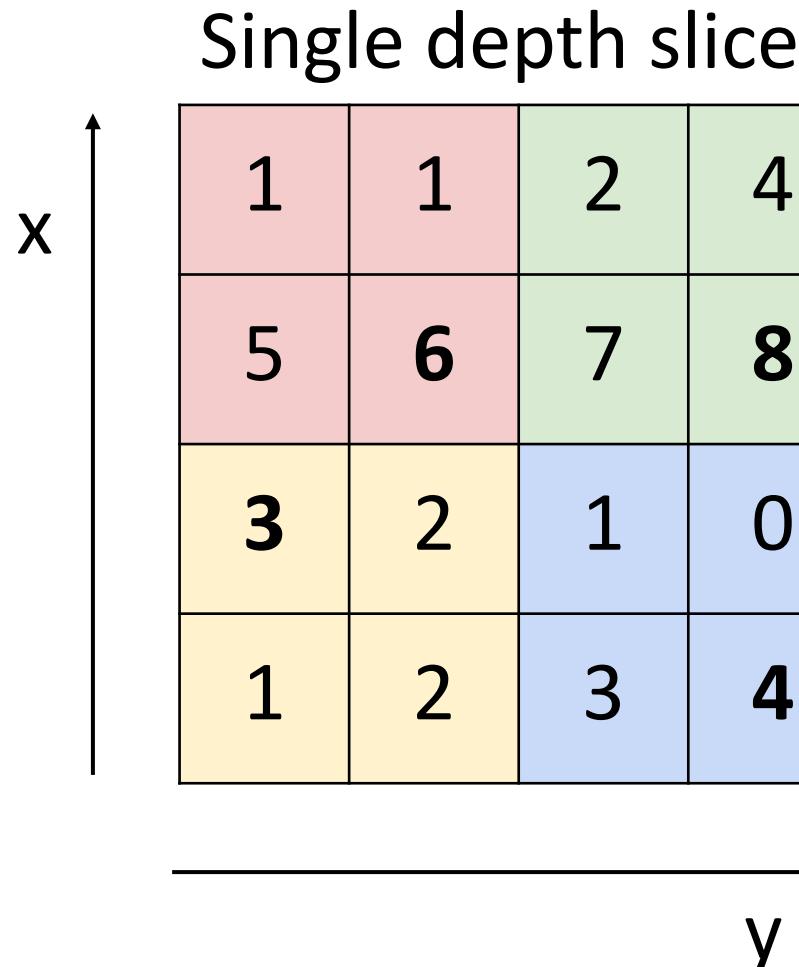
$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}}$$

# Pooling Layers: Another way to downsample

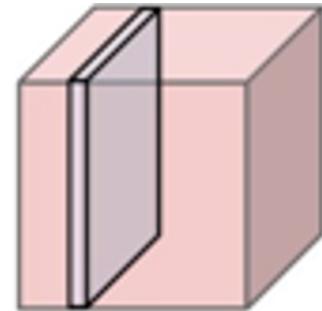


**Hyperparameters:**  
Kernel Size  
Stride  
Pooling function

# Max Pooling



64 x 224 x 224



Max pooling with 2x2  
kernel size and stride 2

6	8
3	4

Introduces **invariance** to  
small spatial shifts  
No learnable parameters!

# Pooling Summary

**Input:**  $C \times H \times W$

**Hyperparameters:**

- Kernel size:  $K$
- Stride:  $S$
- Pooling function (max, avg)

**Output:**  $C \times H' \times W'$  where

- $H' = (H - K) / S + 1$
- $W' = (W - K) / S + 1$

**Learnable parameters:** None!

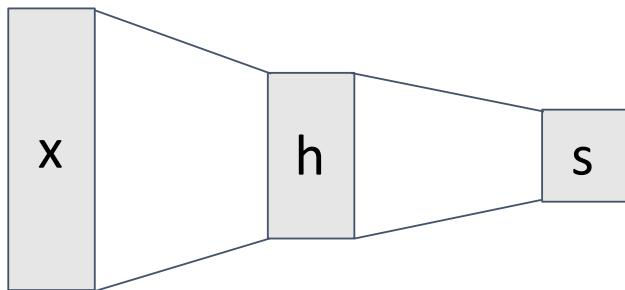
Common settings:

max,  $K = 2, S = 2$

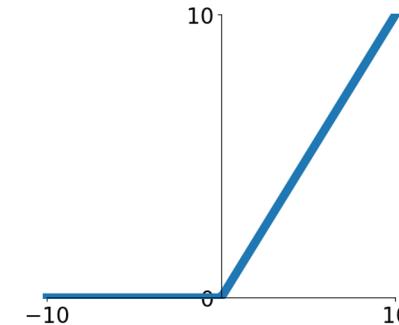
max,  $K = 3, S = 2$  (AlexNet)

# Components of a Convolutional Network

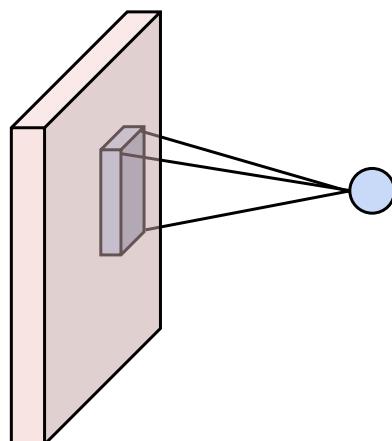
## Fully-Connected Layers



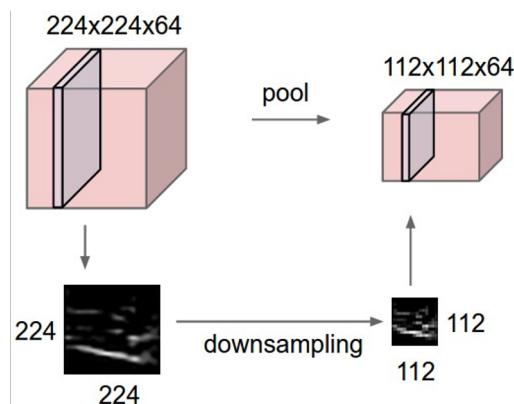
## Activation Function



## Convolution Layers



## Pooling Layers



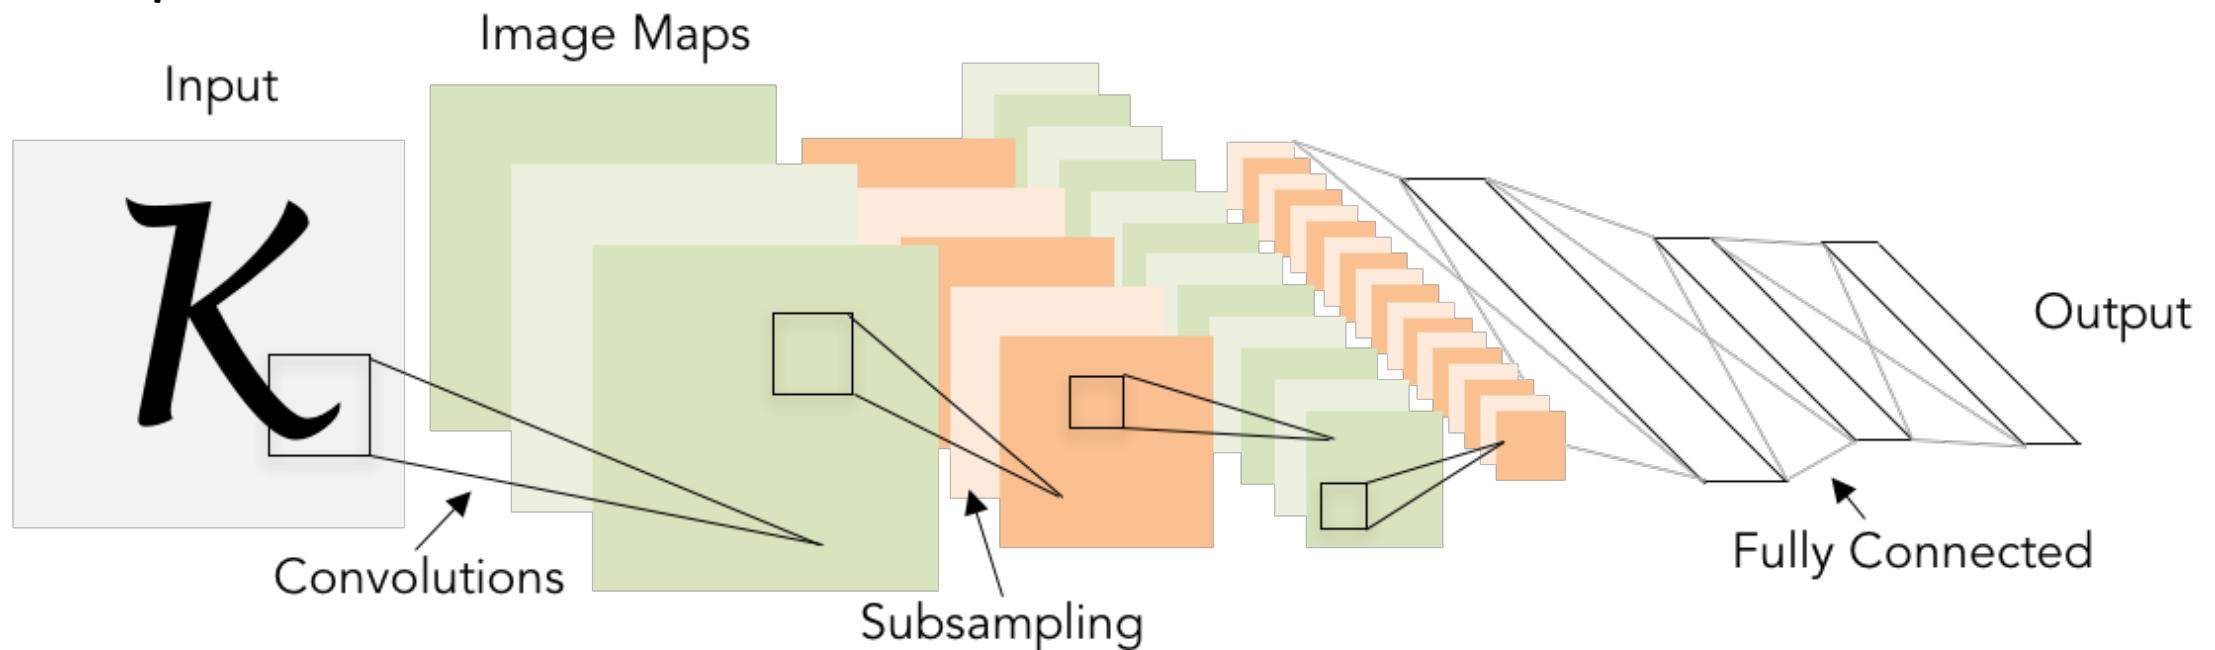
## Normalization

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}}$$

# Convolutional Networks

Classic architecture: [Conv, ReLU, Pool] x N, flatten, [FC, ReLU] x N, FC

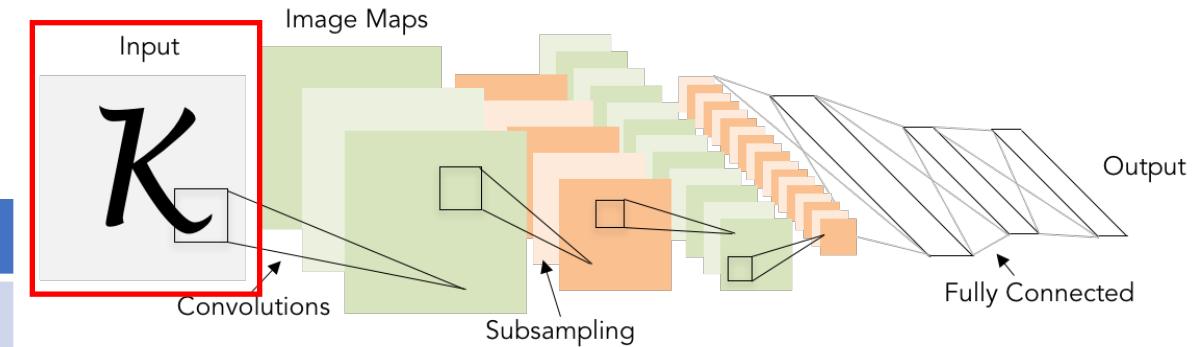
Example: LeNet-5



Lecun et al, "Gradient-based learning applied to document recognition", 1998

# Example: LeNet-5

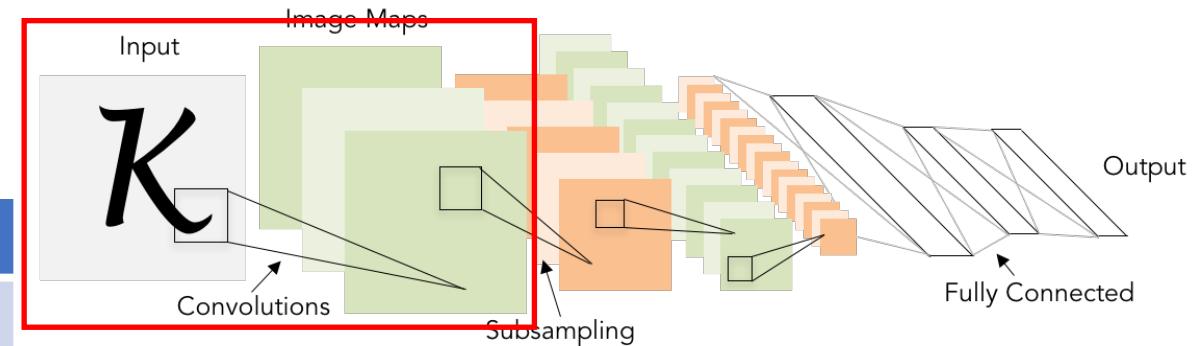
Layer	Output Size	Weight Size
Input	$1 \times 28 \times 28$	



Lecun et al, "Gradient-based learning applied to document recognition", 1998

# Example: LeNet-5

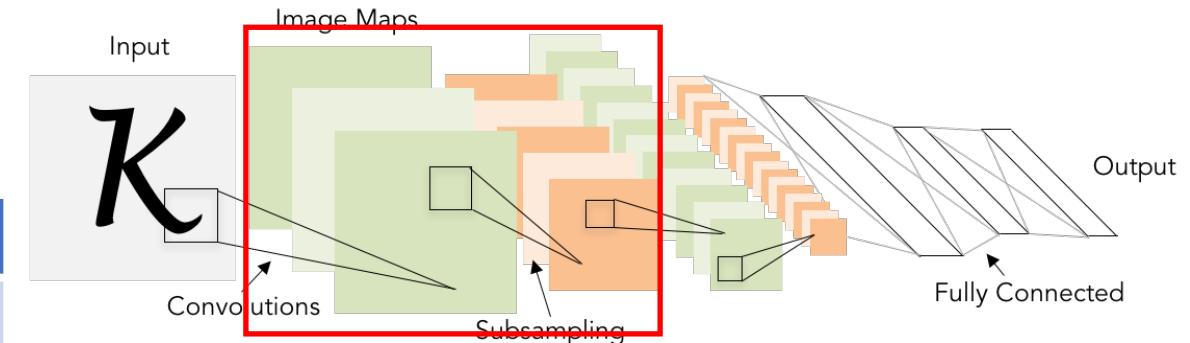
Layer	Output Size	Weight Size
Input	$1 \times 28 \times 28$	
Conv ( $C_{out}=20$ , $K=5$ , $P=2$ , $S=1$ )	$20 \times 28 \times 28$	$20 \times 1 \times 5 \times 5$
ReLU	$20 \times 28 \times 28$	



Lecun et al, "Gradient-based learning applied to document recognition", 1998

# Example: LeNet-5

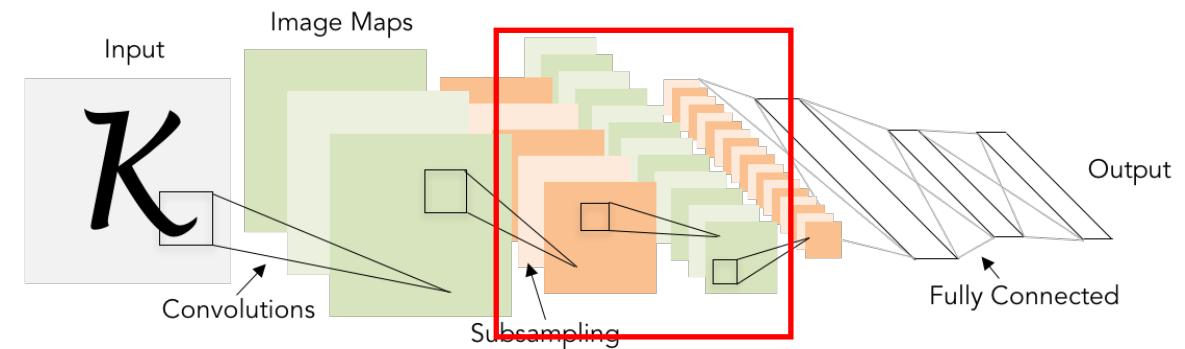
Layer	Output Size	Weight Size
Input	$1 \times 28 \times 28$	
Conv ( $C_{out}=20$ , $K=5$ , $P=2$ , $S=1$ )	$20 \times 28 \times 28$	$20 \times 1 \times 5 \times 5$
ReLU	$20 \times 28 \times 28$	
MaxPool( $K=2$ , $S=2$ )	$20 \times 14 \times 14$	



Lecun et al, "Gradient-based learning applied to document recognition", 1998

# Example: LeNet-5

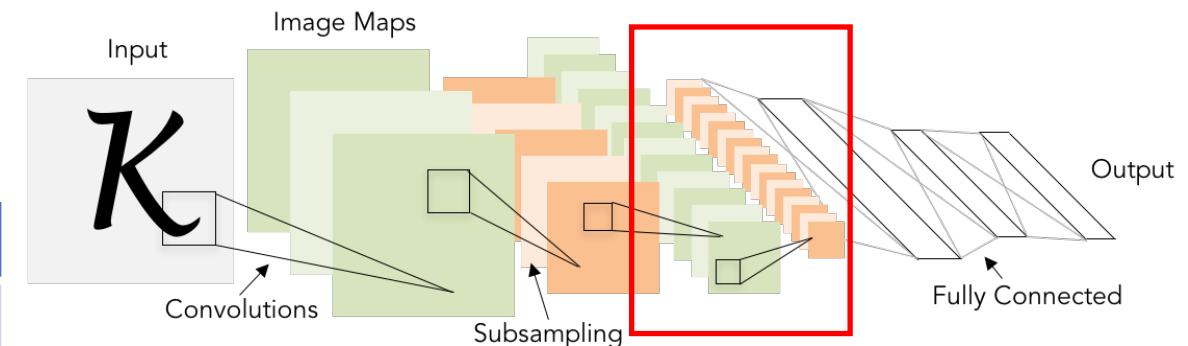
Layer	Output Size	Weight Size
Input	$1 \times 28 \times 28$	
Conv ( $C_{out}=20$ , $K=5$ , $P=2$ , $S=1$ )	$20 \times 28 \times 28$	$20 \times 1 \times 5 \times 5$
ReLU	$20 \times 28 \times 28$	
MaxPool( $K=2$ , $S=2$ )	$20 \times 14 \times 14$	
Conv ( $C_{out}=50$ , $K=5$ , $P=2$ , $S=1$ )	$50 \times 14 \times 14$	$50 \times 20 \times 5 \times 5$
ReLU	$50 \times 14 \times 14$	



Lecun et al, "Gradient-based learning applied to document recognition", 1998

# Example: LeNet-5

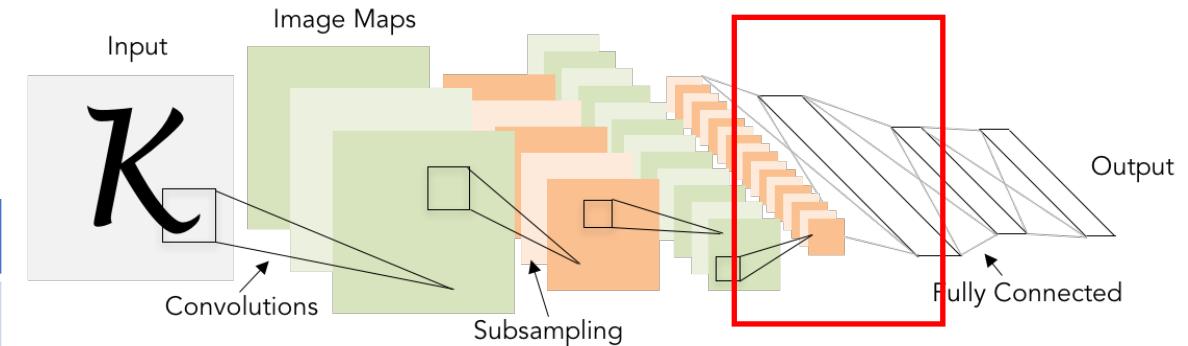
Layer	Output Size	Weight Size
Input	$1 \times 28 \times 28$	
Conv ( $C_{out}=20$ , $K=5$ , $P=2$ , $S=1$ )	$20 \times 28 \times 28$	$20 \times 1 \times 5 \times 5$
ReLU	$20 \times 28 \times 28$	
MaxPool( $K=2$ , $S=2$ )	$20 \times 14 \times 14$	
Conv ( $C_{out}=50$ , $K=5$ , $P=2$ , $S=1$ )	$50 \times 14 \times 14$	$50 \times 20 \times 5 \times 5$
ReLU	$50 \times 14 \times 14$	
MaxPool( $K=2$ , $S=2$ )	$50 \times 7 \times 7$	



Lecun et al, "Gradient-based learning applied to document recognition", 1998

# Example: LeNet-5

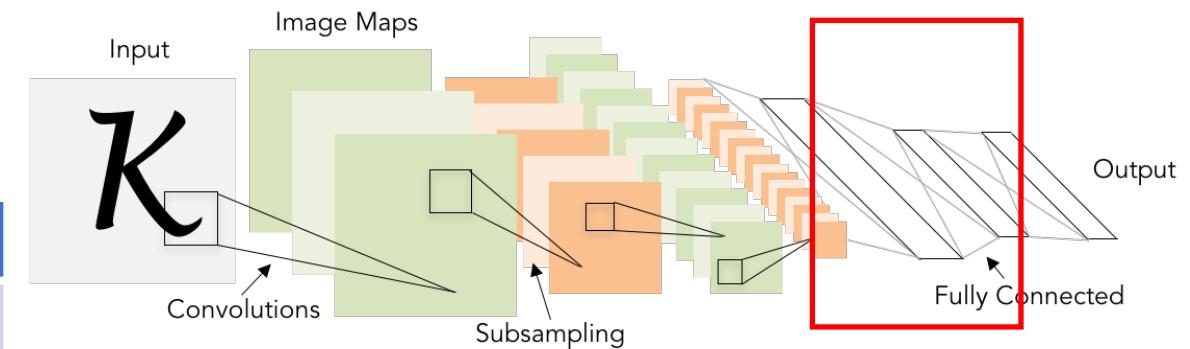
Layer	Output Size	Weight Size
Input	$1 \times 28 \times 28$	
Conv ( $C_{out}=20$ , $K=5$ , $P=2$ , $S=1$ )	$20 \times 28 \times 28$	$20 \times 1 \times 5 \times 5$
ReLU	$20 \times 28 \times 28$	
MaxPool( $K=2$ , $S=2$ )	$20 \times 14 \times 14$	
Conv ( $C_{out}=50$ , $K=5$ , $P=2$ , $S=1$ )	$50 \times 14 \times 14$	$50 \times 20 \times 5 \times 5$
ReLU	$50 \times 14 \times 14$	
MaxPool( $K=2$ , $S=2$ )	$50 \times 7 \times 7$	
Flatten	2450	



Lecun et al, "Gradient-based learning applied to document recognition", 1998

# Example: LeNet-5

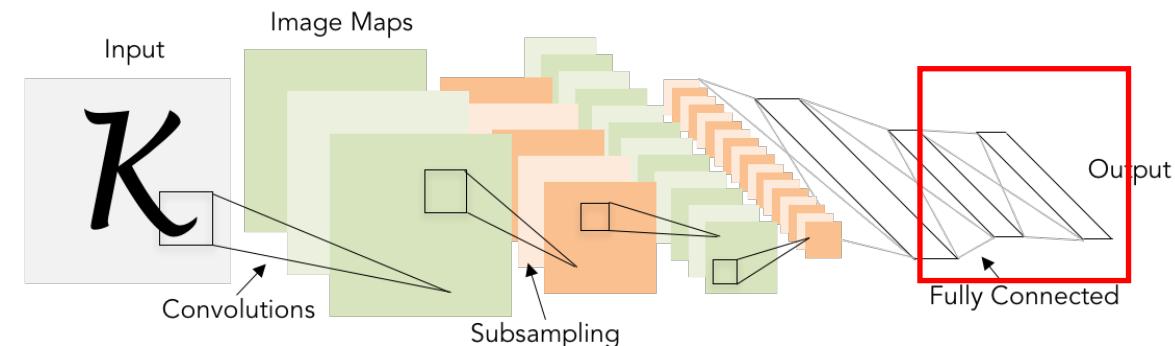
Layer	Output Size	Weight Size
Input	$1 \times 28 \times 28$	
Conv ( $C_{out}=20$ , $K=5$ , $P=2$ , $S=1$ )	$20 \times 28 \times 28$	$20 \times 1 \times 5 \times 5$
ReLU	$20 \times 28 \times 28$	
MaxPool( $K=2$ , $S=2$ )	$20 \times 14 \times 14$	
Conv ( $C_{out}=50$ , $K=5$ , $P=2$ , $S=1$ )	$50 \times 14 \times 14$	$50 \times 20 \times 5 \times 5$
ReLU	$50 \times 14 \times 14$	
MaxPool( $K=2$ , $S=2$ )	$50 \times 7 \times 7$	
Flatten	2450	
Linear (2450 $\rightarrow$ 500)	500	$2450 \times 500$
ReLU	500	



Lecun et al, "Gradient-based learning applied to document recognition", 1998

# Example: LeNet-5

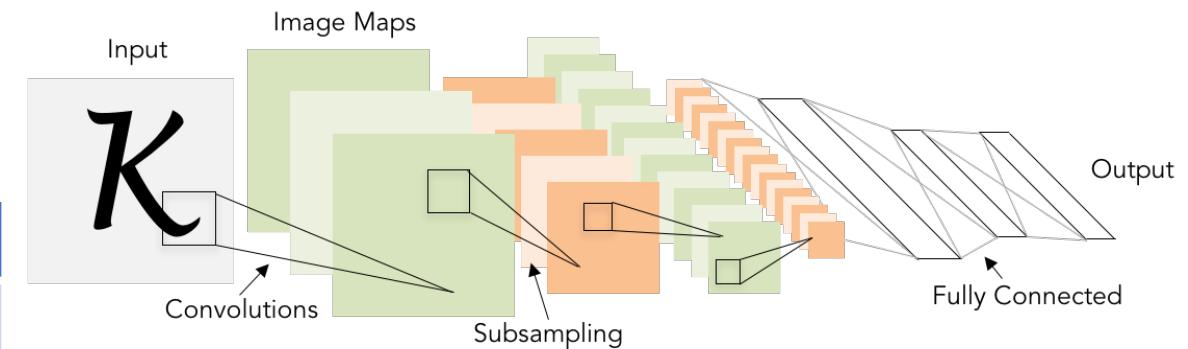
Layer	Output Size	Weight Size
Input	$1 \times 28 \times 28$	
Conv ( $C_{out}=20$ , $K=5$ , $P=2$ , $S=1$ )	$20 \times 28 \times 28$	$20 \times 1 \times 5 \times 5$
ReLU	$20 \times 28 \times 28$	
MaxPool( $K=2$ , $S=2$ )	$20 \times 14 \times 14$	
Conv ( $C_{out}=50$ , $K=5$ , $P=2$ , $S=1$ )	$50 \times 14 \times 14$	$50 \times 20 \times 5 \times 5$
ReLU	$50 \times 14 \times 14$	
MaxPool( $K=2$ , $S=2$ )	$50 \times 7 \times 7$	
Flatten	2450	
Linear (2450 $\rightarrow$ 500)	500	$2450 \times 500$
ReLU	500	
Linear (500 $\rightarrow$ 10)	10	$500 \times 10$



Lecun et al, "Gradient-based learning applied to document recognition", 1998

# Example: LeNet-5

Layer	Output Size	Weight Size
Input	$1 \times 28 \times 28$	
Conv ( $C_{out}=20$ , $K=5$ , $P=2$ , $S=1$ )	$20 \times 28 \times 28$	$20 \times 1 \times 5 \times 5$
ReLU	$20 \times 28 \times 28$	
MaxPool( $K=2$ , $S=2$ )	$20 \times 14 \times 14$	
Conv ( $C_{out}=50$ , $K=5$ , $P=2$ , $S=1$ )	$50 \times 14 \times 14$	$50 \times 20 \times 5 \times 5$
ReLU	$50 \times 14 \times 14$	
MaxPool( $K=2$ , $S=2$ )	$50 \times 7 \times 7$	
Flatten	2450	
Linear (2450 $\rightarrow$ 500)	500	$2450 \times 500$
ReLU	500	
Linear (500 $\rightarrow$ 10)	10	$500 \times 10$



As we go through the network:

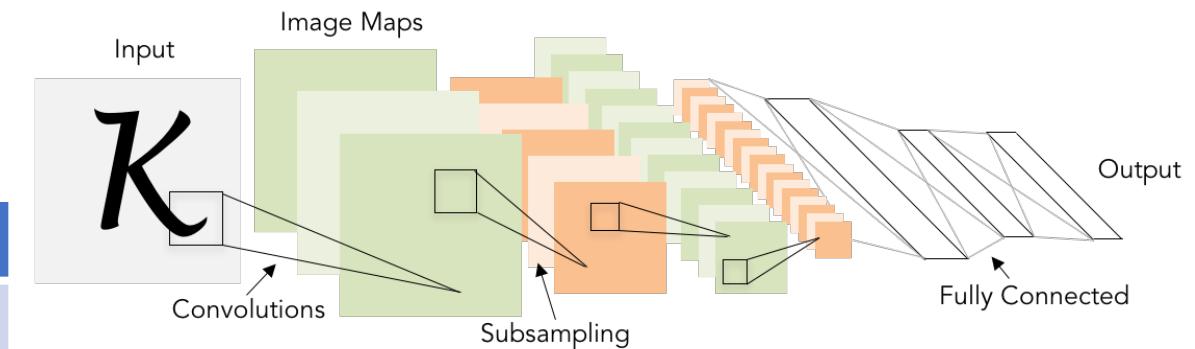
Spatial size **decreases**  
(using pooling or strided conv)

Number of channels **increases**  
(total “volume” is preserved!)

Lecun et al, “Gradient-based learning applied to document recognition”, 1998

# Example: LeNet-5

Layer	Output Size	Weight Size
Input	$1 \times 28 \times 28$	
Conv ( $C_{out}=20$ , $K=5$ , $P=2$ , $S=1$ )	$20 \times 28 \times 28$	$20 \times 1 \times 5 \times 5$
ReLU	$20 \times 28 \times 28$	
MaxPool( $K=2$ , $S=2$ )	$20 \times 14 \times 14$	
Conv ( $C_{out}=50$ , $K=5$ , $P=2$ , $S=1$ )	$50 \times 14 \times 14$	$50 \times 20 \times 5 \times 5$
ReLU	$50 \times 14 \times 14$	
MaxPool( $K=2$ , $S=2$ )	$50 \times 7 \times 7$	
Flatten	2450	
Linear (2450 $\rightarrow$ 500)	500	$2450 \times 500$
ReLU	500	
Linear (500 $\rightarrow$ 10)	10	$500 \times 10$



As we go through the network:

Spatial size **decreases**  
(using pooling or strided conv)

Number of channels **increases**  
(total “volume” is preserved!)

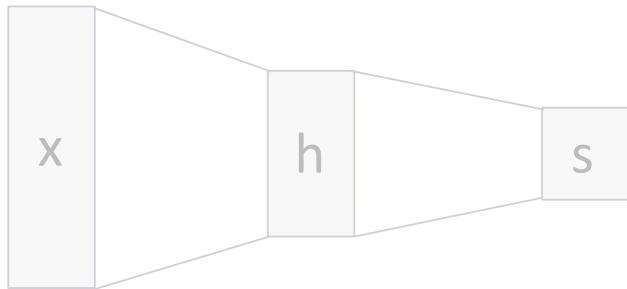
Some modern architectures  
break this trend -- stay tuned!

Lecun et al, "Gradient-based learning applied to document recognition", 1998

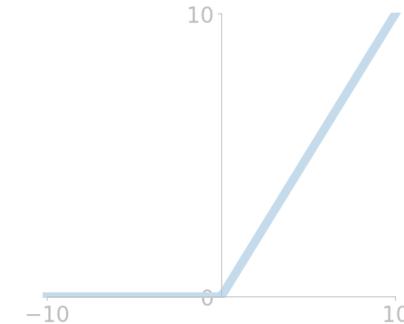
Problem: Deep Networks very hard to train!

# Components of a Convolutional Network

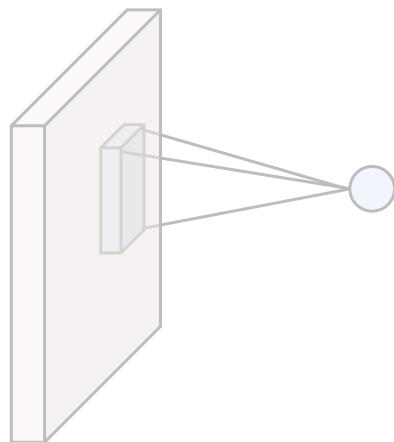
Fully-Connected Layers



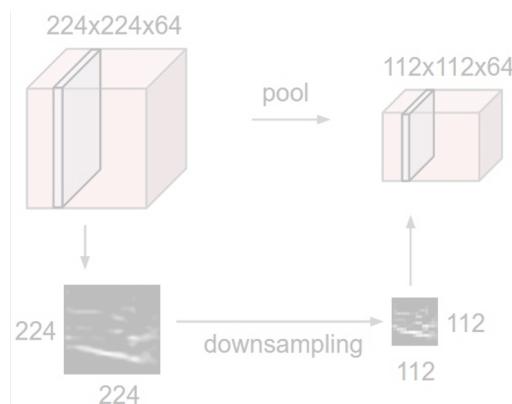
Activation Function



Convolution Layers



Pooling Layers



Normalization

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}}$$

# Batch Normalization

Idea: “Normalize” the outputs of a layer so they have zero mean and unit variance

Why? Helps reduce “internal covariate shift”, improves optimization

We can normalize a batch of activations like this:

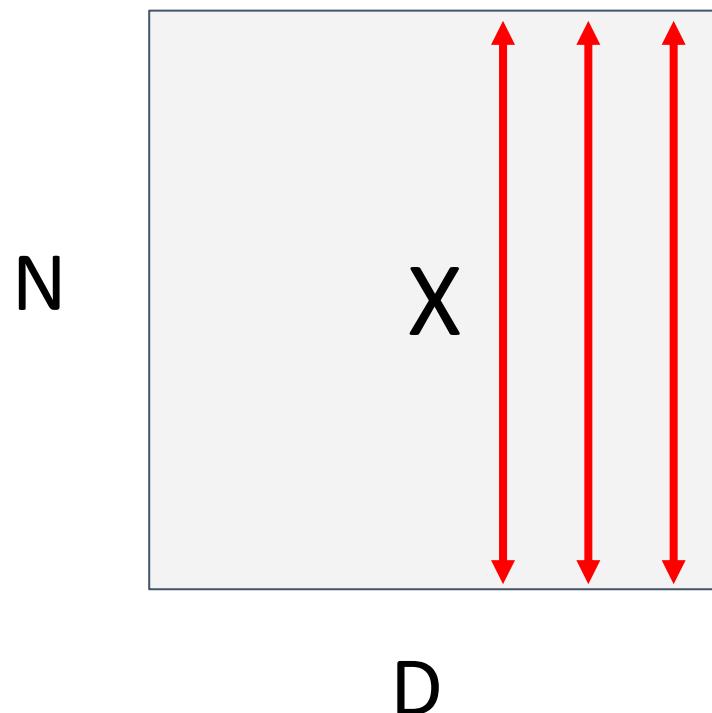
$$\hat{x} = \frac{x - E[x]}{\sqrt{Var[x]}}$$

This is a **differentiable function**, so we can use it as an operator in our networks and backprop through it!

Ioffe and Szegedy, “Batch normalization: Accelerating deep network training by reducing internal covariate shift”, ICML 2015

# Batch Normalization

**Input:**  $x \in \mathbb{R}^{N \times D}$



$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

Per-channel  
mean, shape is D

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

Per-channel  
std, shape is D

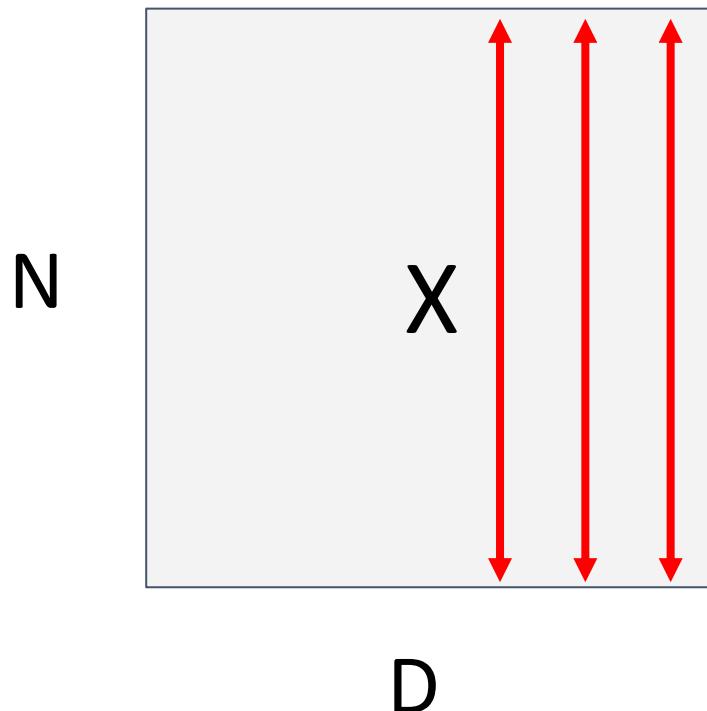
$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

Normalized x,  
Shape is N x D

Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

# Batch Normalization

**Input:**  $x \in \mathbb{R}^{N \times D}$



$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

Per-channel  
mean, shape is D

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

Per-channel  
std, shape is D

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

Normalized x,  
Shape is N x D

Problem: What if zero-mean, unit  
variance is too hard of a constraint?

Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

# Batch Normalization

**Input:**  $x \in \mathbb{R}^{N \times D}$

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

Per-channel  
mean, shape is D

**Learnable scale and shift parameters:**

$$\gamma, \beta \in \mathbb{R}^D$$

Learning  $\gamma = \sigma$ ,  $\beta = \mu$  will recover the identity function (in expectation)

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

Per-channel std, shape is D

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

Normalized x,  
Shape is N x D

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$

Output,  
Shape is N x D

# Batch Normalization

**Problem:** Estimates depend on minibatch; can't do this at test-time!

**Input:**  $x \in \mathbb{R}^{N \times D}$

**Learnable scale and shift parameters:**

$$\gamma, \beta \in \mathbb{R}^D$$

Learning  $\gamma = \sigma$ ,  $\beta = \mu$  will recover the identity function (in expectation)

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

Per-channel mean, shape is D

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

Per-channel std, shape is D

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

Normalized x,  
Shape is N x D

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$

Output,  
Shape is N x D

# Batch Normalization: Test-Time

**Input:**  $x \in \mathbb{R}^{N \times D}$

$\mu_j$  = (Running) average of values seen during training Per-channel mean, shape is D

**Learnable scale and shift parameters:**

$\gamma, \beta \in \mathbb{R}^D$

$\sigma_j^2$  = (Running) average of values seen during training Per-channel std, shape is D

Learning  $\gamma = \sigma, \beta = \mu$  will recover the identity function (in expectation)

$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$  Normalized x, Shape is N x D

$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$  Output, Shape is N x D

# Batch Normalization: Test-Time

**Input:**  $x \in \mathbb{R}^{N \times D}$

$\mu_j$  = (Running) average of values seen during training

Per-channel mean, shape is D

**Learnable scale and shift parameters:**

$\gamma, \beta \in \mathbb{R}^D$

Learning  $\gamma = \sigma, \beta = \mu$  will recover the identity function (in expectation)

$$\mu_j^{test} = 0$$

For each training iteration:

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

$$\mu_j^{test} = 0.99 \mu_j^{test} + 0.01 \mu_j$$

(Similar for  $\sigma$ )

# Batch Normalization: Test-Time

**Input:**  $x \in \mathbb{R}^{N \times D}$

$\mu_j$  = (Running) average of values seen during training Per-channel mean, shape is D

**Learnable scale and shift parameters:**

$\gamma, \beta \in \mathbb{R}^D$

$\sigma_j^2$  = (Running) average of values seen during training Per-channel std, shape is D

Learning  $\gamma = \sigma, \beta = \mu$  will recover the identity function (in expectation)

$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$  Normalized x, Shape is N x D

$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$  Output, Shape is N x D

# Batch Normalization: Test-Time

**Input:**  $x \in \mathbb{R}^{N \times D}$

$\mu_j$  = (Running) average of values seen during training  
Per-channel mean, shape is D

**Learnable scale and shift parameters:**

$\gamma, \beta \in \mathbb{R}^D$

During testing batchnorm becomes a linear operator!

Can be fused with the previous fully-connected or conv layer

$\sigma_j^2$  = (Running) average of values seen during training  
Per-channel std, shape is D

$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$   
Normalized x,  
Shape is N x D

$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$   
Output,  
Shape is N x D

# Batch Normalization for ConvNets

Batch Normalization for  
**fully-connected** networks

$$x : N \times D$$

Normalize

$$\mu, \sigma : 1 \times D$$

$$\gamma, \beta : 1 \times D$$

$$y = \frac{(x - \mu)}{\sigma} \gamma + \beta$$

Batch Normalization for  
**convolutional** networks  
(Spatial Batchnorm, BatchNorm2D)

$$x : N \times C \times H \times W$$

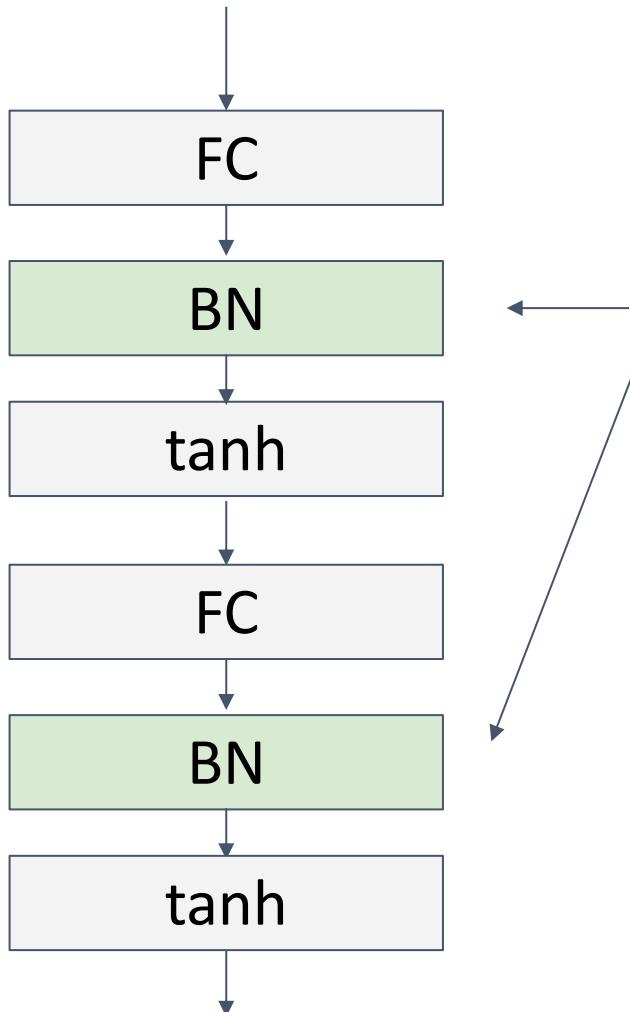
Normalize

$$\mu, \sigma : 1 \times C \times 1 \times 1$$

$$\gamma, \beta : 1 \times C \times 1 \times 1$$

$$y = \frac{(x - \mu)}{\sigma} \gamma + \beta$$

# Batch Normalization

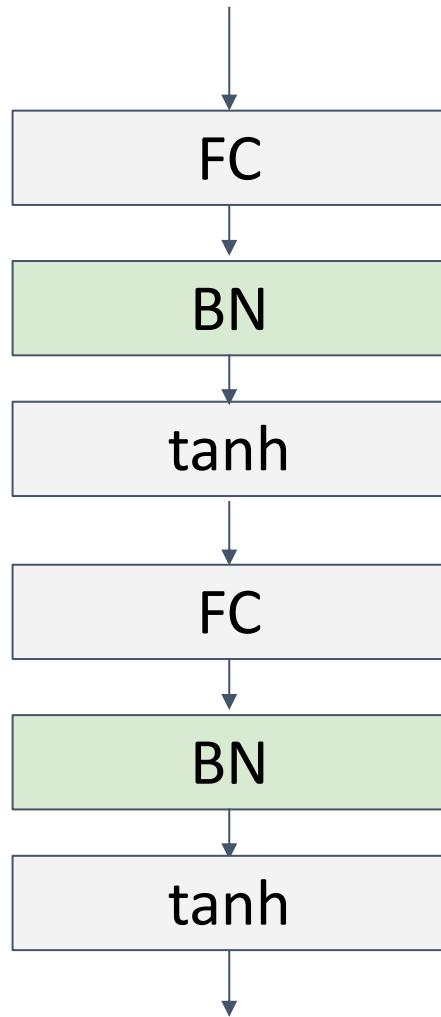


Usually inserted after Fully Connected or Convolutional layers, and before nonlinearity.

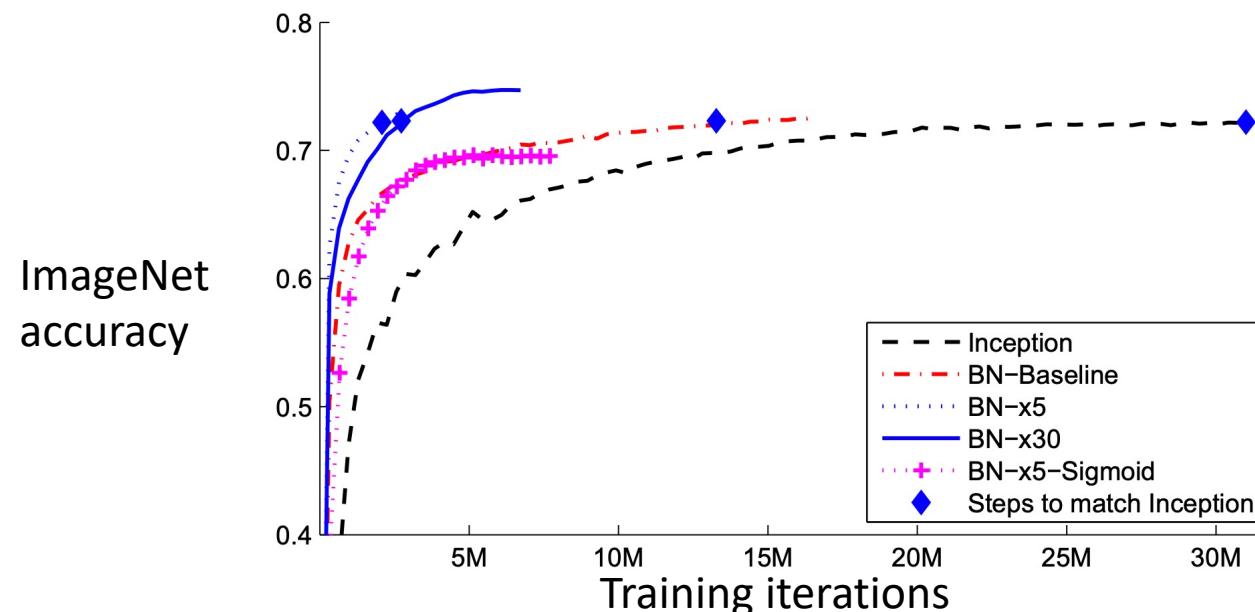
$$\hat{x} = \frac{x - E[x]}{\sqrt{Var[x]}}$$

Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

# Batch Normalization

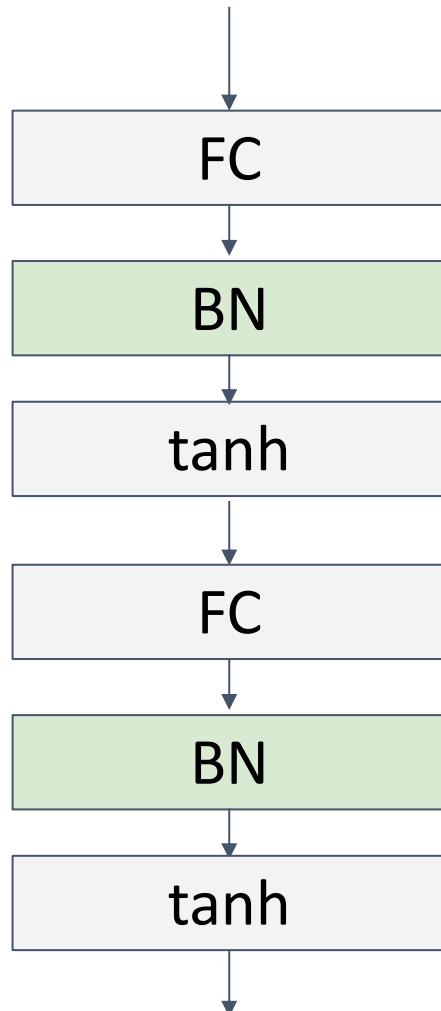


- Makes deep networks **much** easier to train!
- Allows higher learning rates, faster convergence
- Networks become more robust to initialization
- Acts as regularization during training
- Zero overhead at test-time: can be fused with conv!



Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

# Batch Normalization



- Makes deep networks **much** easier to train!
- Allows higher learning rates, faster convergence
- Networks become more robust to initialization
- Acts as regularization during training
- Zero overhead at test-time: can be fused with conv!
- Not well-understood theoretically (yet)
- Behaves differently during training and testing: this is a **very common source of bugs!**

Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

# Layer Normalization

Batch Normalization for  
**fully-connected** networks

$$x : N \times D$$

Normalize

$$\mu, \sigma : 1 \times D$$

$\gamma, \beta : 1 \times D$

$$y = \frac{(x - \mu)}{\sigma} \gamma + \beta$$

## Layer Normalization for fully-

connected networks

Same behavior at train and test!

Used in RNNs, Transformers

$$x : N \times D$$

Normalize

$$\mu, \sigma : N \times 1$$

$\gamma, \beta : 1 \times D$

$$y = \frac{(x - \mu)}{\sigma} \gamma + \beta$$

# Instance Normalization

**Batch Normalization** for convolutional networks

$$x : N \times C \times H \times W$$

Normalize

$$\mu, \sigma : 1 \times C \times 1 \times 1$$

$$\gamma, \beta : 1 \times C \times 1 \times 1$$

$$y = \frac{(x - \mu)}{\sigma} \gamma + \beta$$

**Instance Normalization** for convolutional networks

$$x : N \times C \times H \times W$$

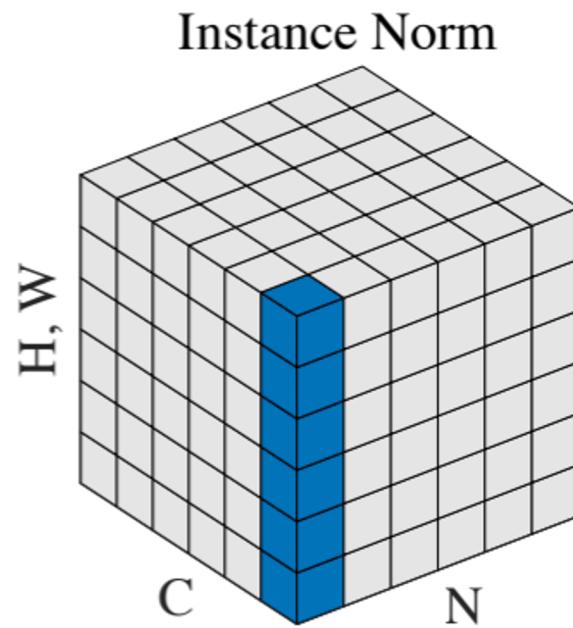
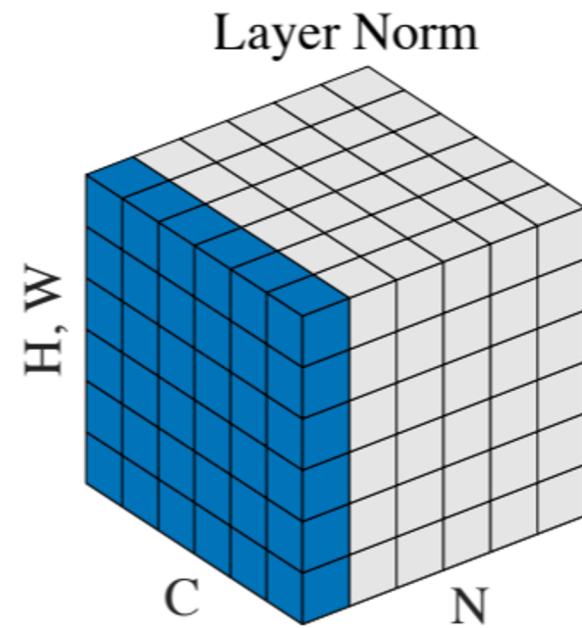
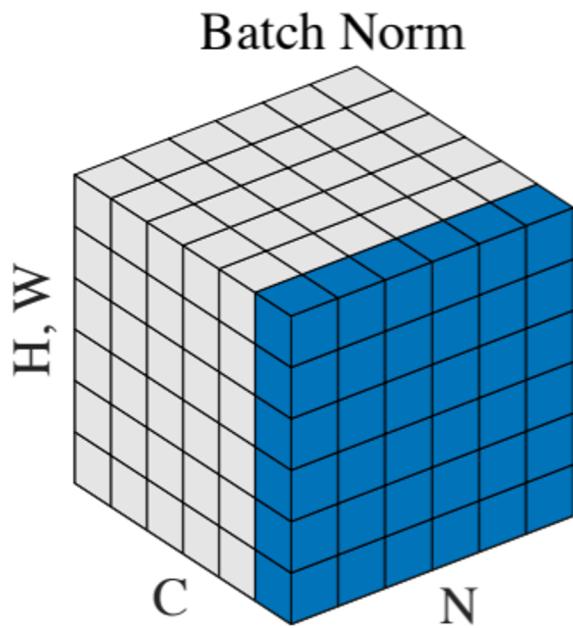
Normalize

$$\mu, \sigma : N \times C \times 1 \times 1$$

$$\gamma, \beta : 1 \times C \times 1 \times 1$$

$$y = \frac{(x - \mu)}{\sigma} \gamma + \beta$$

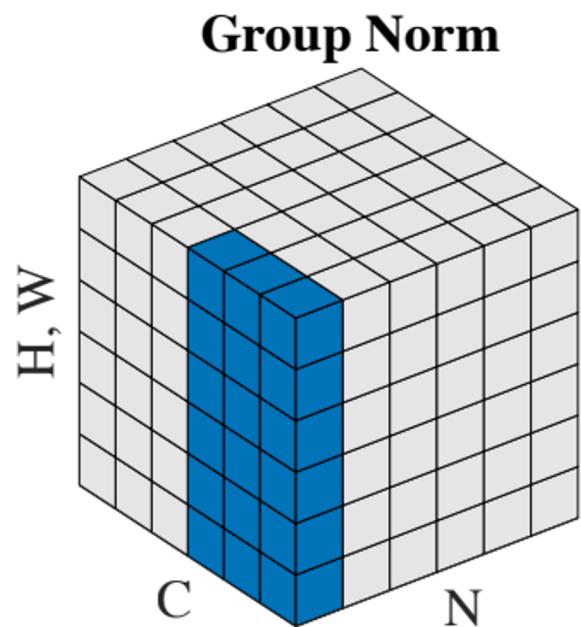
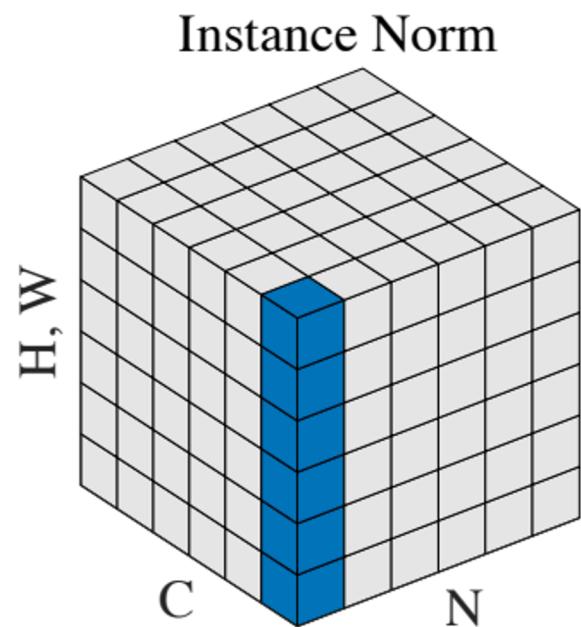
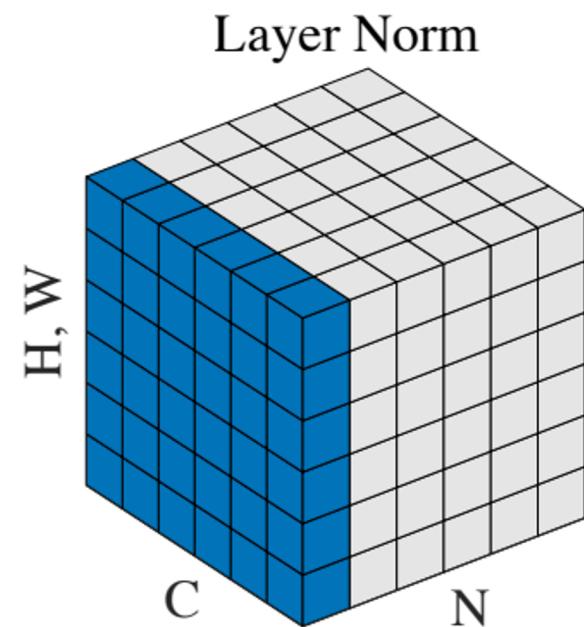
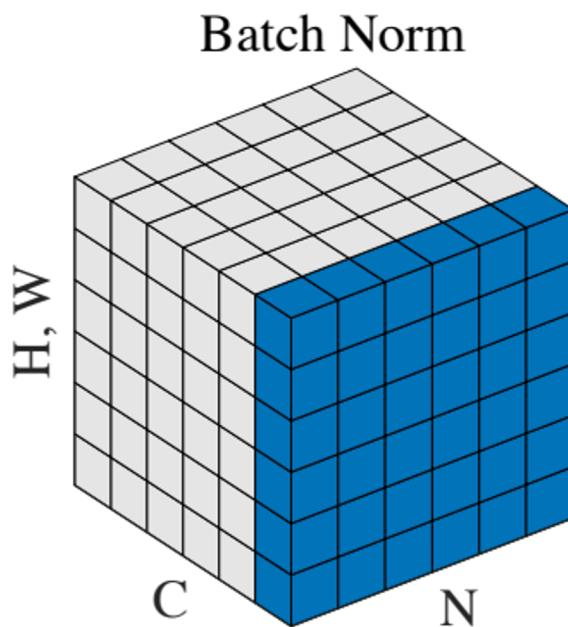
# Comparison of Normalization Layers



---

Wu and He, "Group Normalization", ECCV 2018

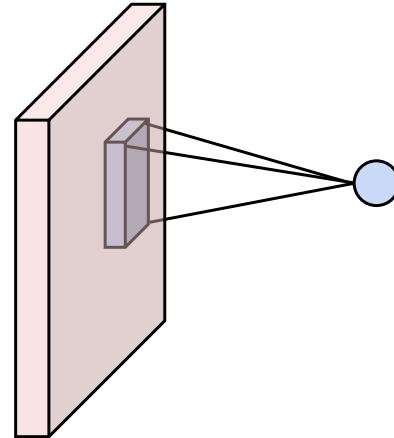
# Group Normalization



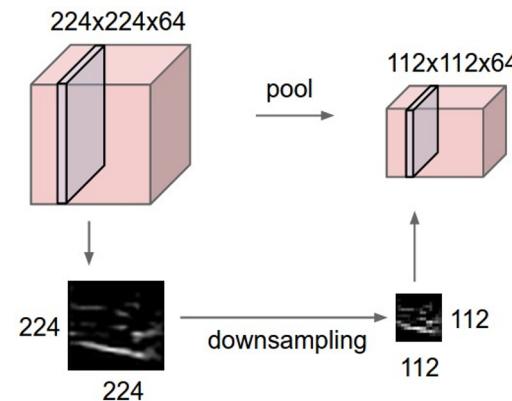
Wu and He, "Group Normalization", ECCV 2018

# Components of a Convolutional Network

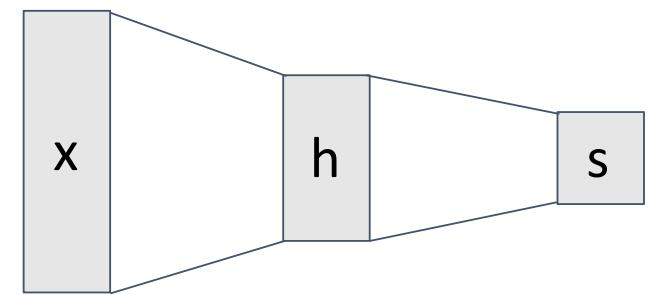
## Convolution Layers



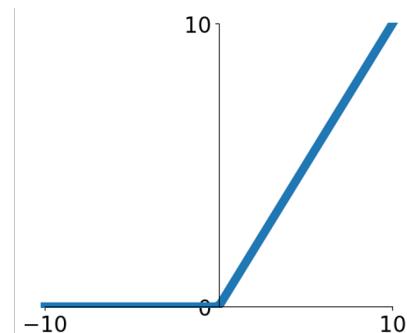
## Pooling Layers



## Fully-Connected Layers



## Activation Function

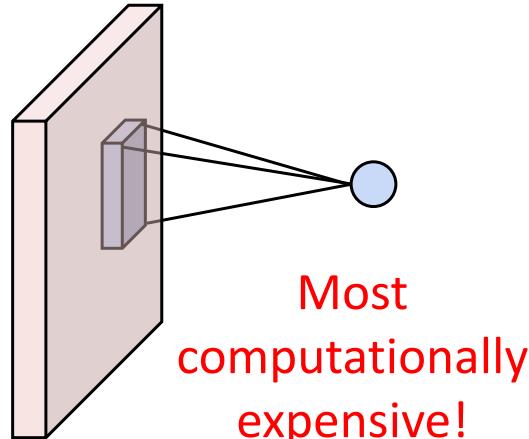


## Normalization

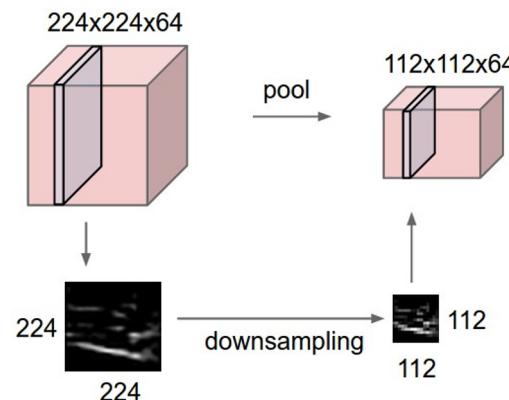
$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}}$$

# Components of a Convolutional Network

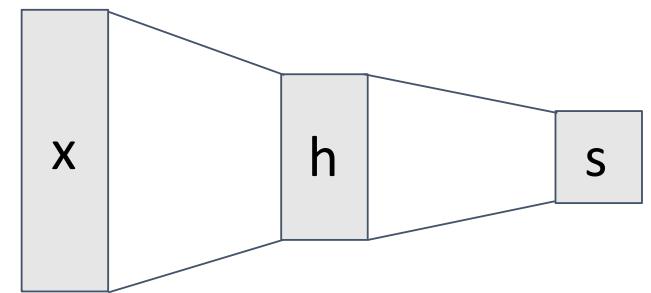
## Convolution Layers



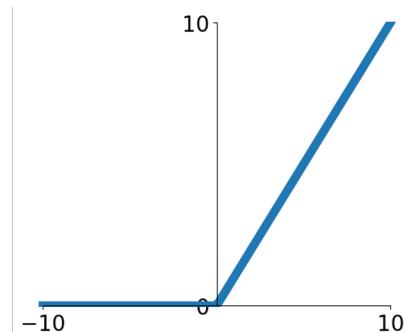
## Pooling Layers



## Fully-Connected Layers



## Activation Function

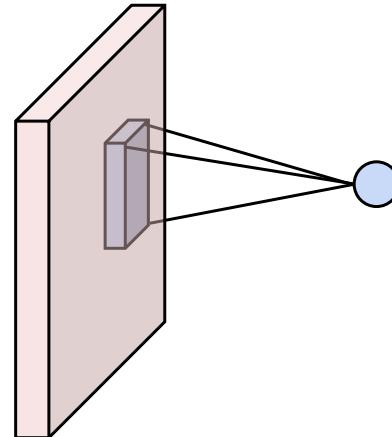


## Normalization

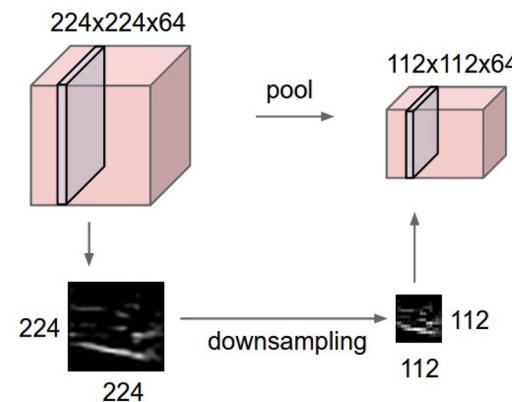
$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}}$$

# Summary: Components of a Convolutional Network

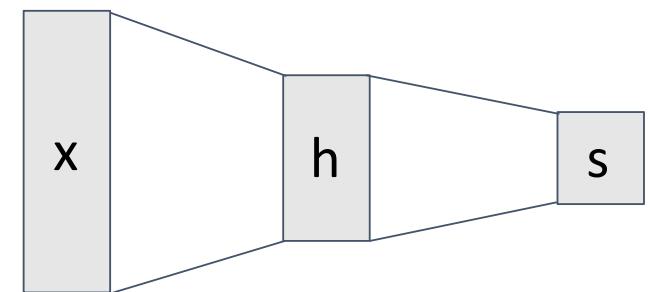
## Convolution Layers



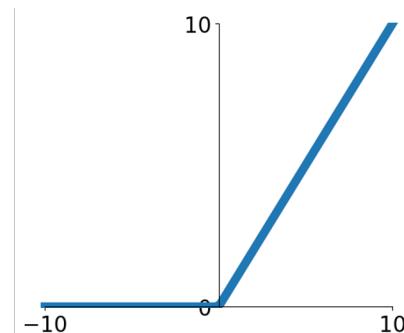
## Pooling Layers



## Fully-Connected Layers



## Activation Function

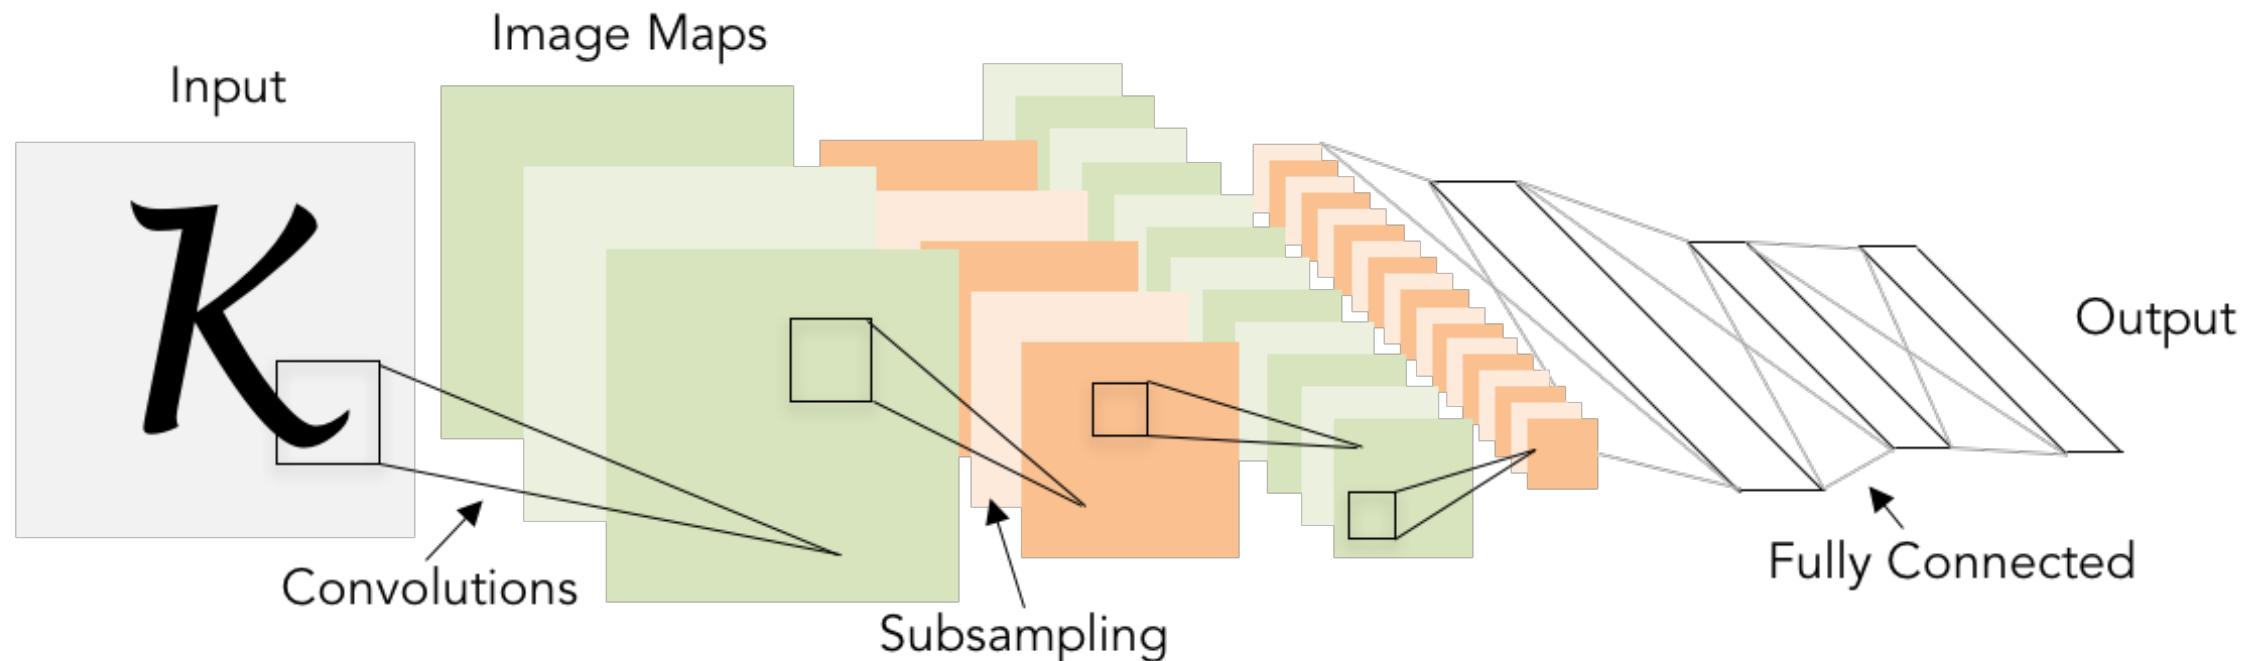


## Normalization

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}}$$

# Summary: Components of a Convolutional Network

**Problem:** What is the right way to combine all these components?



Next time:  
CNN Architectures